

# About Me

## How I became a Bitcoin Developer

- Joined the Seoul Bitcoin Meetup in 2017  
([meetup.com/seoulbitcoin](https://www.meetup.com/seoulbitcoin))

# About Me

## How I became a Bitcoin Developer

- Joined the Seoul Bitcoin Meetup in 2017  
([meetup.com/seoulbitcoin](https://www.meetup.com/seoulbitcoin))
- Started contributing in 2019

# About Me

## How I became a Bitcoin Developer

- Joined the Seoul Bitcoin Meetup in 2017  
([meetup.com/seoulbitcoin](https://www.meetup.com/seoulbitcoin))
- Started contributing in 2019
- Got funding in 2020

*"The most important book about technology today,  
with implications that go far beyond programming."*  
—Guy Kawasaki

Revised & Expanded

# THE CATHEDRAL & THE BAZAAR

MUSINGS ON LINUX AND OPEN SOURCE  
BY AN ACCIDENTAL REVOLUTIONARY



ERIC S. RAYMOND

WITH A FOREWORD BY BOB YOUNG, CHAIRMAN & CEO OF RED HAT, INC.

- No single roadmap

- No single roadmap
- Differing views on how to scale Bitcoin

- No single roadmap
- Differing views on how to scale Bitcoin
- I'll be sharing my take

- Anthony Towns - [www.erisian.com.au/wordpress/  
2023/06/21/putting-the-b-in-btc](http://www.erisian.com.au/wordpress/2023/06/21/putting-the-b-in-btc)
- James O'Berine - [jameso.be/2023/08/14/scaling-  
bitcoin.html](http://jameso.be/2023/08/14/scaling-bitcoin.html)

# **Scaling Bitcoin with Utreexo**

**Current limitations and overcoming them**

<https://github.com/kcalvinalvin/2023-12-02-Taipei-Meetup>



**“We very, very much need such a system, but  
the way I understand your proposal, it does  
not seem to scale to the required size.”**

<https://www.mail-archive.com/cryptography@metzdowd.com/msg09963.html>

# **What makes Bitcoin special**

**Something that can't  
be negotiated**

# Trustlessness

# **What is trustlessness?**

**Only trusting myself**

**“Of Bitcoin’s many properties, trustlessness, or the ability to use Bitcoin without trusting anything but the open-source software you run, is, by far, king”**

**<https://bluematt.bitcoin.ninja/2017/02/28/bitcoin-trustlessness/>**

- Decentralization
- P2P
- Self-custody

**Things that hurt  
trustlessness is rejected**

# Using Bitcoin Trustlessly

# Run a full node

Trustless usage

- Full node = fully validating



# Run a full node

## Trustless usage

- Full node = fully validating
- Validating transactions and blocks



# **Validation Process**

**What all full nodes go through**

- Download 548GB of blocks

# **Validation Process**

**What all full nodes go through**

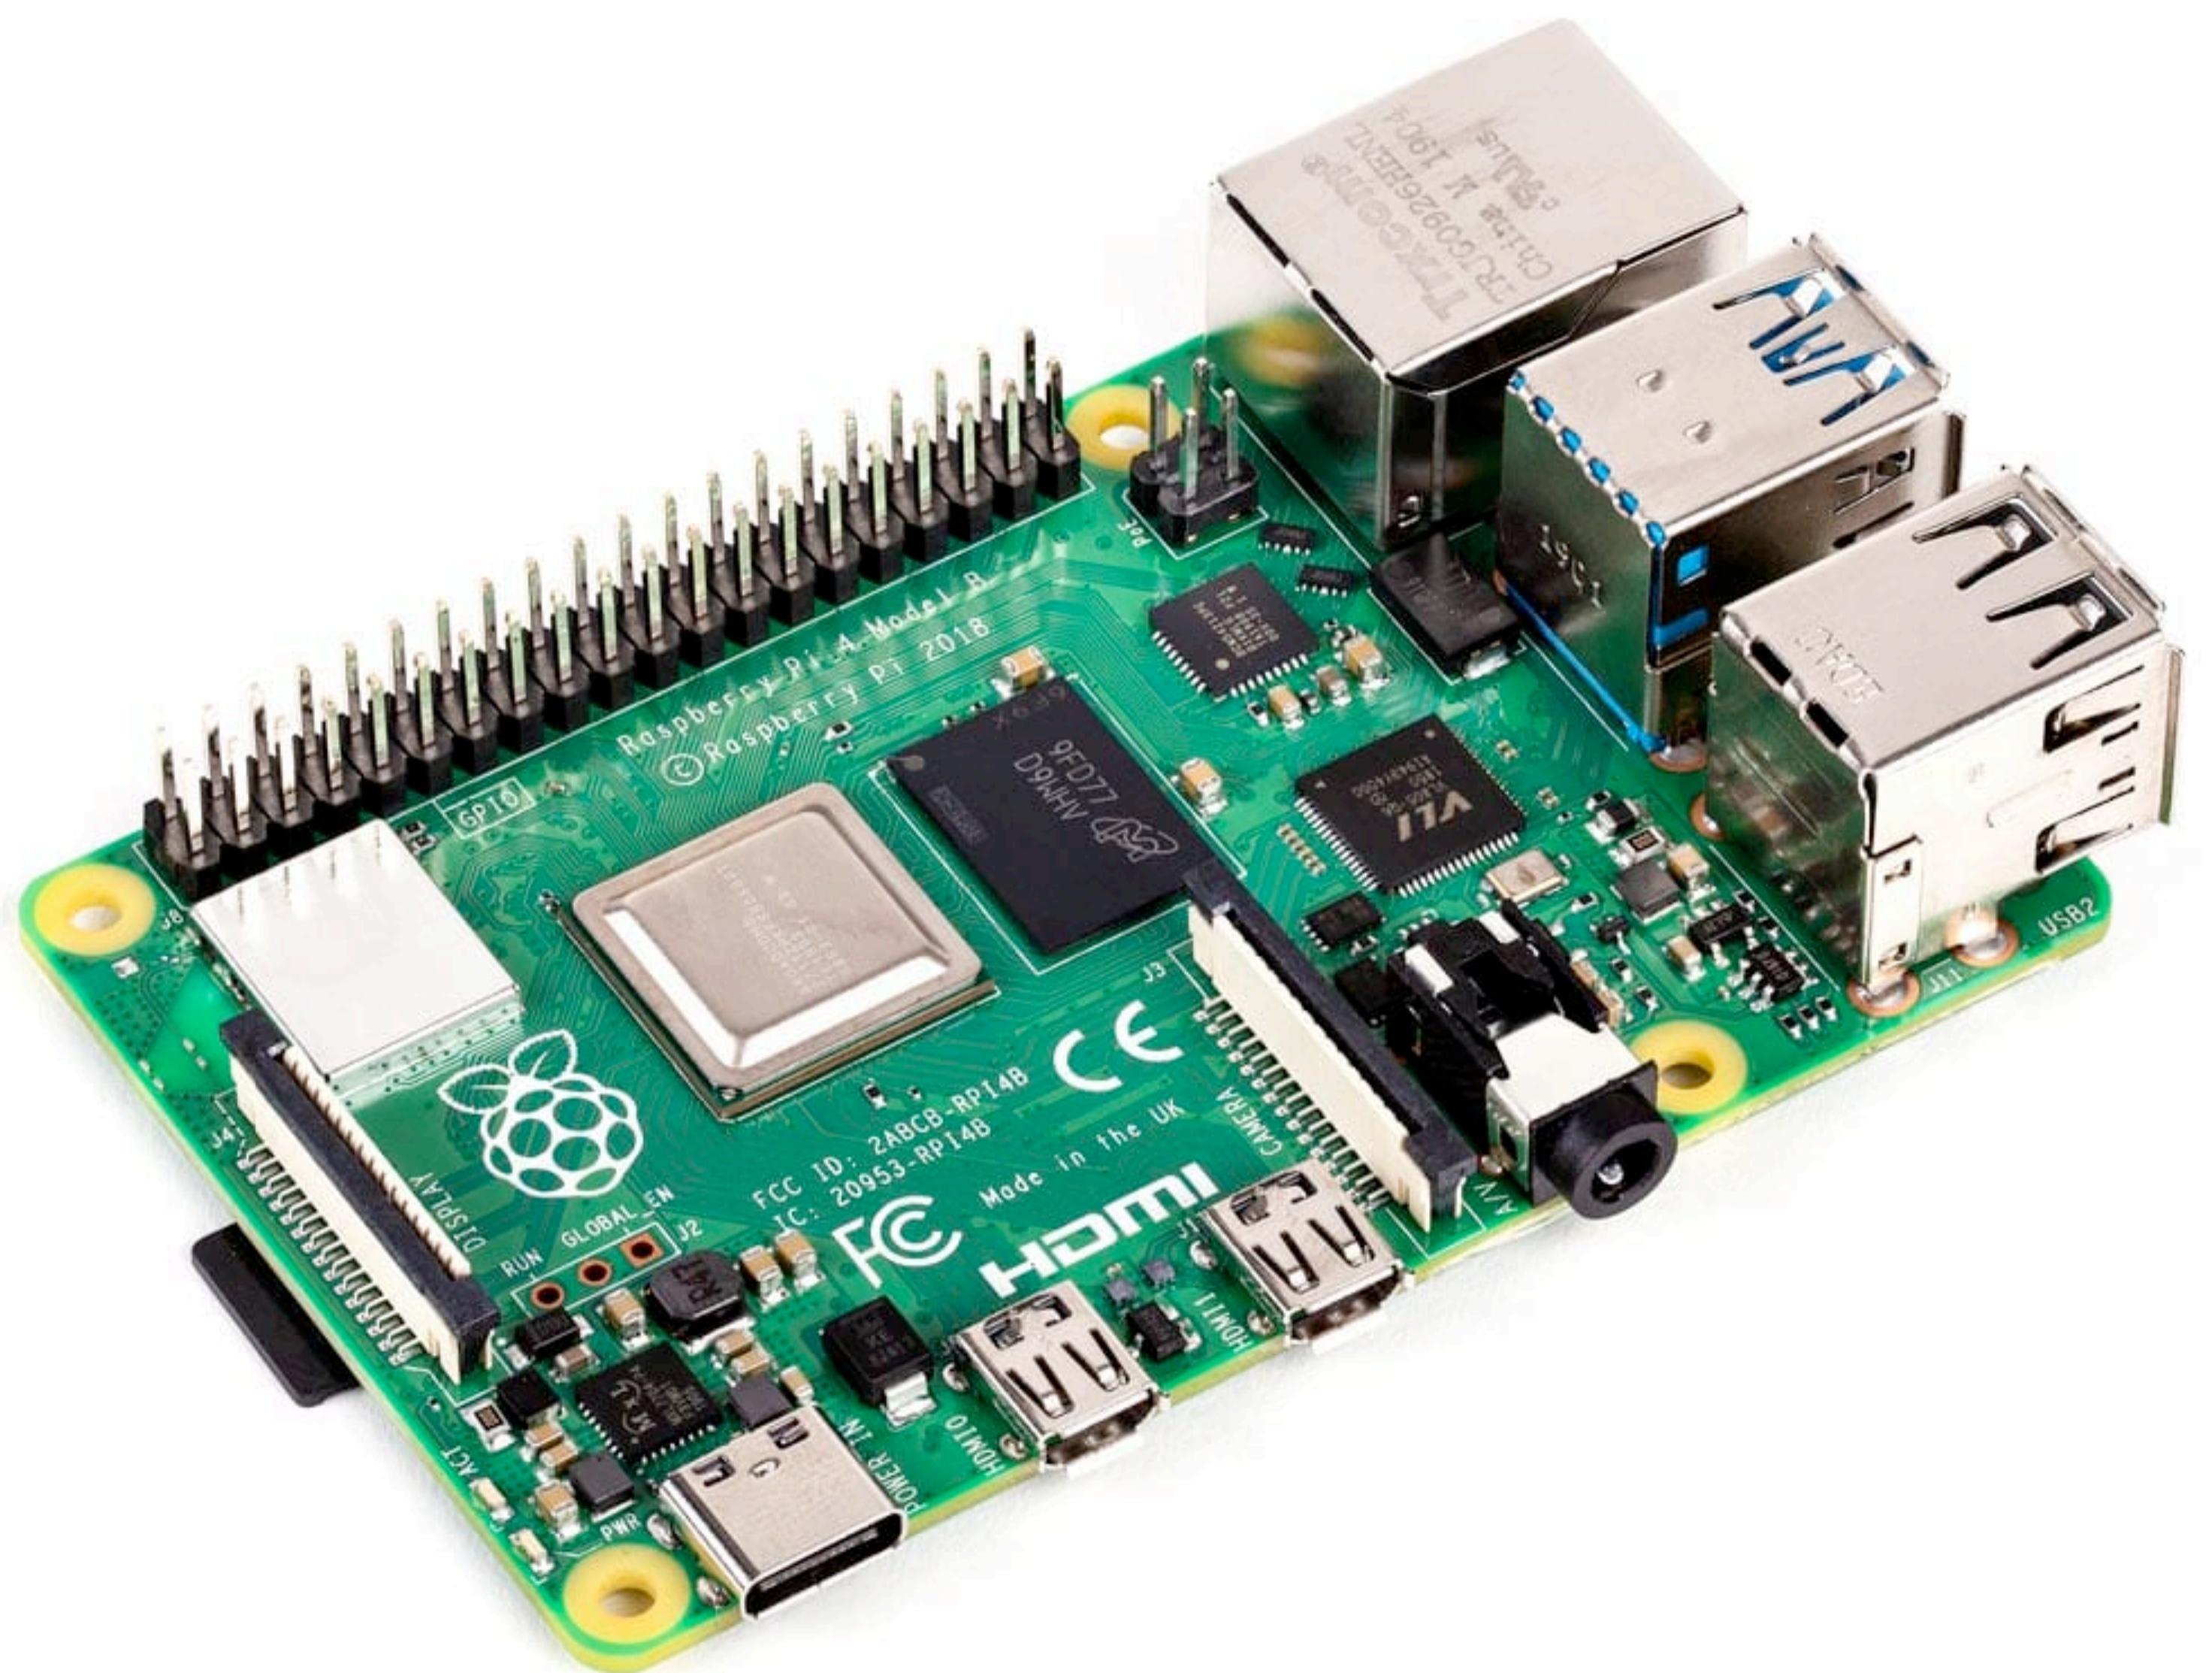
- Download 548GB of blocks
- Update the UTXO set every block

# **Validation Process**

**What all full nodes go through**

- Download 548GB of blocks
- Update the UTXO set every block
- Validate the signature of the tx

**Can everyone use it trustlessly?**



# Things limiting scalability

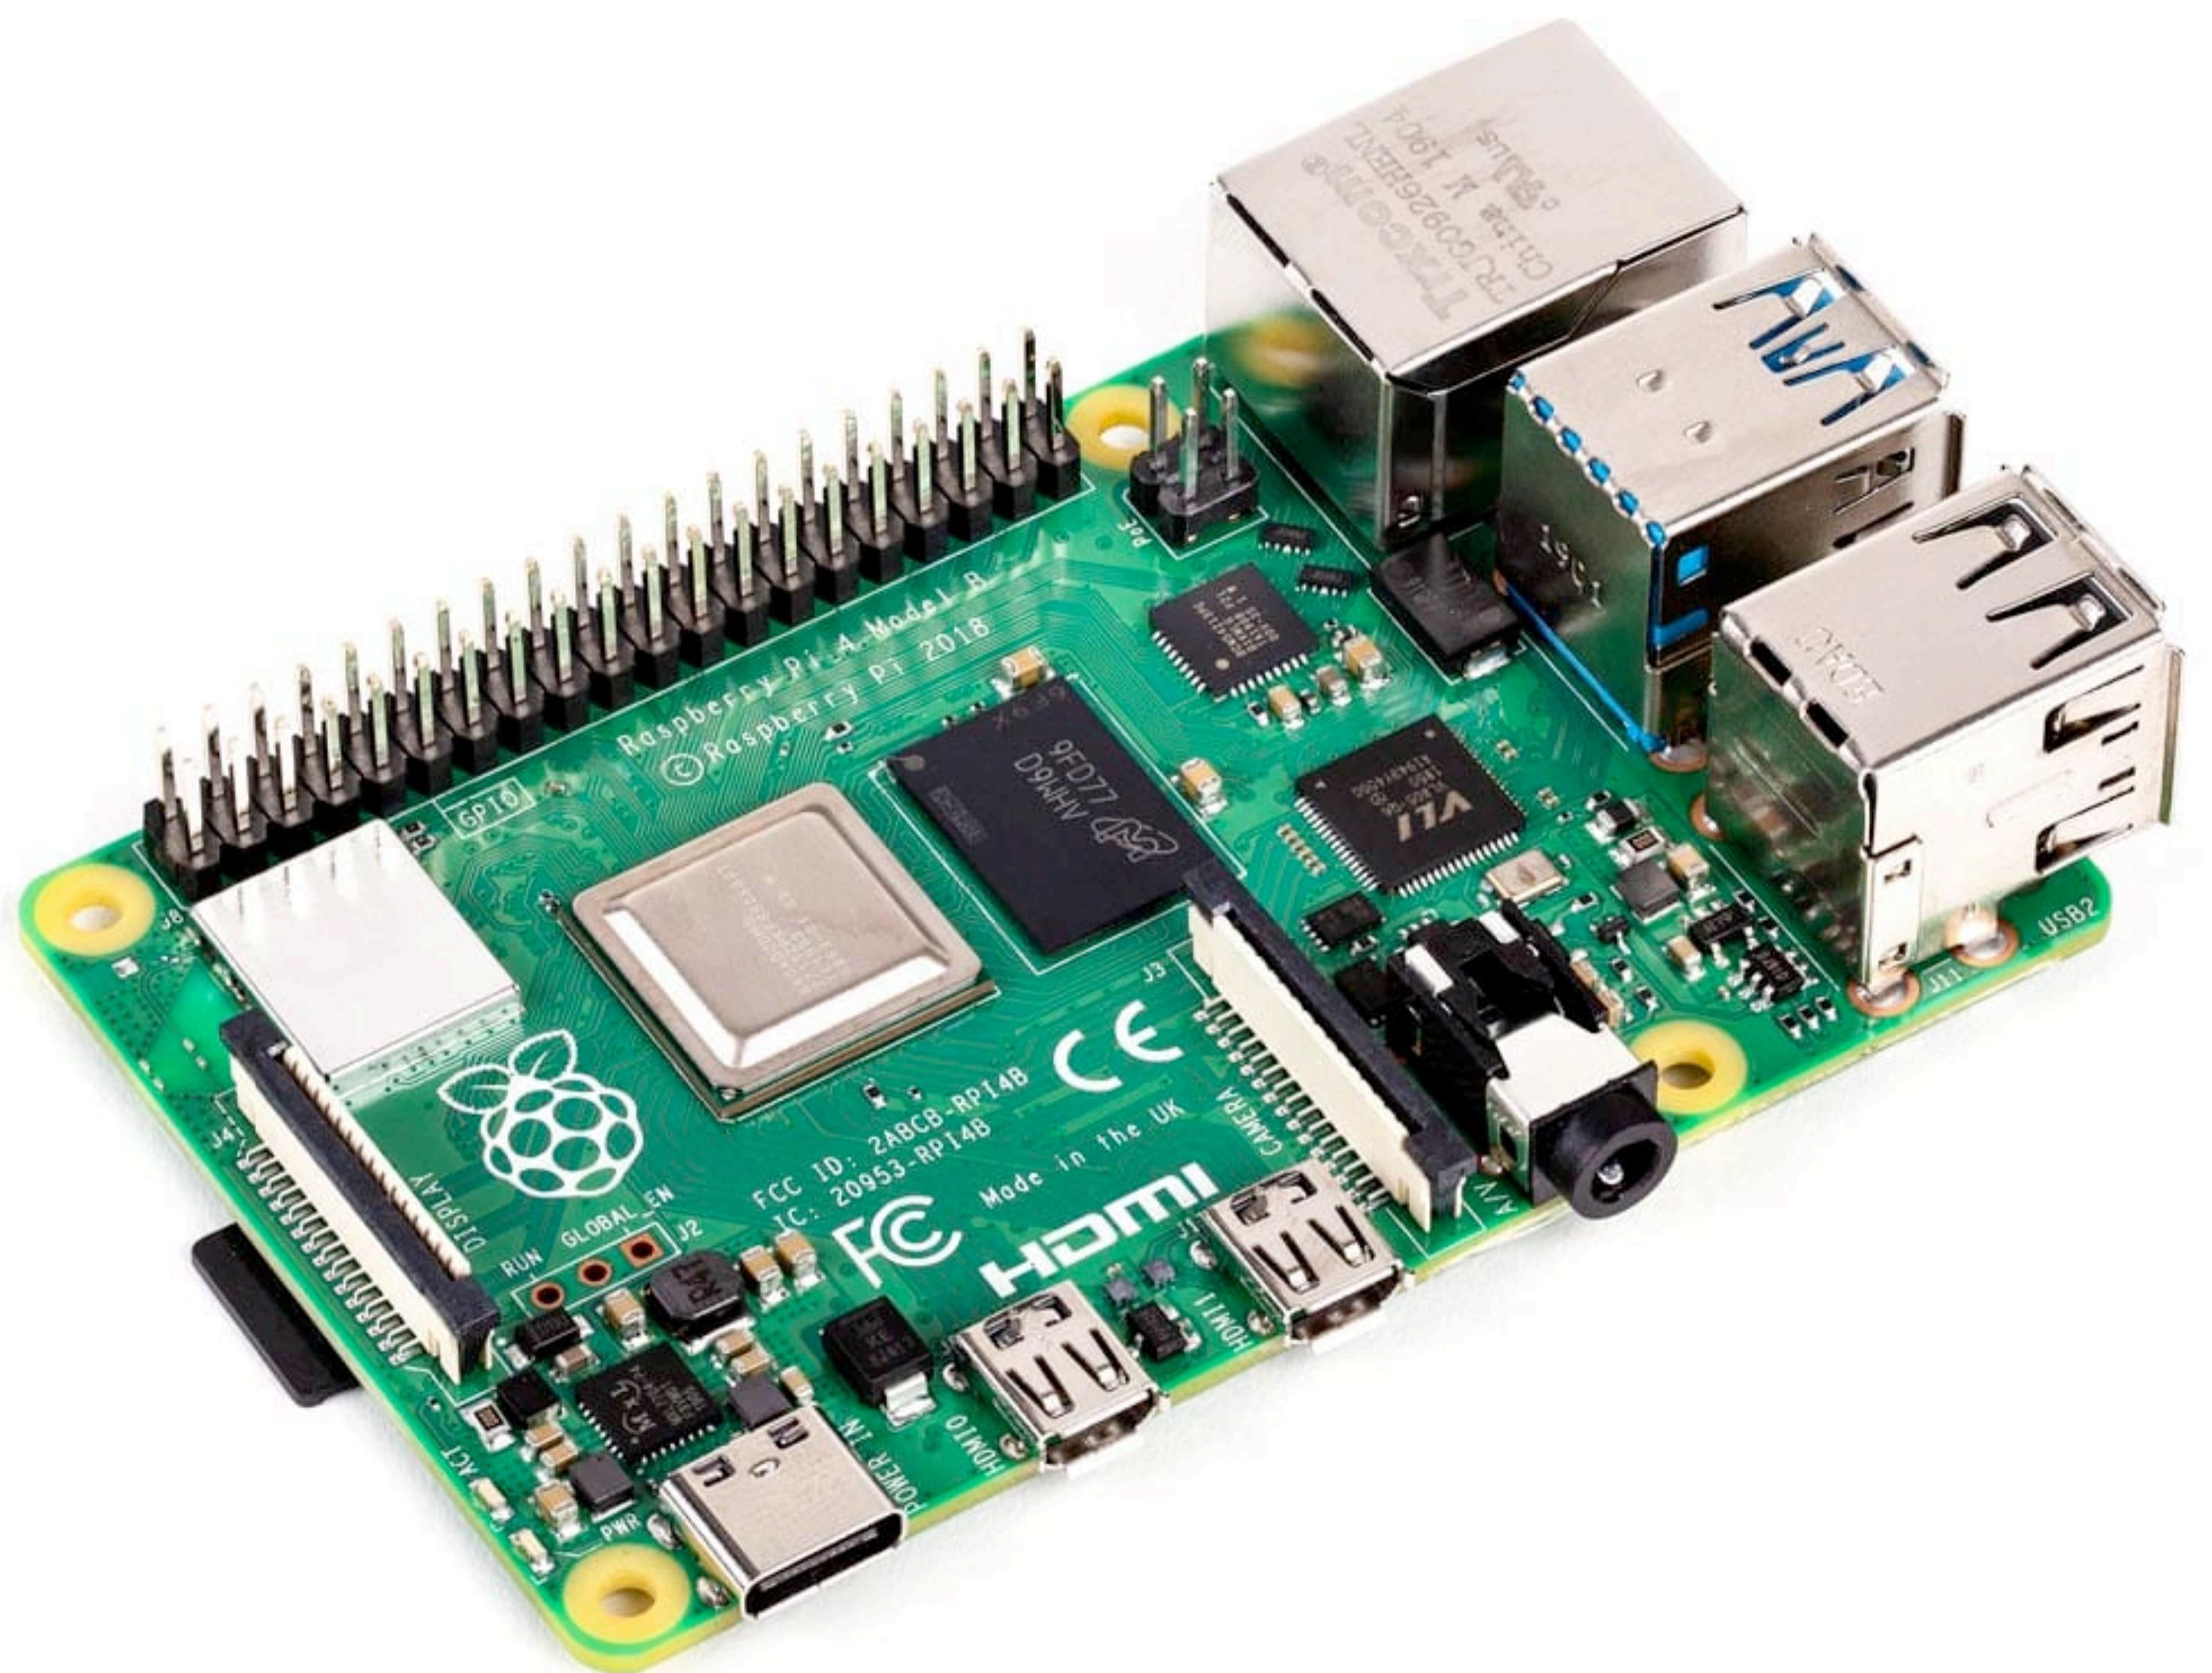
# Scalability?

**How much txs the system is able to process in a given time**

**~2MB block every  
10minutes**

**2000-3000 transaction**

**Limited by the slowest  
computer**



# **How to scale Bitcoin**

## **While keeping the trustlessness**

- Use off-chain protocols

# **How to scale Bitcoin**

## **While keeping the trustlessness**

- Use off-chain protocols
- Increase the block size

# Off-chain Protocols

## The gist

- Increase the bang-for-the-buck

# Off-chain Protocols

## The gist

- Increase the bang-for-the-buck
- Multiple transfers within a single Bitcoin tx

# Protocols with security tradeoffs

## Off-chain Protocols

- Liquid Network

# Protocols with security tradeoffs

## Off-chain Protocols

- Liquid Network
- Statechains

# Protocols with security tradeoffs

## Off-chain Protocols

- Liquid Network
- Statechains
- Fedimint/cashu

# **Protocols without security tradeoffs**

## **Off-chain Protocols**

- Lightning Network

# Lightning Network

## UX downsides during a high fee environment

- Need to make an expensive on-chain tx when getting started

# **Lightning Network**

## **UX downsides during a high fee environment**

- Need to make an expensive on-chain tx when getting started
- Need to make an expensive on-chain tx when there's liquidity issues

# Protocols without security tradeoffs

## Off-chain Protocols

- Lightning Network
- Ark

# **Ark**

## **UX downsides during a high fee environment**

- Exiting the Ark system is tricky when the fee is higher than the value of the VTXO

**High fees make L2s  
suck**

**My highly opinionated conclusion**

**Eventually increase  
the block size**

**Currently blocked by the technical reasons**

# Why can't we increase the block size

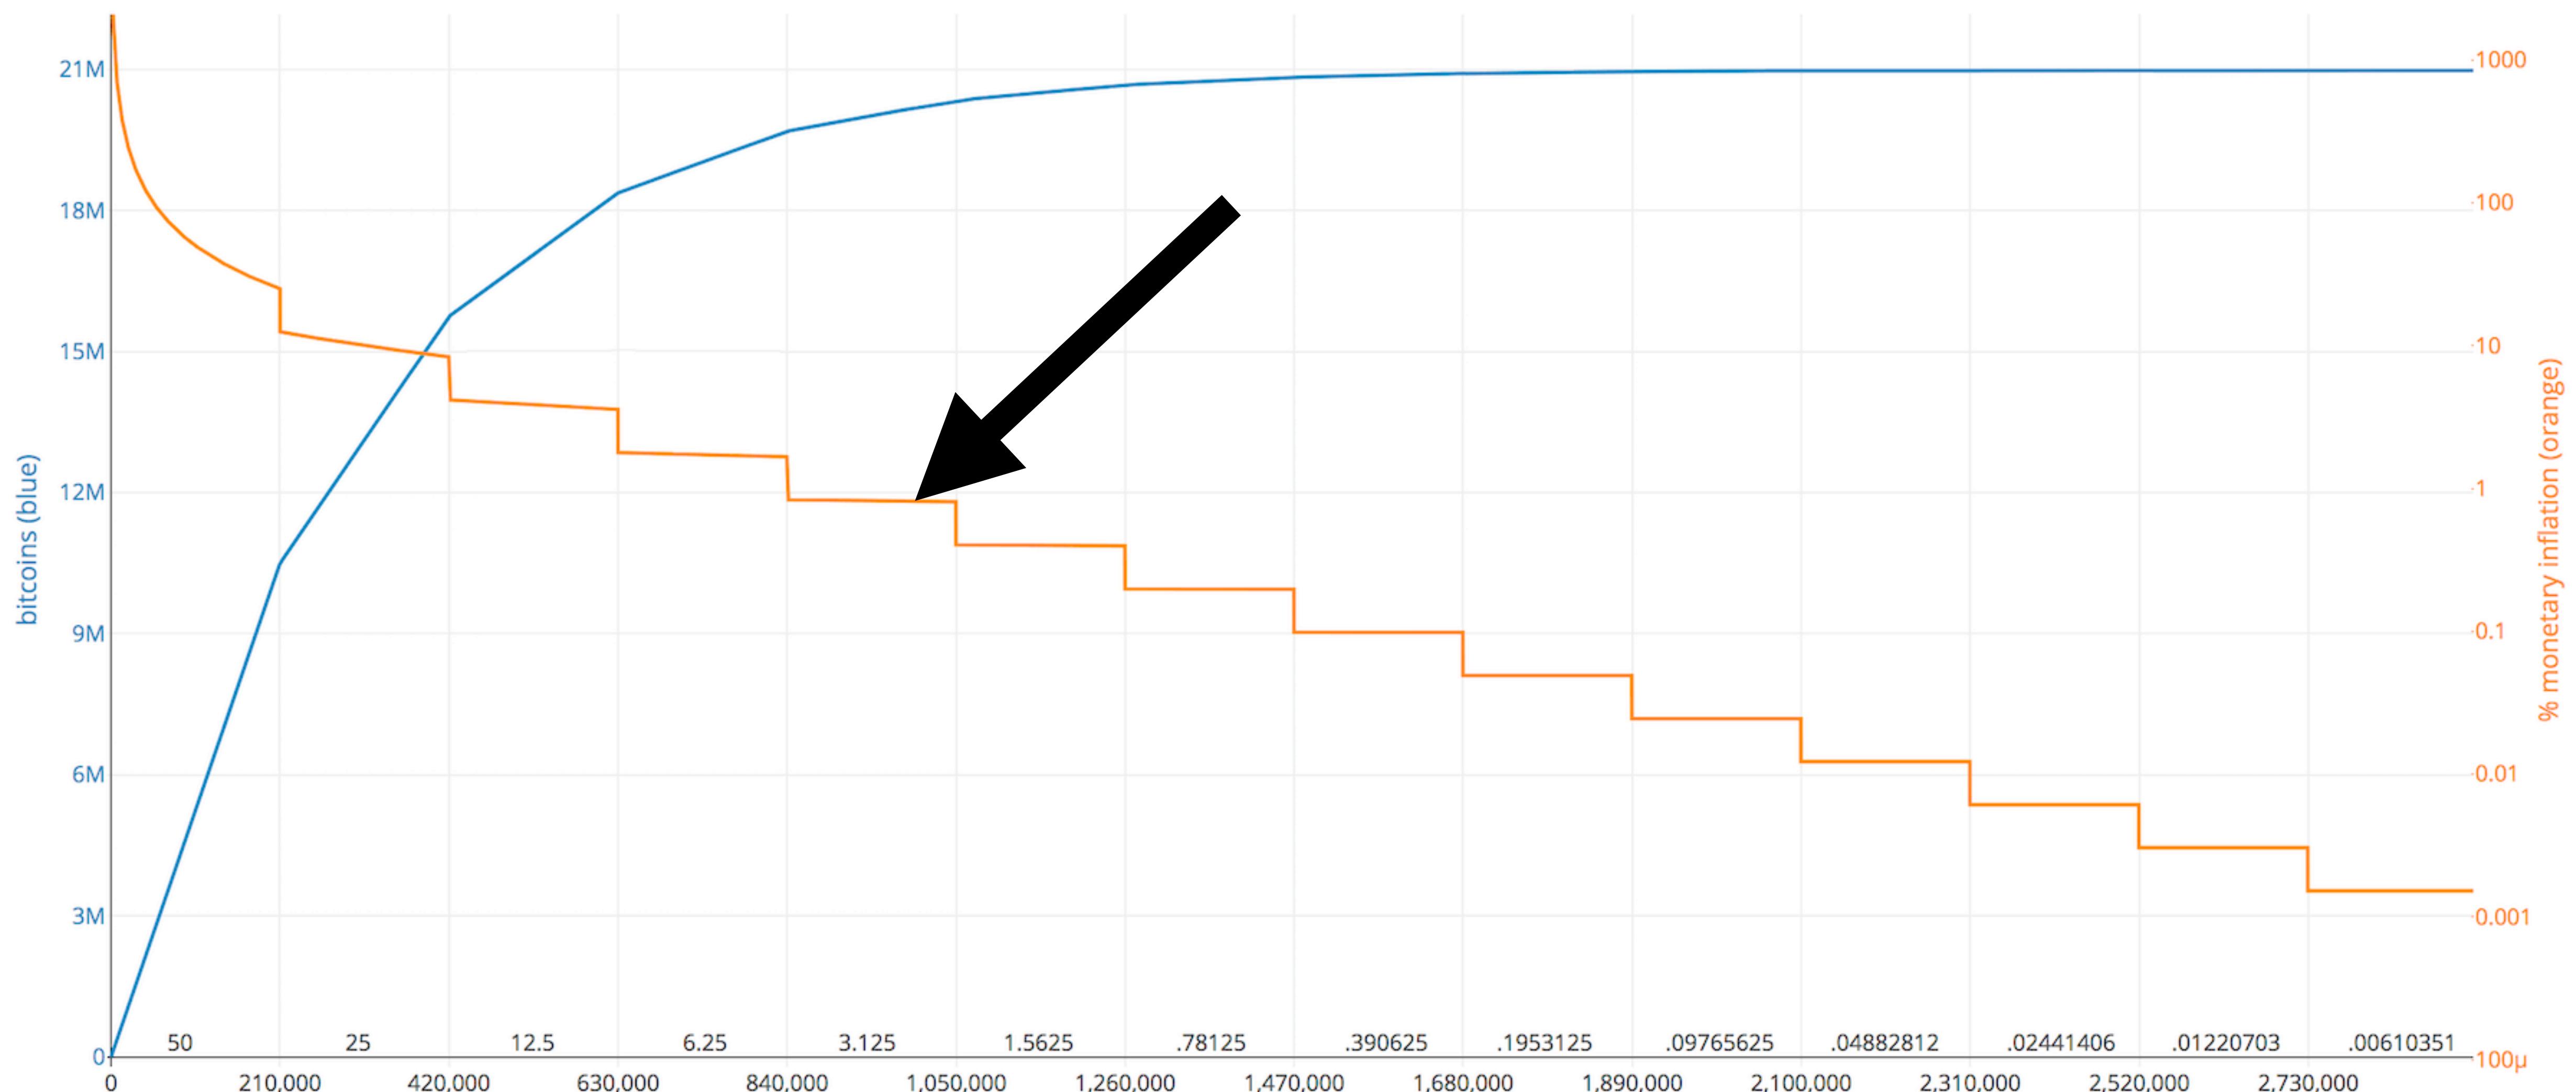
2 reasons

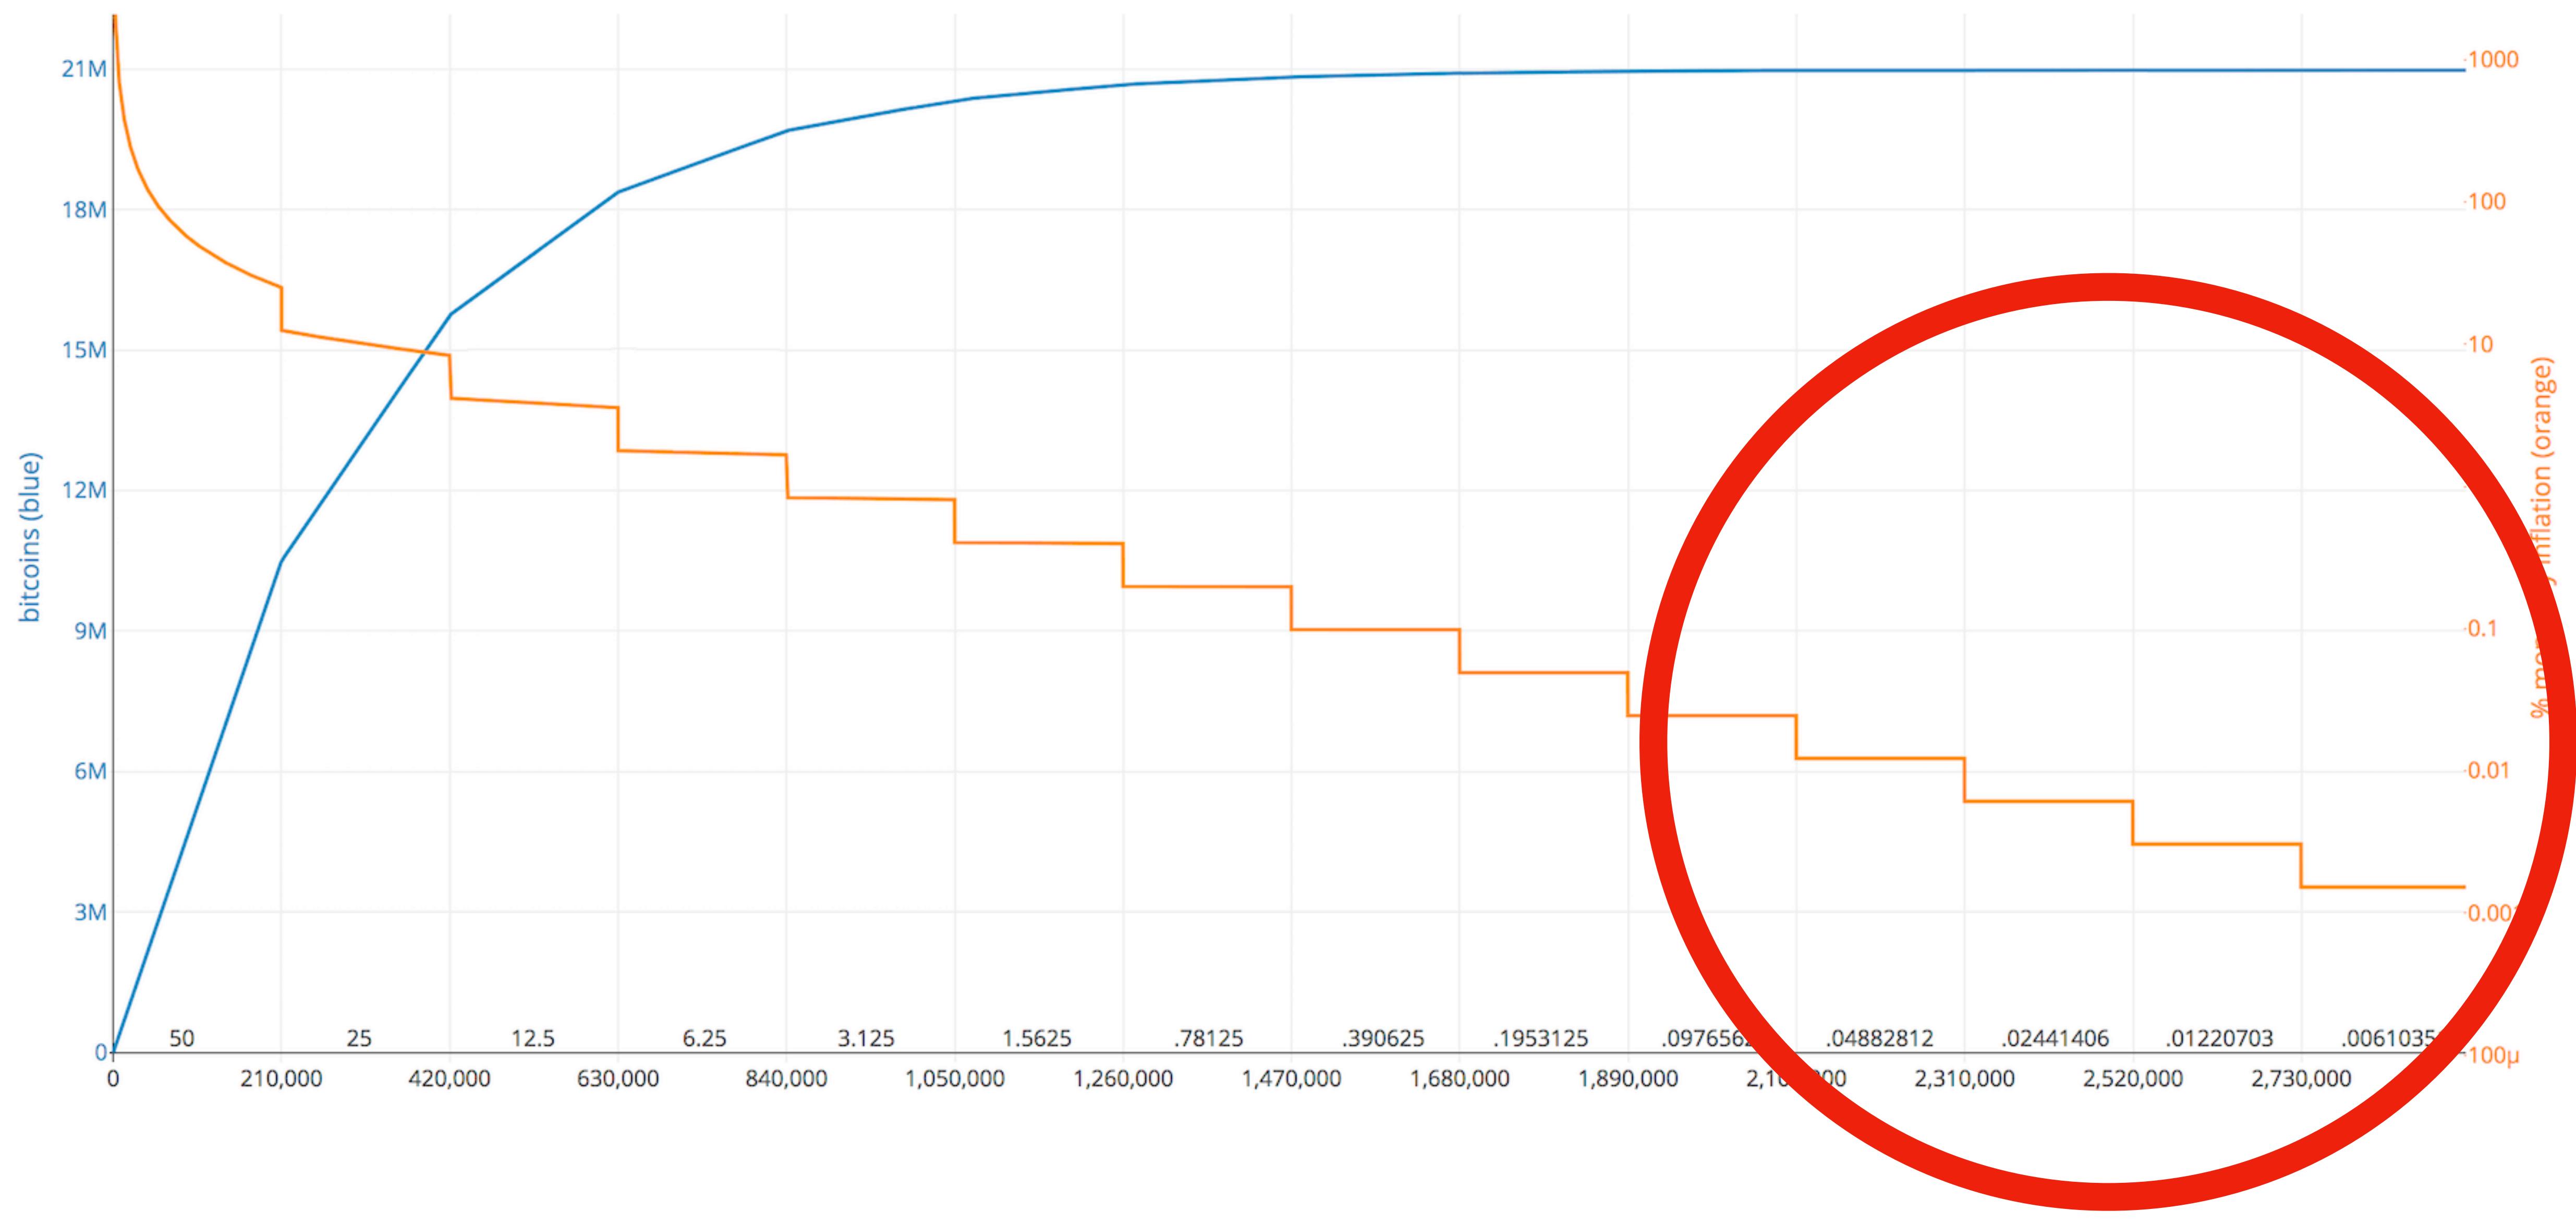
- Economical
- Technical

# Economical reasons

## Supply and demand

- Are there going to be enough fees?





# Economical reasons

## Supply and demand

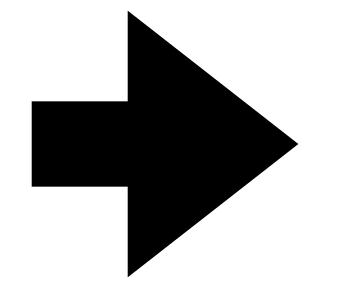
- Current double spend deterrents are mostly from the block rewards

# Economical reasons

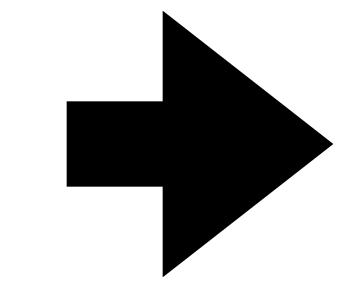
## Supply and demand

- Current double spend deterrents are mostly from the block rewards
- As rewards go down, fees have to go up

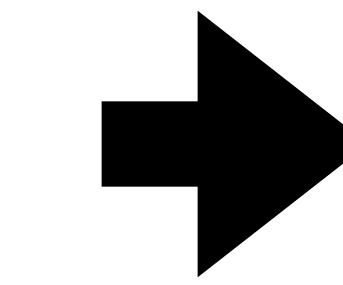
1



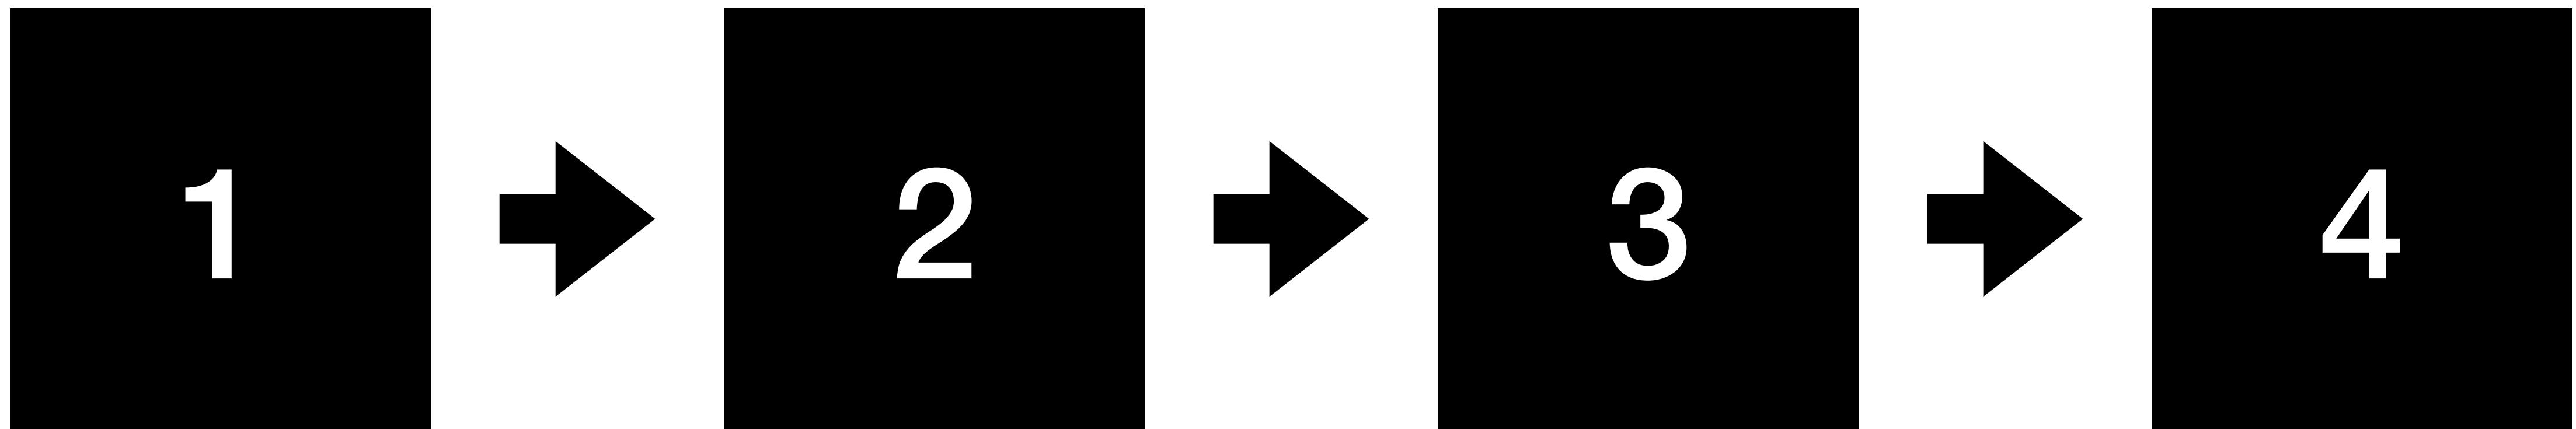
2



3



4



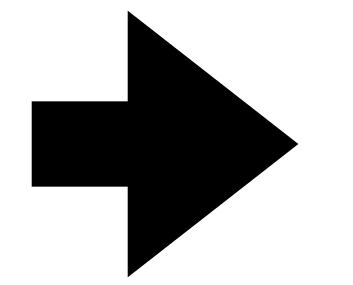
50 btc transaction

- Alice buys 50 btc with gold

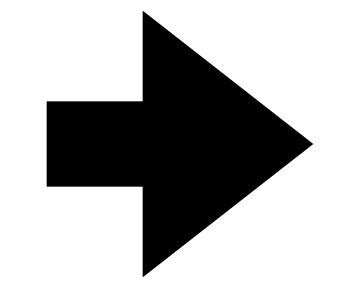
- Alice buys 50 btc with gold
- To revert the 50 btc tx, Alice bribes miner Bob with 10 btc

- Alice buys 50 btc with gold
- To revert the 50 btc tx, Alice bribes miner Bob with 10 btc
- Alice's profit - 40 btc + gold
- Bob's profit - 10 btc + block rewards

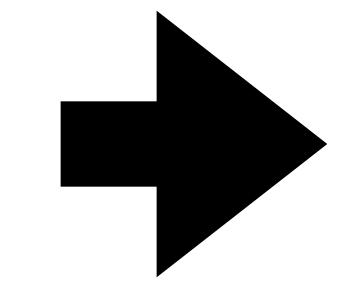
1



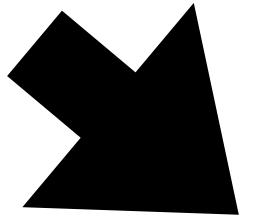
2



3

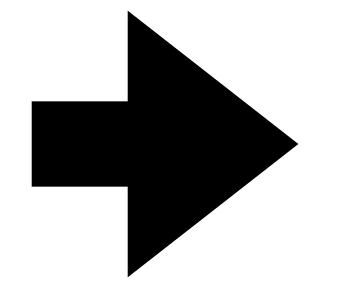


4

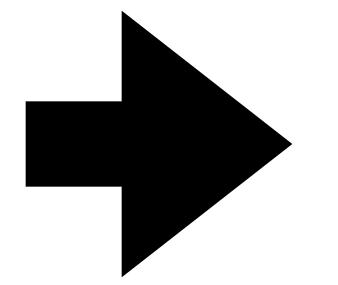


4b

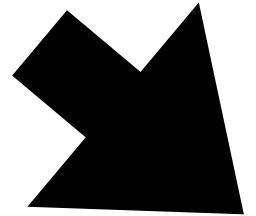
2



3

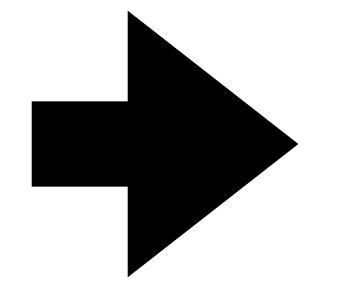


4

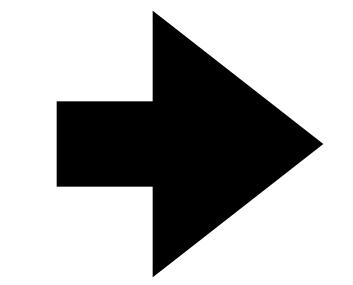


4b

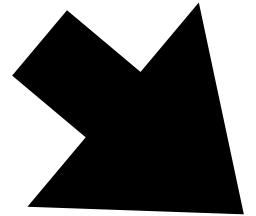
2



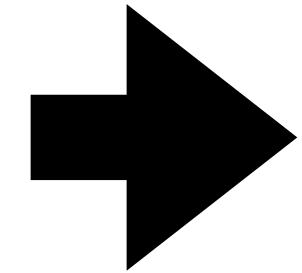
3



4



4b



5

# Economical reasons

## Supply and demand

- Enough fees need to exist to secure Bitcoin

# Economical reasons

## Supply and demand

- Enough fees need to exist to secure Bitcoin
- Blocksize ↑, Fees ↓

No research ongoing for  
economic reasons

# **Technical reasons**

## **Trustlessness**

- Can everyone use Bitcoin trustlessly?

# Technical reasons

## Trustlessness

- Can everyone use Bitcoin trustlessly?
- Blocksize ↑, Cost of running full nodes ↑

# Technical reasons

## Process of Validation

- Download ~548GB of data
- Update the UTXO set per block
- Validate the signature of the tx

# **Technical reasons**

## **Process of Validation**

- Internet
- Disk
- CPU

# **Internet**

## **Block / Transaction Validation**

- Download ~500GB on the first start

# **Internet**

## **Block / Transaction Validation**

- Download ~500GB on the first start
- 100mbps Internet = ~666 minutes

# **Internet**

## **Block / Transaction Validation**

- Download ~500GB on the first start
- 100mbps Internet = ~666 minutes
- **1gbps = ~66 minutes**

# **Internet**

## **Block / Transaction Validation**

- Download ~500GB on the first start
- 100mbps Internet = ~666 minutes
- 1gbps = ~66 minutes
- 10gbps = ~6 minutes

**Internet speeds are enough**

# Disk

## UTXO set read/writes

- UTXO set - All the Bitcoin available to be spent

# What's a UTXO set?

- TXO - Transaction Output

# Typical TX

[8ad63210d059908229f00ad177469d9f16a5d37a7ea64cede9c92684ff7cf06e](#)

DETAILS



#0 [f758ad929fbefac99ef0d1f8b13990fa75efb517ef6c919178c8af7308f145](#) 0.04852489 BTC  
07:1

#0 [17dSwJytZBszAafyWrK8Pue7rsRh9s5xeJ](#) 0.04033963 BTC

#1 [bc1qwqdg6squ38e46795at95yu9atm8azzmyvckulcc7kytlcckxswvvzej](#) 0.00758526 BTC

1 CONFIRMATION 0.04792489 BTC

# Outputs of a TX

[8ad63210d059908229f00ad177469d9f16a5d37a7ea64cede9c92684ff7cf06e](#)

DETAILS



#0 [f758ad929fbefac99ef0d1f8b13990fa75efb517ef6c919178c8af7308f145](#) 0.04852489 BTC  
07:1

#0 [17dSwJytZBszAafyWrK8Pue7rsRh9s5xeJ](#) 0.04033963 BTC

#1 [bc1qwqdg6squ38e46795at95yu9atm8azzmyvckulcc7kytlcckxswvvzej](#) 0.00758526 BTC

1 CONFIRMATION 0.04792489 BTC

# What's a UTXO set?

- TXO - Transaction Output
- UTXO - Unspent TXO. Not yet spent TXO

# What's a UTXO set?

- TXO - Transaction Output
- UTXO - Unspent TXO. Not yet spent TXO
- UTXO set - All the UTXOs

# Disk

## UTXO set read/writes

- UTXO set is stored as a key-value database (leveldb)

# Disk

## UTXO set read/writes

- UTXO set is stored as a key-value database (leveldb)
- leveldb is random read/write

# Disk

## UTXO set read/writes

- Minimum 59B per UTXO

# Disk

## UTXO set read/writes

- Minimum 59B per UTXO
- Typical tx is 1 input, 2 output

# Disk

## UTXO set read/writes

- Minimum 59B per UTXO
- Typical tx is 1 input, 2 output
- For 2000txs: ~2000 read, ~4000 write

기술의 경이로움이 시작된다

## Platinum P41 SSD

科技的奇蹟開始了



SK하이닉스 Platinum P41 SSD는 업계 최고 수준의 성능으로  
차세대 컴퓨팅 시스템을 완성하는 제품입니다. PCIe 4세대 NVMe 인터페이스와  
176단 NAND의 한층 더 강화된 기술력을 토대로 최강의 데이터 전송 속도와  
성능을 제공합니다. 이제 차원이 다른 성능을 경험해 보세요.

## 성능

---

2TB

순차성능  
(최대)

읽기 : 7,000 MB/s  
쓰기 : 6,500 MB/s

랜덤성능  
(최대)

읽기 : 1,400K IOPS  
쓰기 : 1,300K IOPS

TBW

최대 1,200TBW

1TB

읽기 : 7,000 MB/s  
쓰기 : 6,500 MB/s

읽기 : 1,400K IOPS  
쓰기 : 1,300K IOPS

최대 750TBW

500GB

읽기 : 7,000 MB/s  
쓰기 : 4,700 MB/s

읽기 : 960K IOPS  
쓰기 : 1,000K IOPS

최대 500TBW

- 1,400 IOPS = ~5.6MB/s read
- 1,300 IOPS = ~5.2MB/s write

(1,400 \* 4) / 1000. 4 is the OS page size (mb)

- 1,400 IOPS = ~5.6MB/s read
- 1,300 IOPS = ~5.2MB/s write
- Per block, 0.11MB read, 0.23MB write

$$(1,400 * 4) / 1000. \text{ 4 is the OS page size (mb)}$$



- 73 IOPS = ~0.29MB/s read&write
- Per block 0.11MB read, 0.23MB write

# Caching

Current way of reducing the disk burden

- Write is fast (only to RAM)

# Caching

**Current way of reducing the disk burden**

- Write is fast (only to RAM)
- Read is sped-up (if cached)

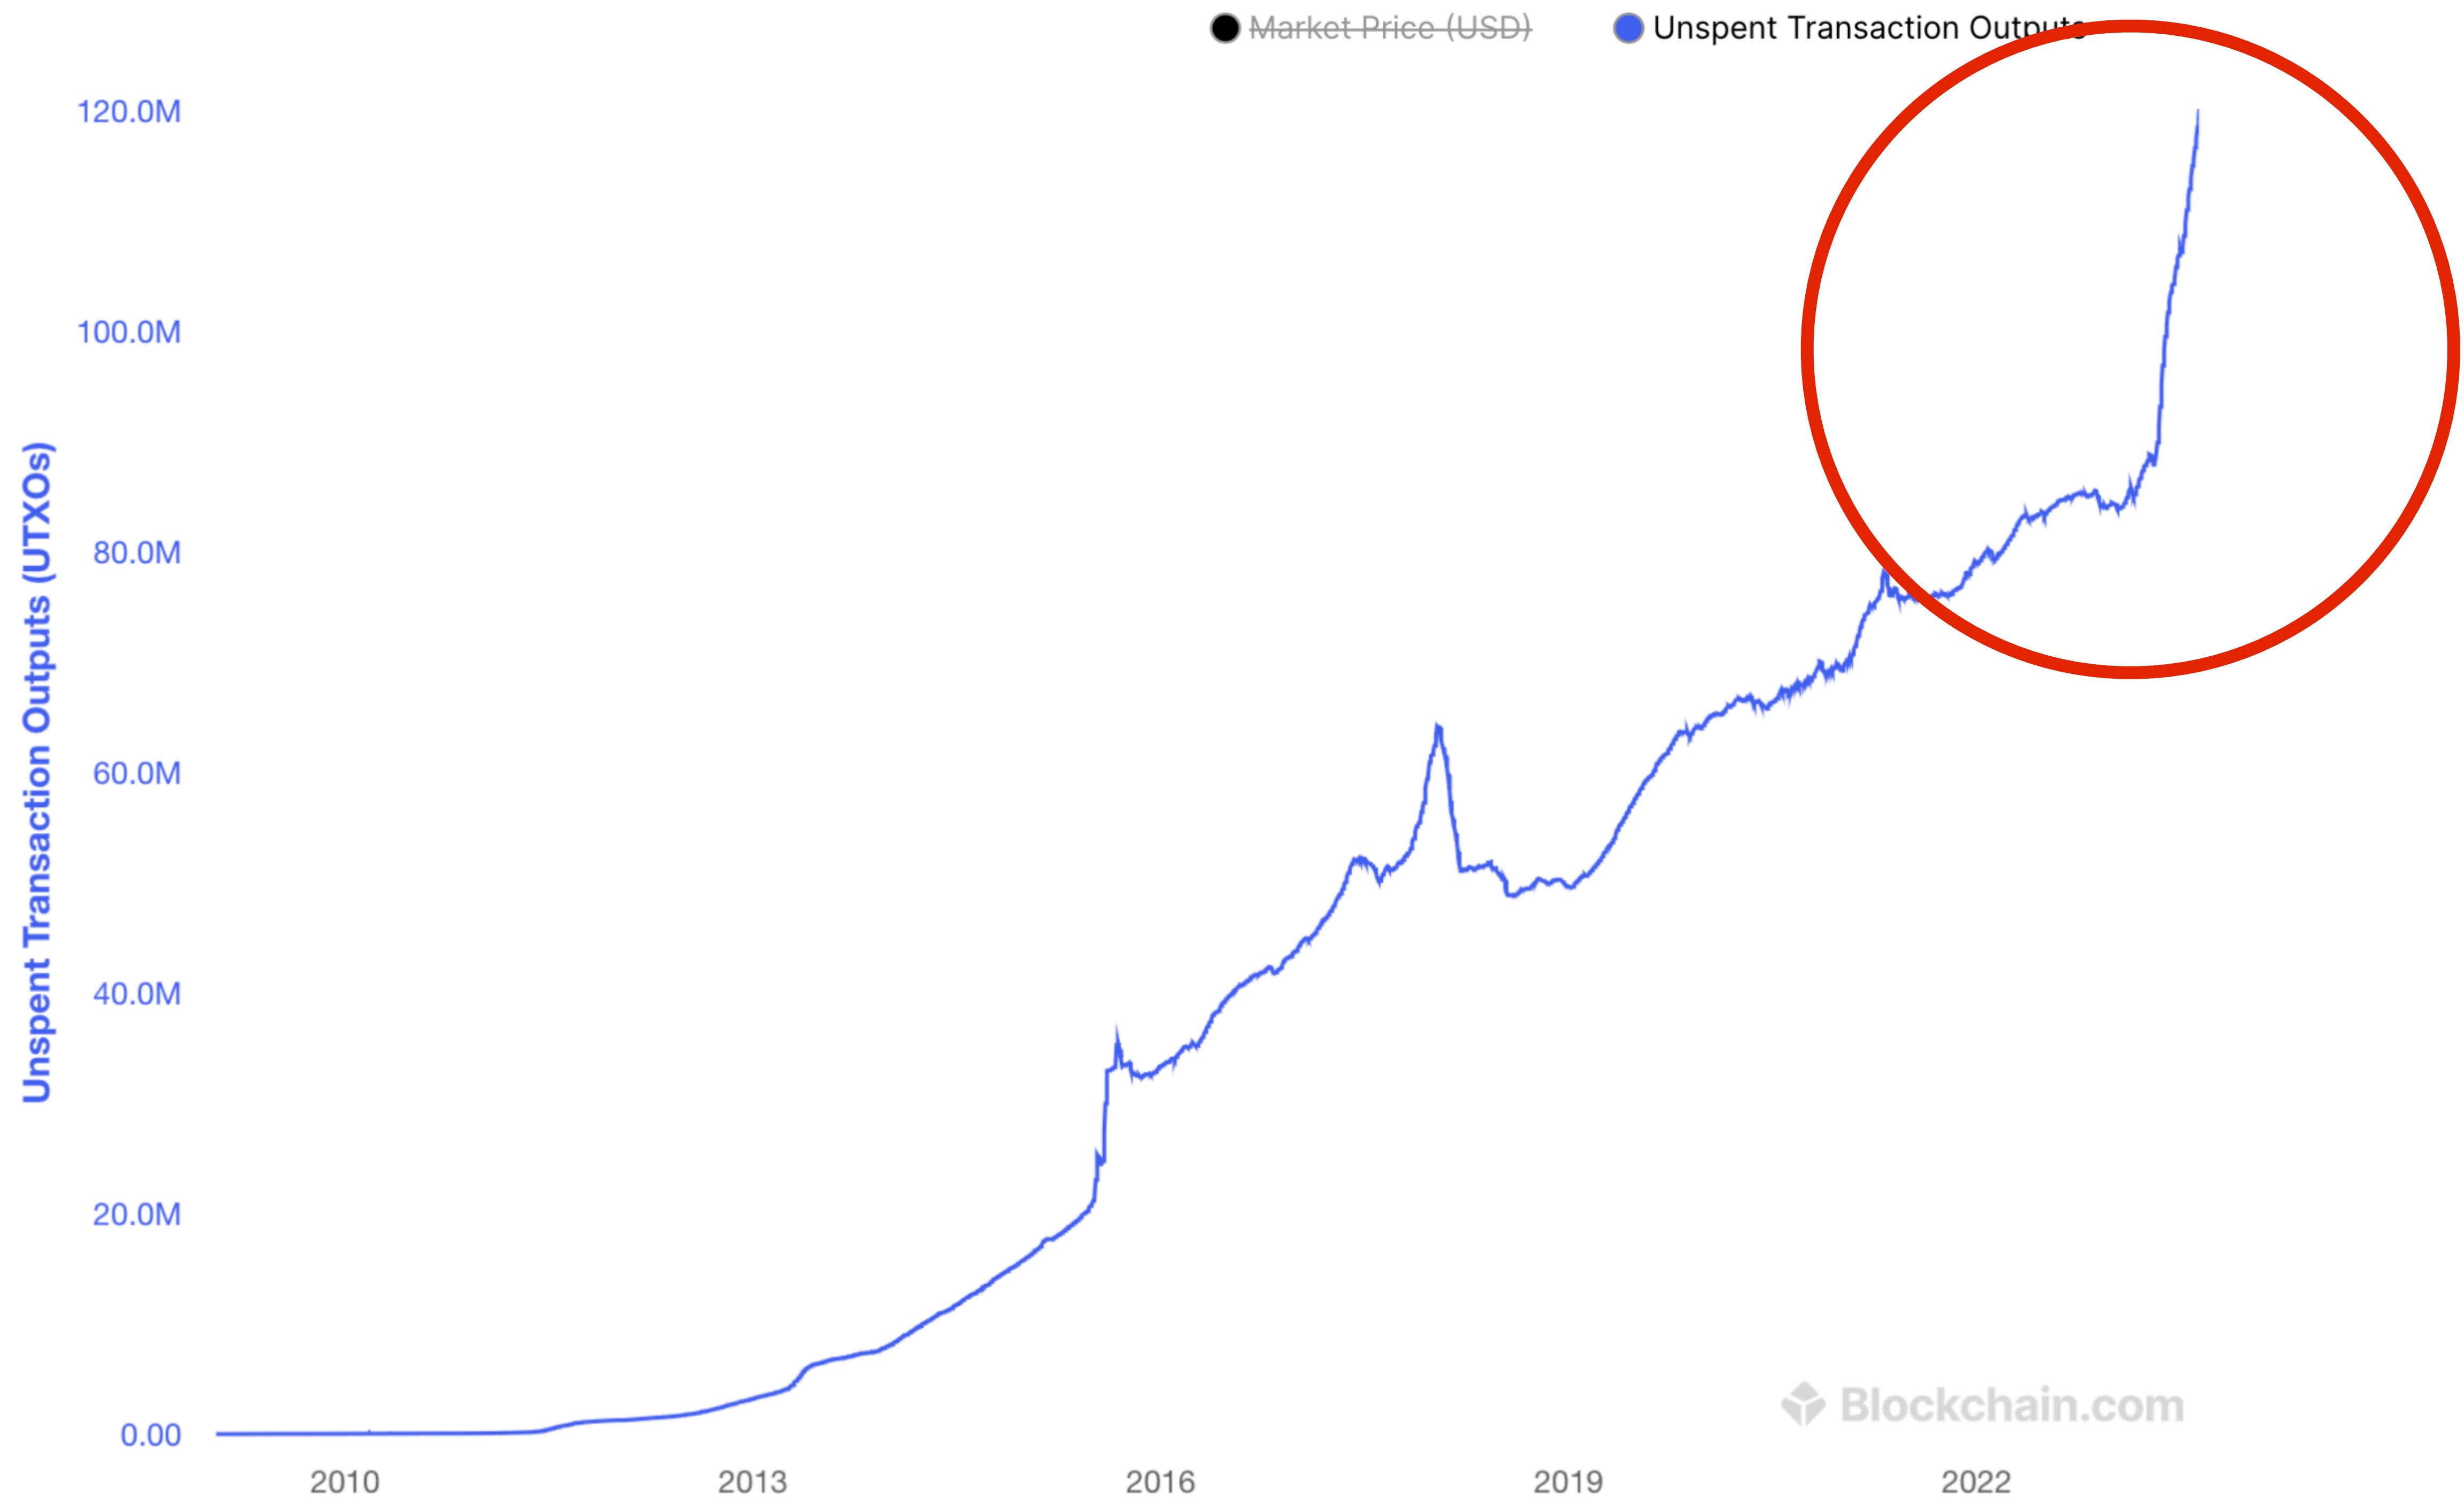
# Caching

Current way of reducing the disk burden

- As the UTXO set grows, the RAM size limit will make optimization less effective

● Market Price (USD) ● Unspent Transaction Outputs





**As the UTXO set grows, it  
gets slower**

**UTXO set size has no  
bounds**

**8 billion UTXO, 59B each**

**472GB**

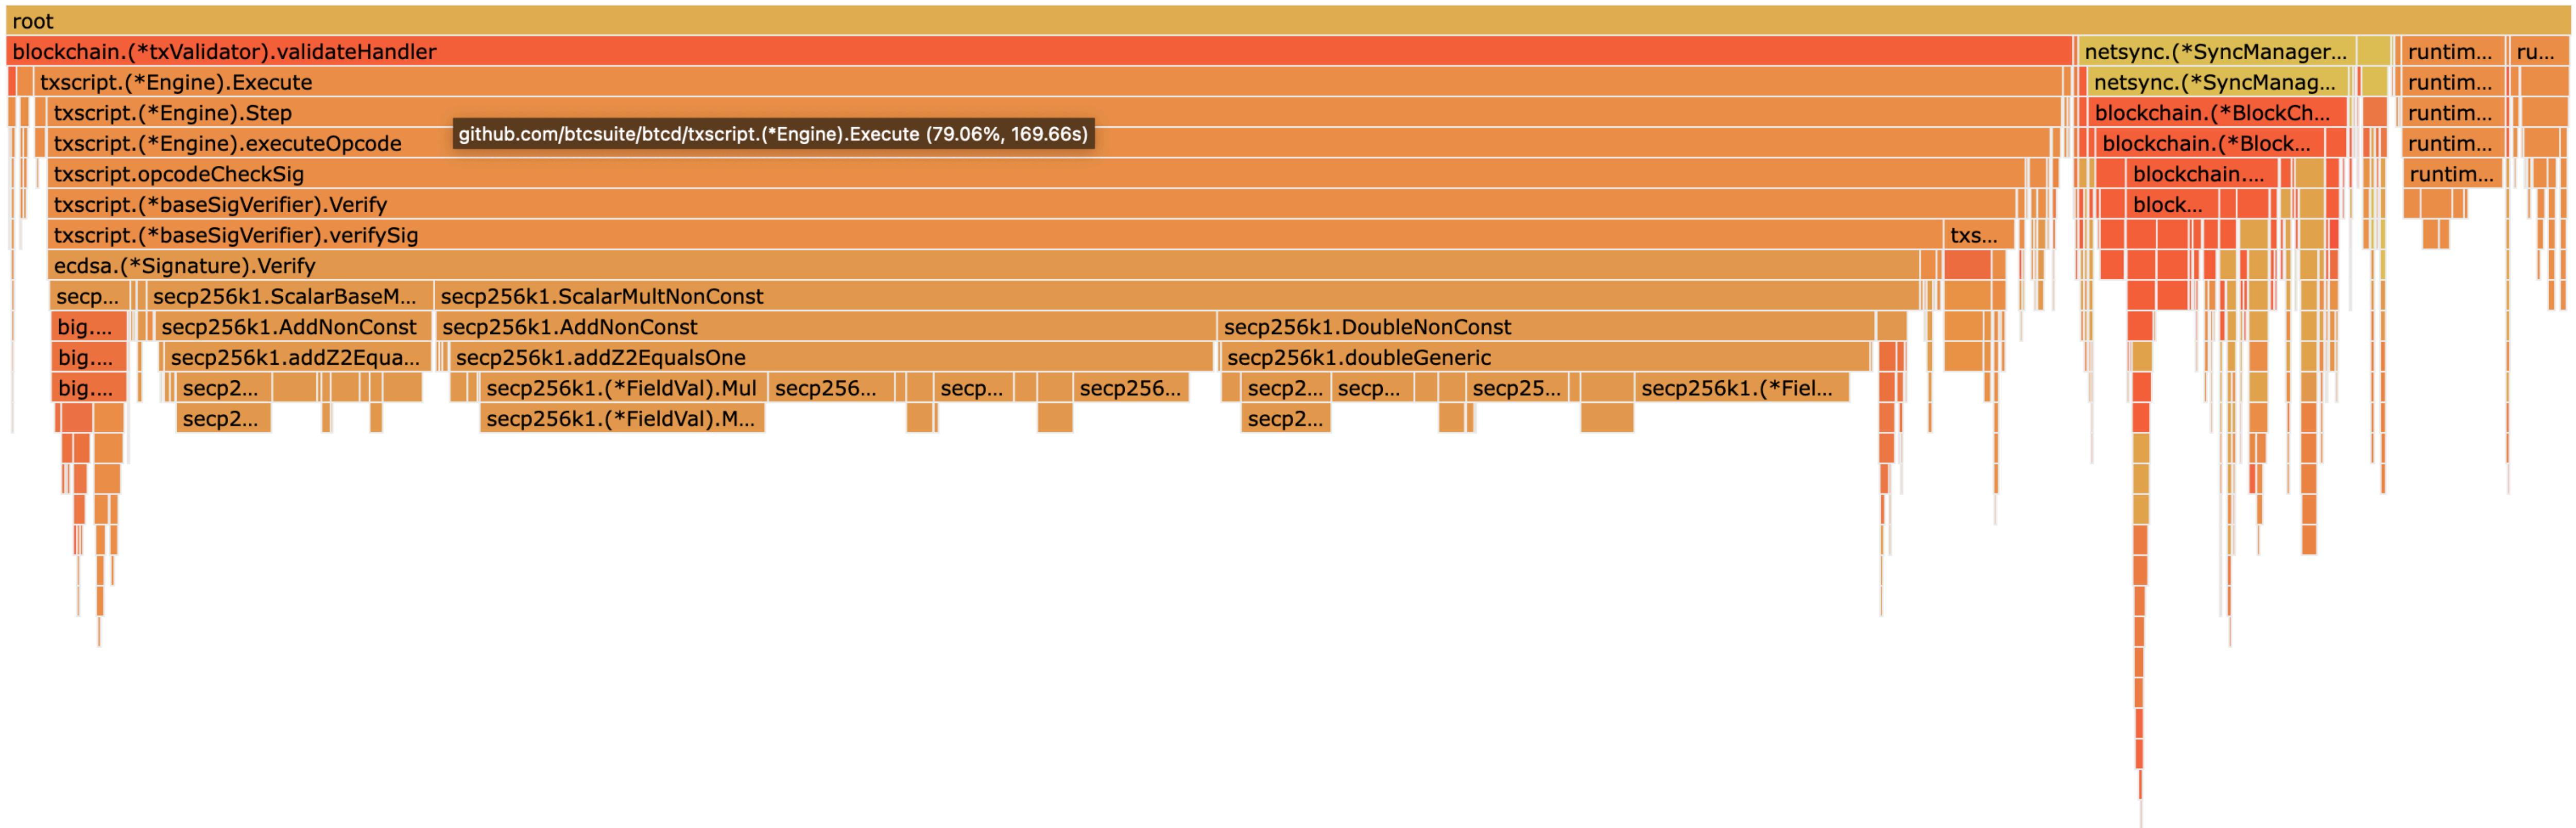
**Already bottleneck for HDD**

# CPU

## Transaction signature validation

- Signature operations take up a bulk of the block validation

github.com/btcsuite/btcd/txscript.(\*Engine).Execute (79.06%, 169.66s)



**80% time on sig  
validation**

**For btcd**

# CPU

## Transaction signature validation

- CPU cores increasing
- Taproot txs can be batch verified

# libsecp256k1

[build](#) [passing](#) [dependencies](#) [none](#) [irc.libera.chat](#) [#secp256k1](#)

Optimized C library for ECDSA signatures and secret/public key operations on curve secp256k1.

This library is intended to be the highest quality publicly available library for cryptography on the secp256k1 curve. However, the primary focus of its development has been for usage in the Bitcoin system and usage unlike Bitcoin's may be less well tested, verified, or suffer from a less well thought out interface. Correct usage requires some care and consideration that the library is fit for your application's purpose.

## Features:

- secp256k1 ECDSA signing/verification and key generation.
- Additive and multiplicative tweaking of secret/public keys.
- Serialization/parsing of secret keys, public keys, signatures.
- Constant time, constant memory access signing and public key generation.
- Derandomized ECDSA (via RFC6979 or with a caller provided function.)
- Very efficient implementation.
- Suitable for embedded systems.
- No runtime dependencies.
- Optional module for public key recovery.
- Optional module for ECDH key exchange.
- Optional module for Schnorr signatures according to [BIP-340](#).

**Already CPU bottlenecked  
but rooms for improvement**

**What do we need to increase the  
block size?**

**Signature validation ↑**  
**UTXO read/write ↑**

# **How to improve read/write?**

# **Utreexo**

**The layer 1 scaling solution!**

# **Merkelize the UTXO set**

**472GB** →



→ **1 KB**

**~Same performance on  
HDD and SSD**

# **disk read/write X**

**No more random reads**

# Testing the performance

## Test methodology

- Use an option that does a lot of read/writes and compare

**36% slower**

**On current nodes**

**3% slower**

**On Utreexo nodes**

# The writes that happen for Utreexo

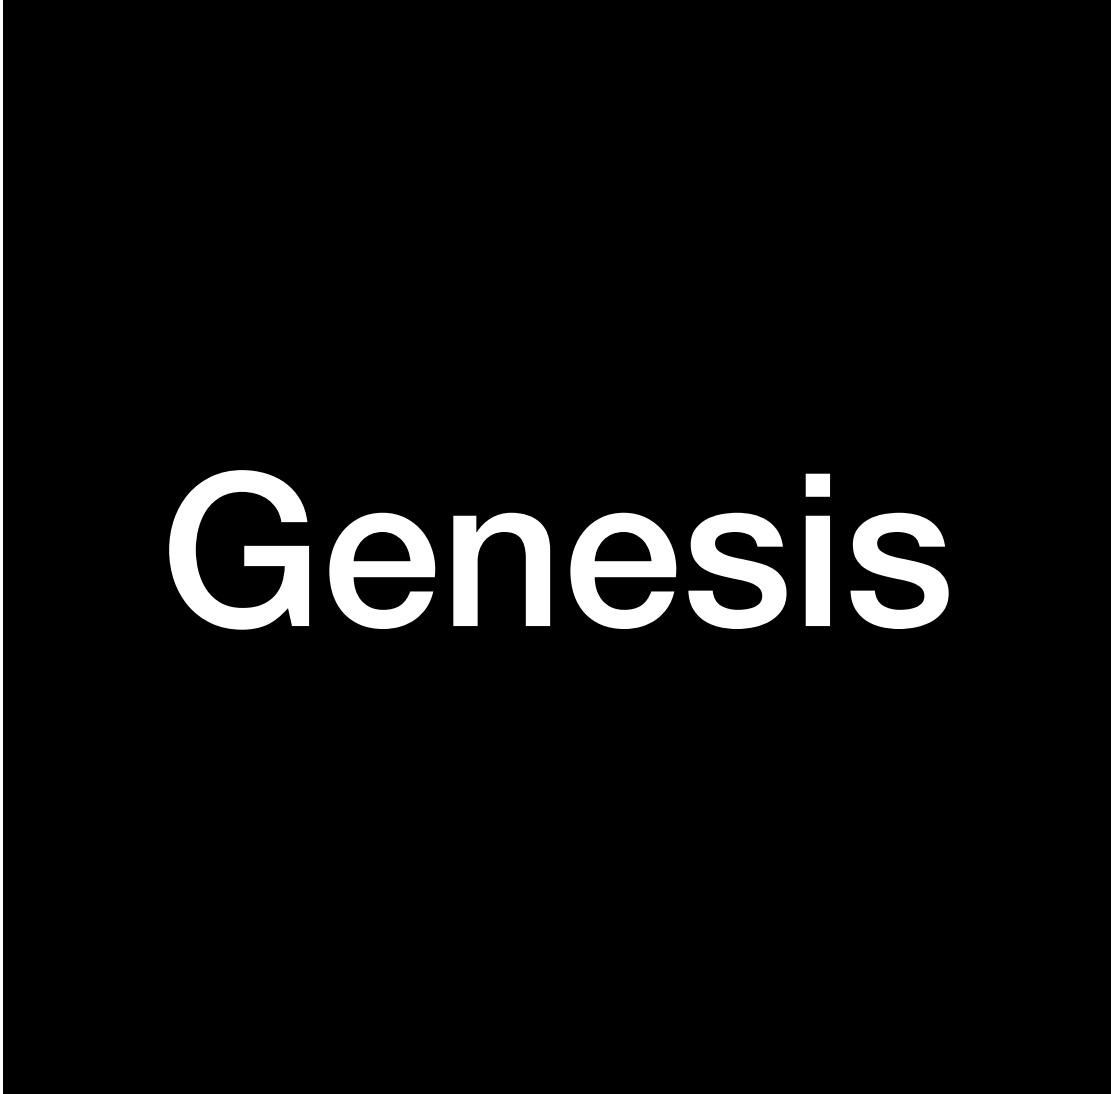
## No reads

```
{443000, newHashFromStr("000000000000000047efc3c126c3398cfdef1609682d1492cd2909e3c0a17f"),  
    43243051, []*chainhash.Hash{  
        newLeafHashFromStr("df31c496a54a8021c38c77809da15a8bd3968970315f5a56c07f0fd9763262ff"),  
        newLeafHashFromStr("755f4cfbf5727c3e14aa6dcc6e5ab150c61101cead06fade5a886ee8d1dd377"),  
        newLeafHashFromStr("c572a92c7cd9e14ce5d7c94473ef90bee21856654eed5535314b3c117006bc83"),  
        newLeafHashFromStr("9e85567347ef0ea38c1797d8ed8e399ed0d99a07025388f6e06791ff6b88643c"),  
        newLeafHashFromStr("ab4efc738bd7b8281689e899979d44ed0ca8fe1489efa0f5a09f5b1c276fe792"),  
        newLeafHashFromStr("1c763d8f490baef0b756597dfc2da5a31d0aee324da6314b52f21707351cf51a"),  
        newLeafHashFromStr("29fc9989939a17e5d113fb1f247f3bd05606488bcb48c9cfad55a57006f2248"),  
        newLeafHashFromStr("e79f840fe1bcd6b638a53b637885dda7ea3da2b2b5f1f5447ebf4d54fc0a2e"),  
        newLeafHashFromStr("02f33c3ce0d684e644410fa35fee4c7572454facc710389ea0f69a00e27c1780"),  
        newLeafHashFromStr("4e29ad940088ab42991bd99d7943a68ea38d498f8435c4bea21b8252788b5a7d"),  
        newLeafHashFromStr("feabfc6565e918827df8ade12d0885e7da28bb81410c71a0fc92da28e5538e"),  
        newLeafHashFromStr("42275481edcd150c0891f8f7486941e7ecc9dbec0d1554d1ae3027d8523571b4"),  
        newLeafHashFromStr("84dd84bbb564ab9e878ea774d66a2cb818c2b74680974e01a07dde68bd064bfd"),  
        newLeafHashFromStr("19a64f2f77403d4a1e3482b3a1801aa360120f6429f5ed47261fc8676920d9d"),  
    },  
},  
},
```

# **Out of Order Block Validation**

# **Current Block Validation**

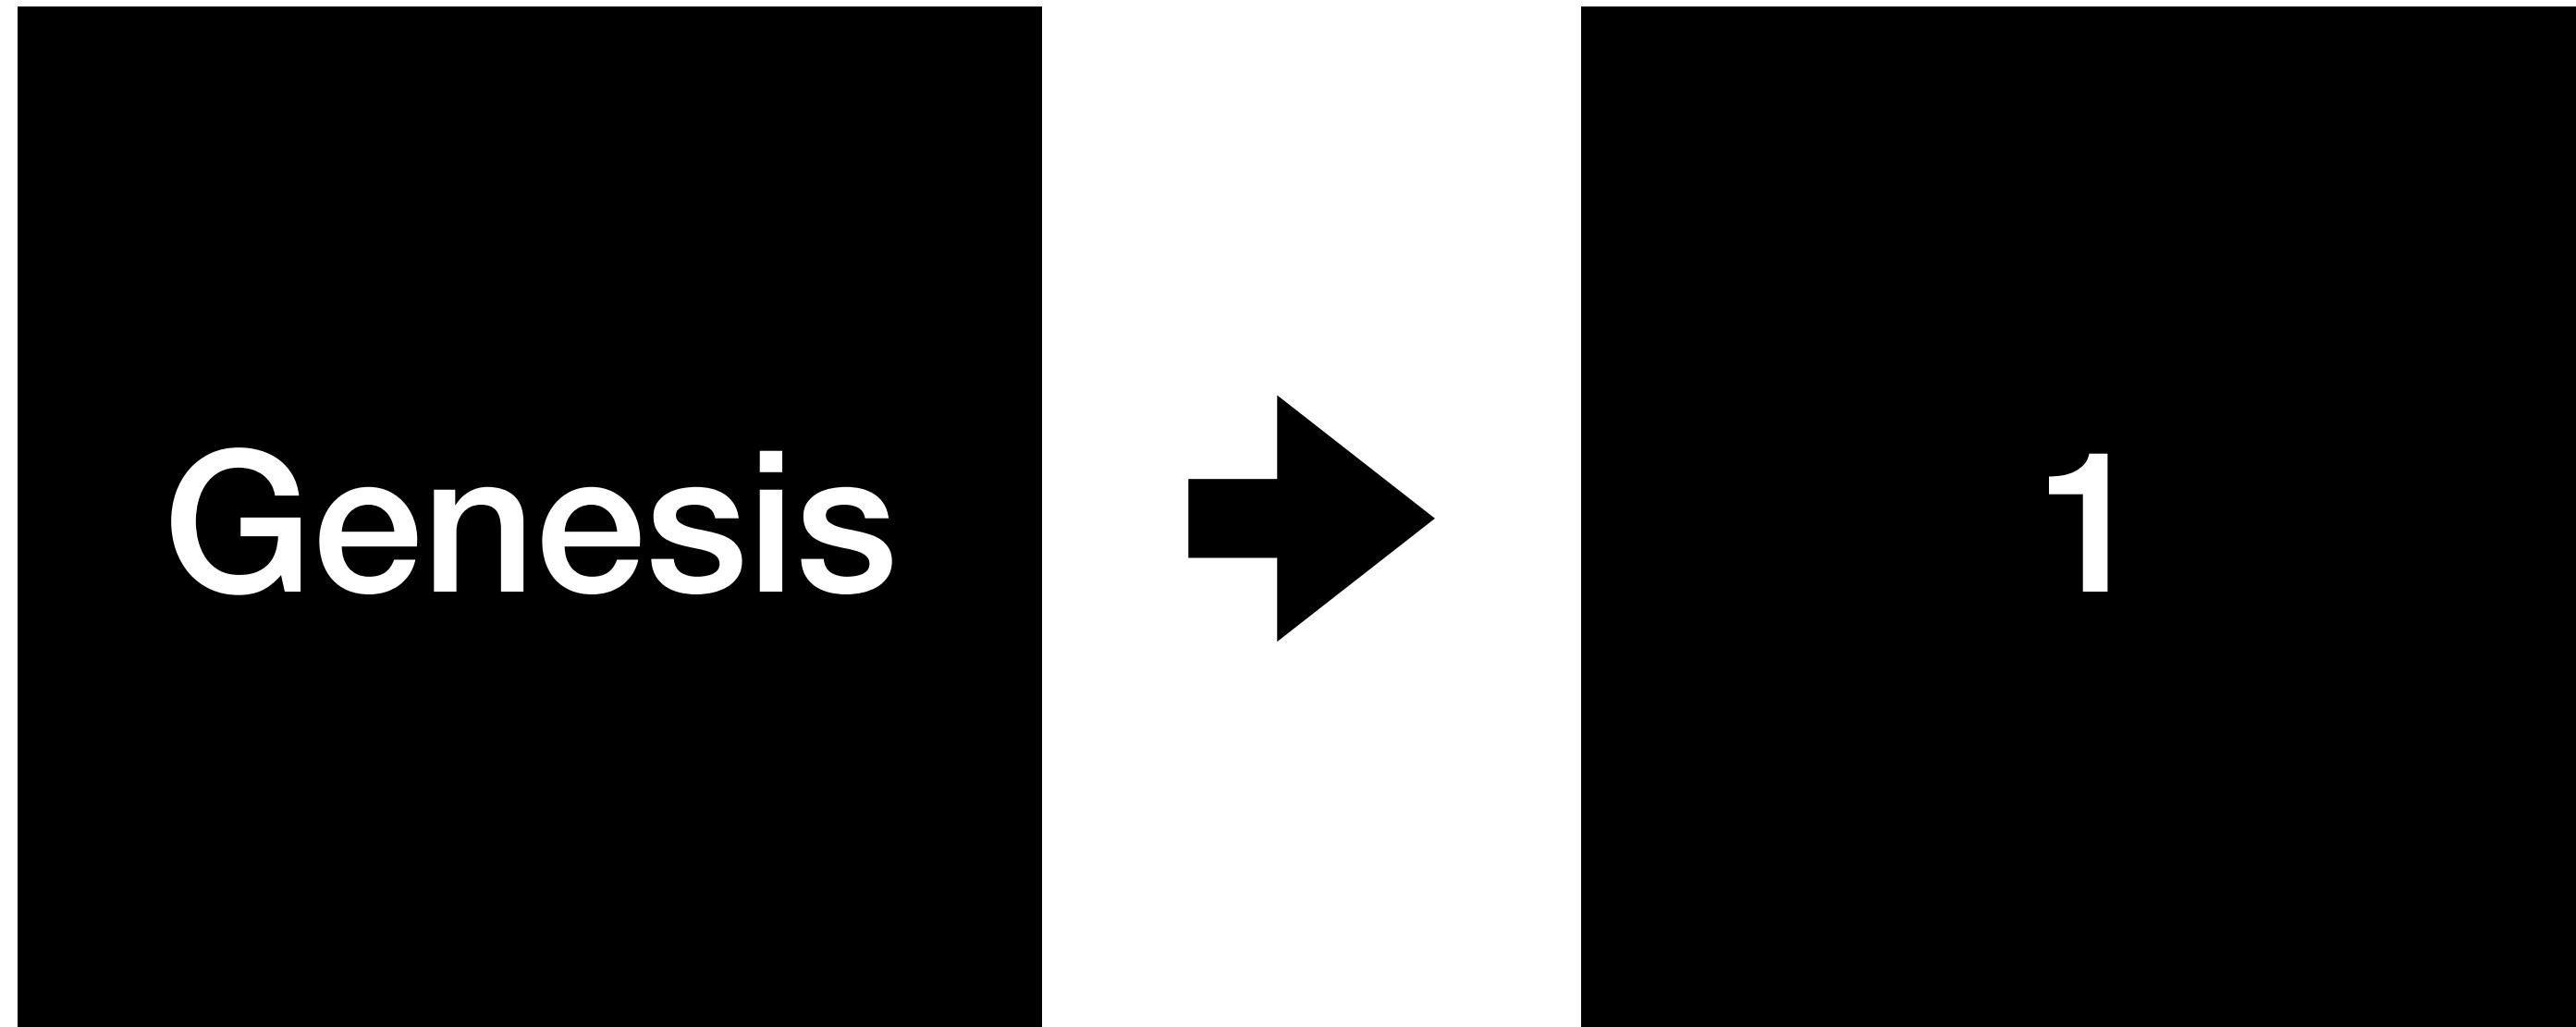
## **Sequential**



**Genesis**

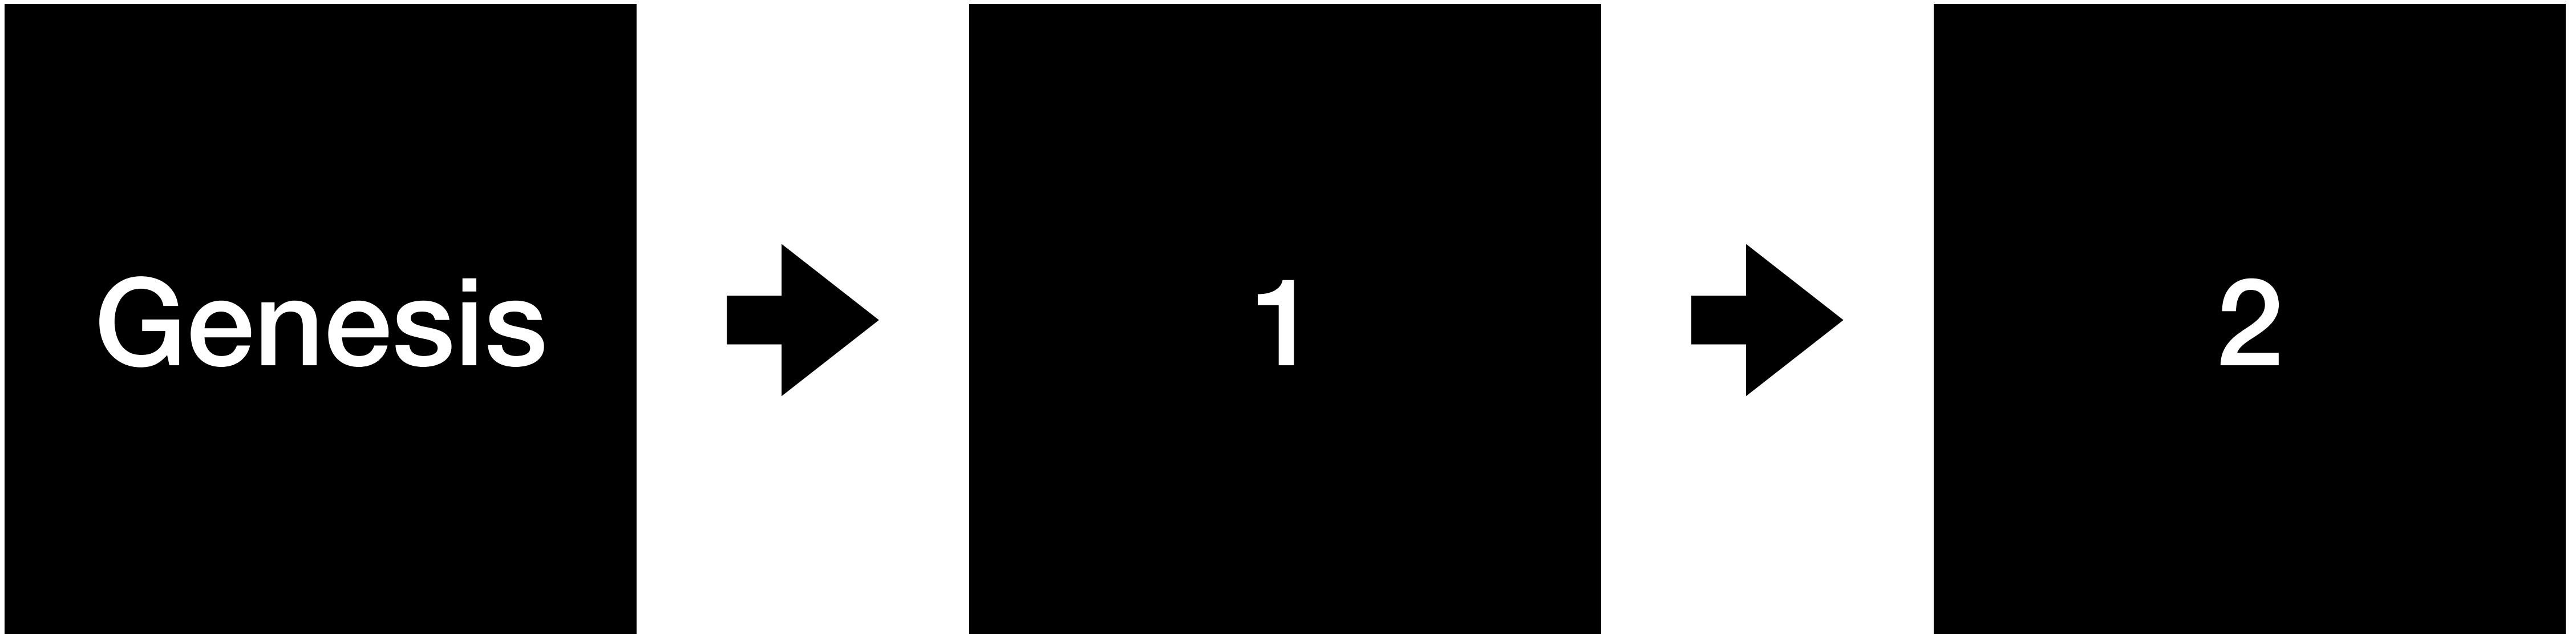
# Current Block Validation

## Sequential



# Current Block Validation

Sequential



# Out of Order Block Validation

Doesn't have to be sequential

- Provide the merkle roots of the Utreexo at each block

# **Out of Order Block Validation**

**Doesn't have to be sequential**

- Provide the merkle roots of the Utreexo at each block
- Use them to validate the next block in line

# Out of Order Validation

Genesis

100,000

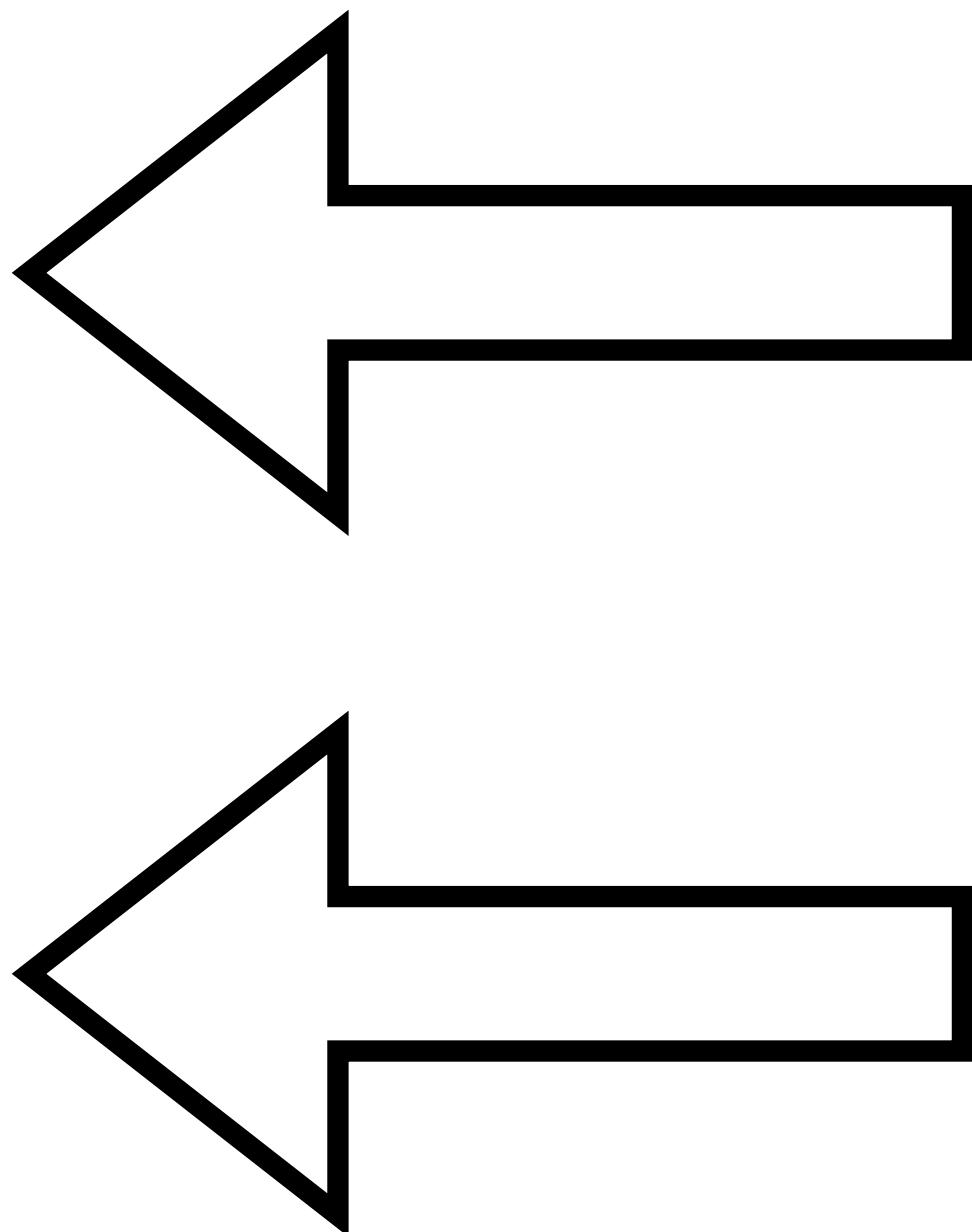
200,000

# Out of Order Validation

Genesis

100,000

200,000



Roots committed in  
binary

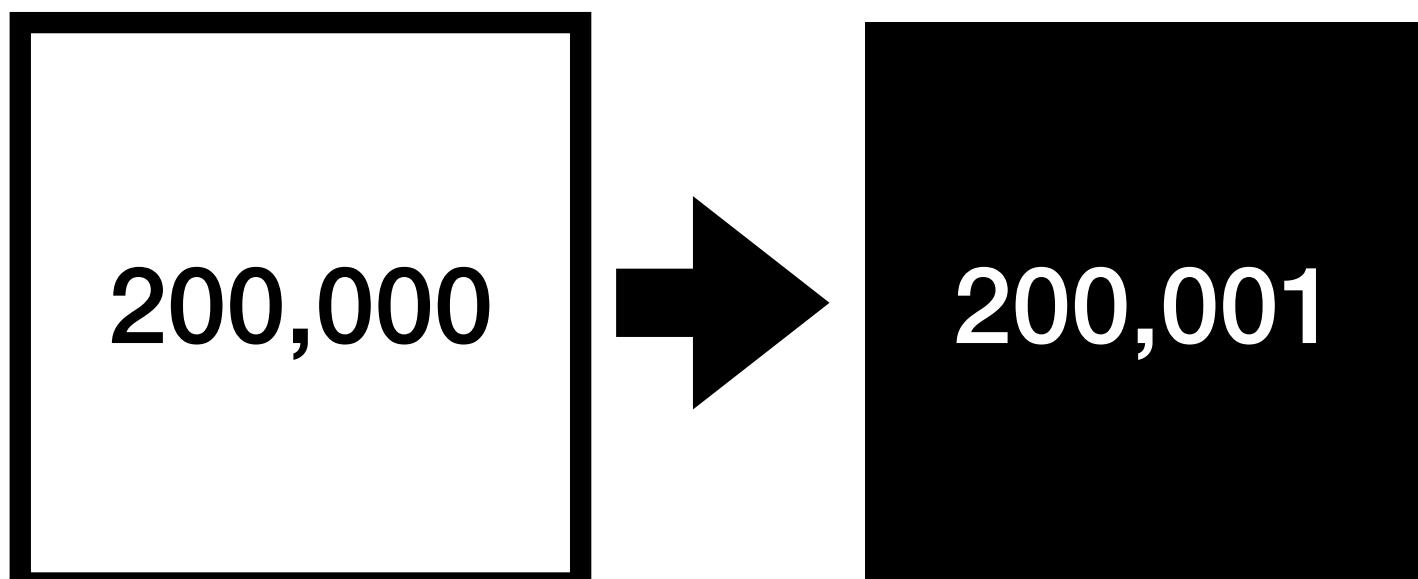
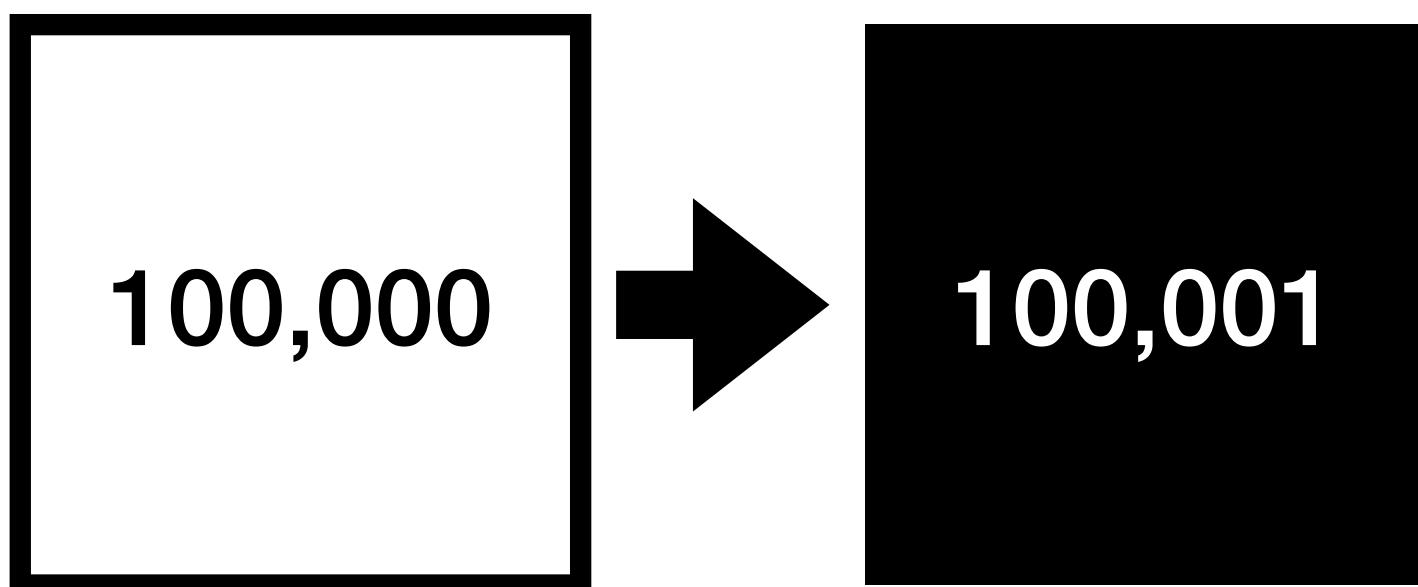
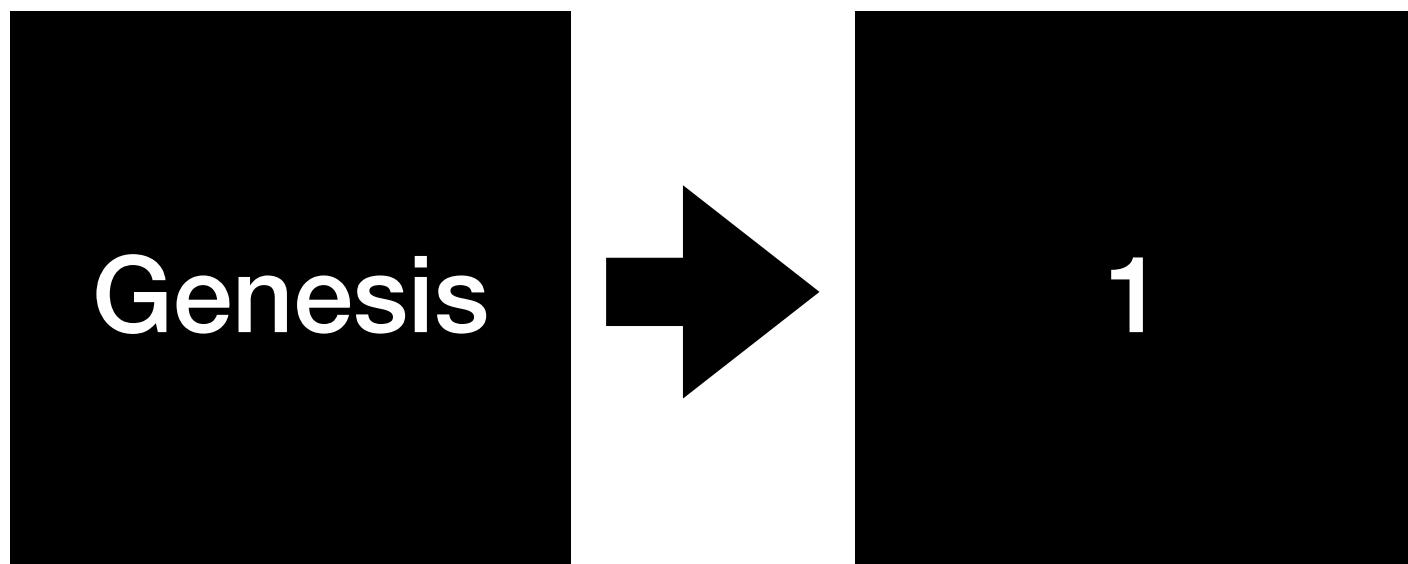
# Out of Order Validation

Genesis

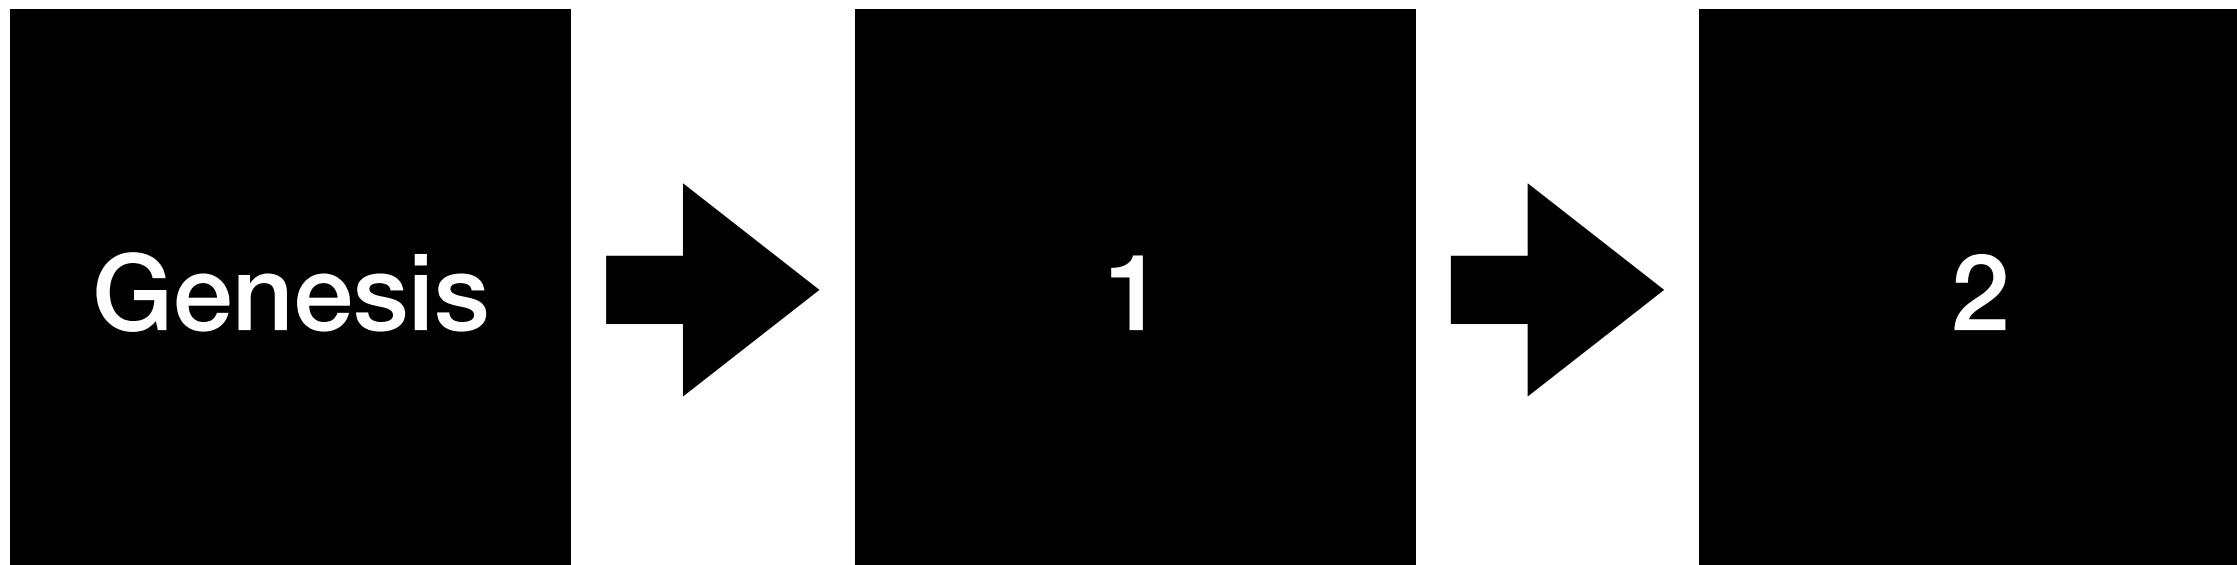
100,000

200,000

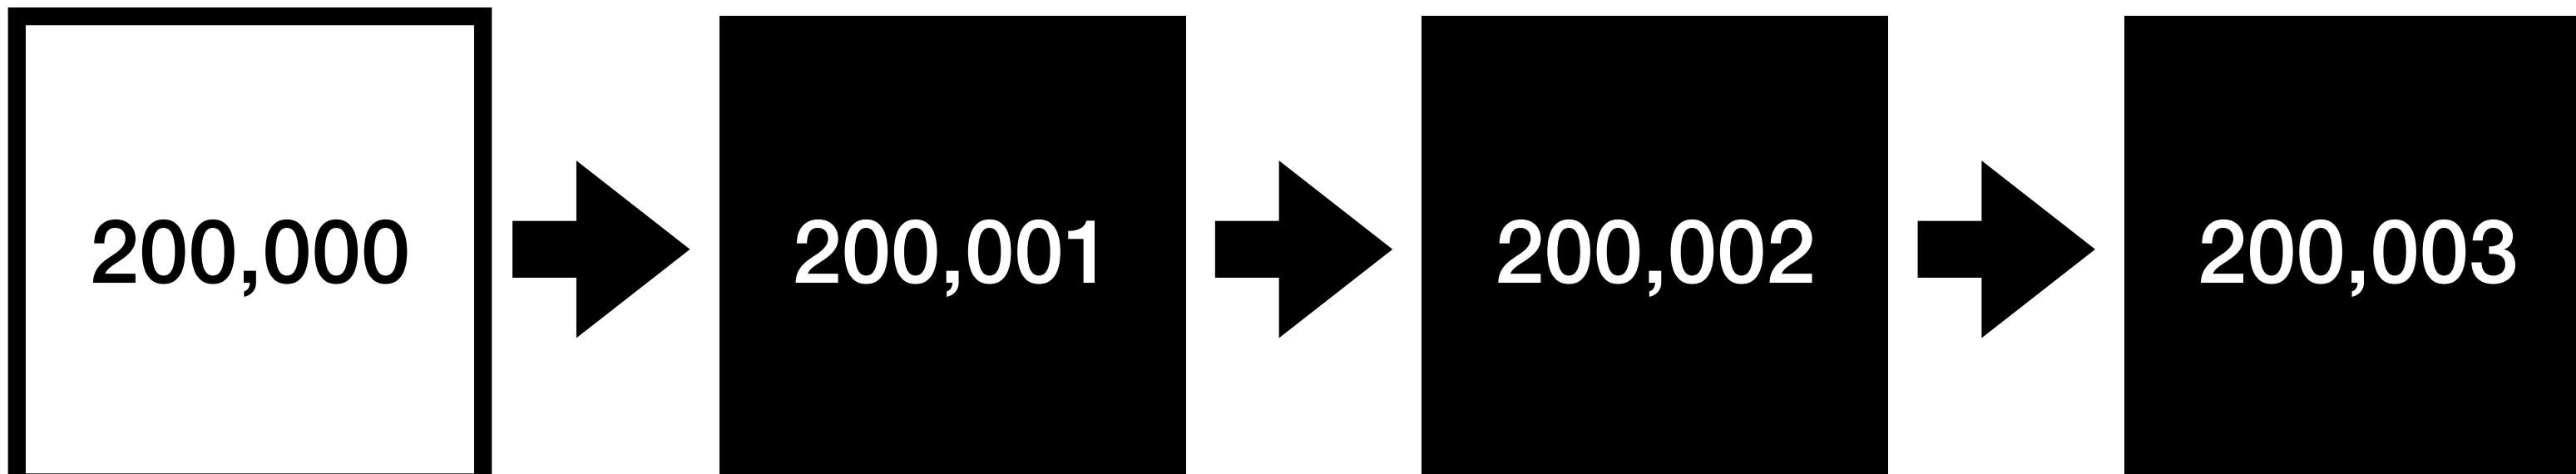
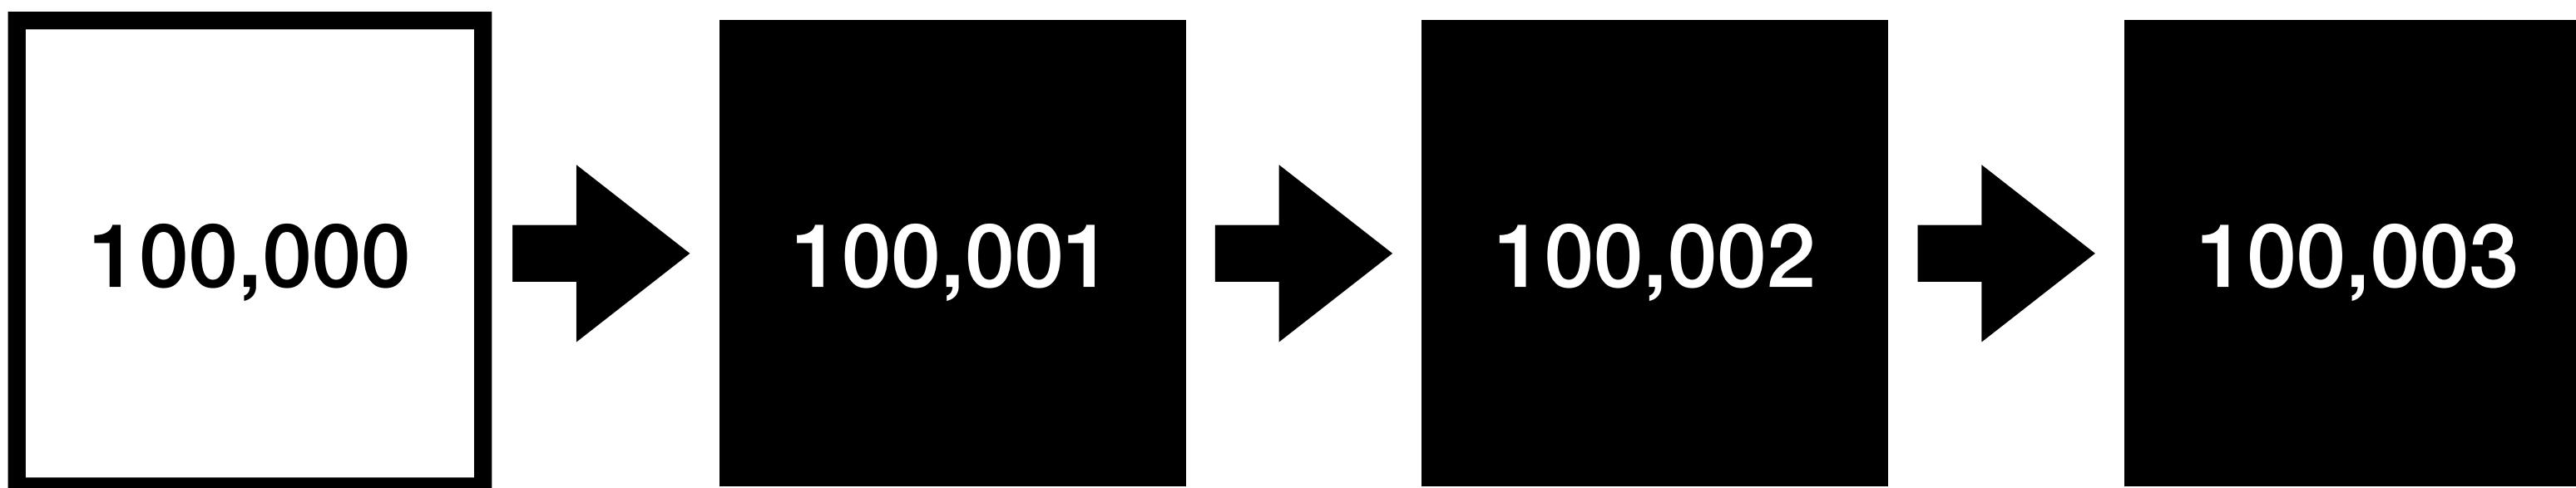
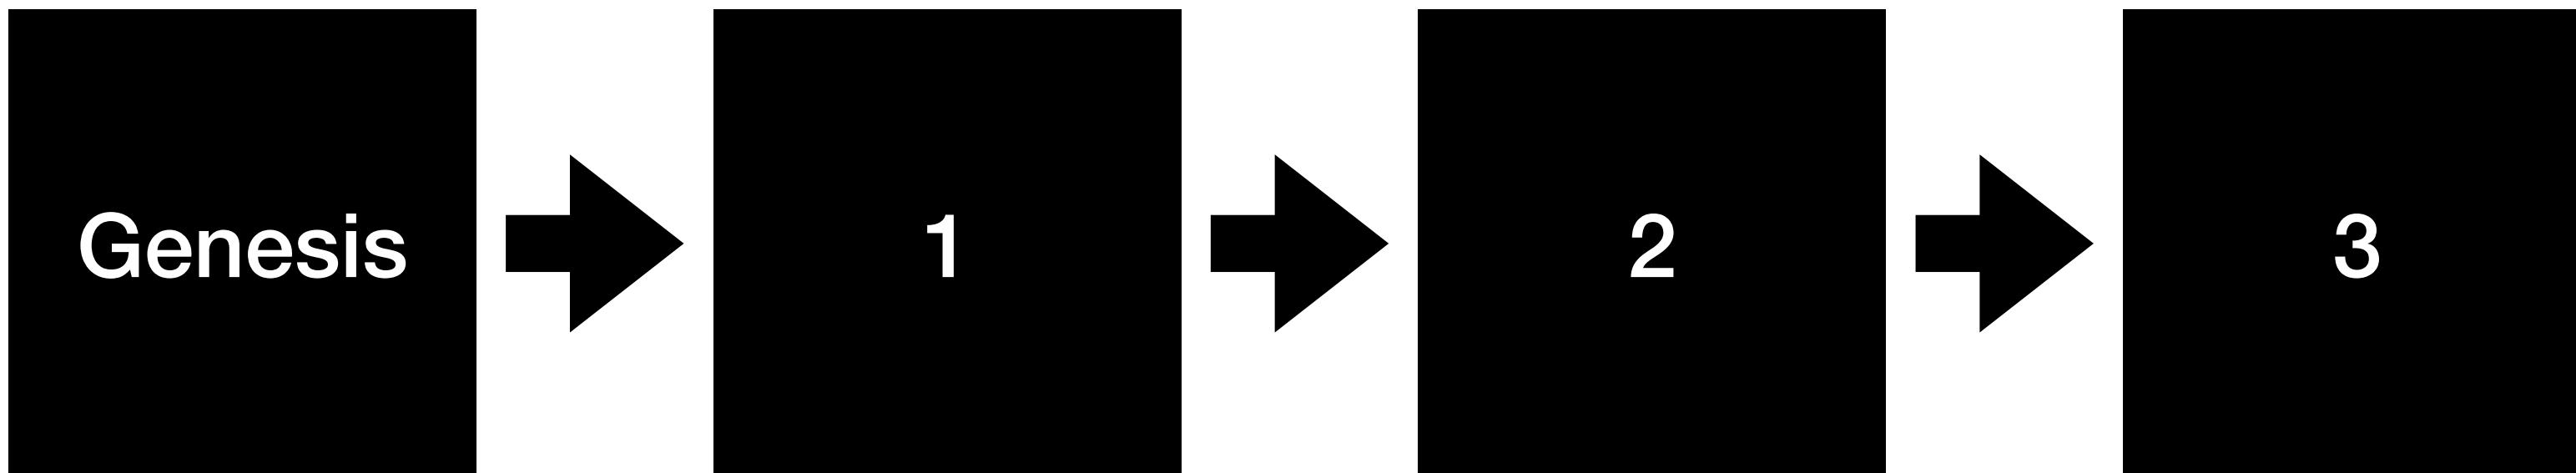
# Out of Order Validation



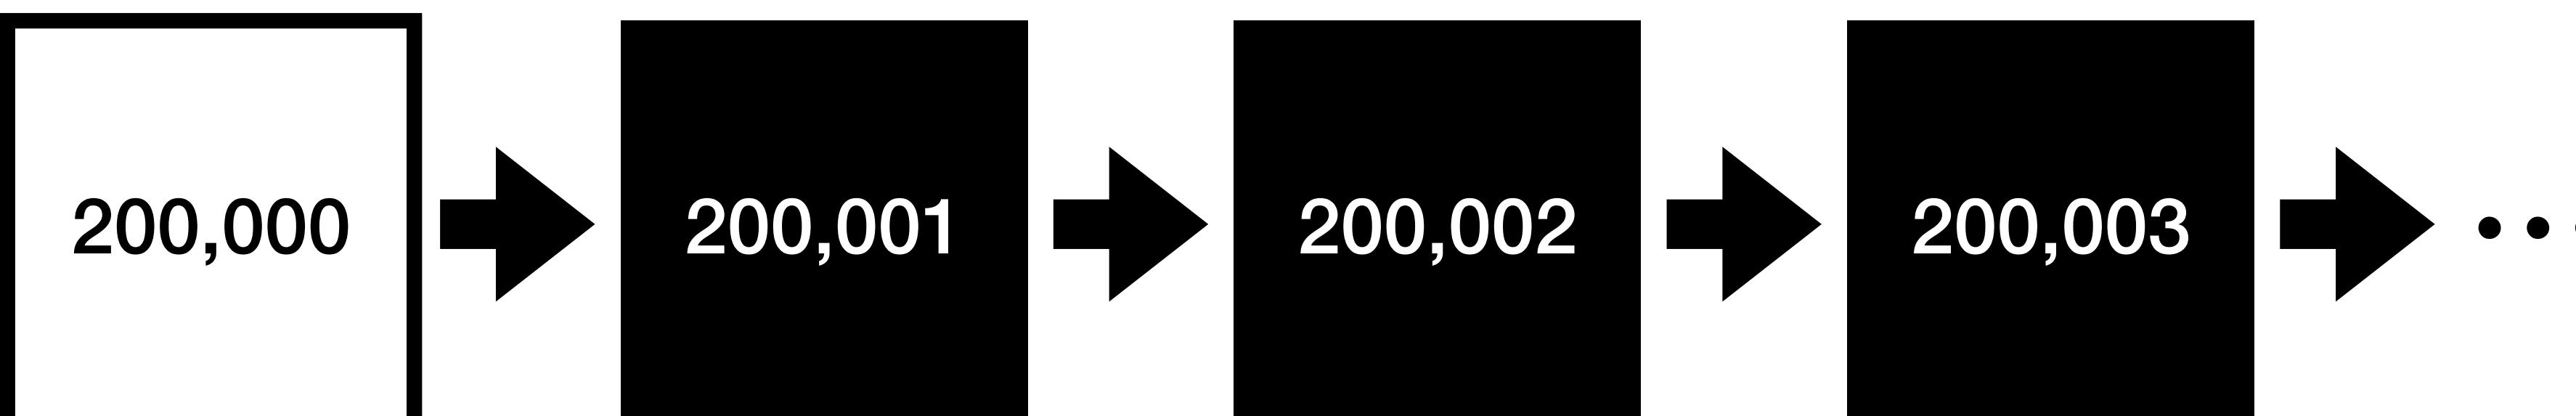
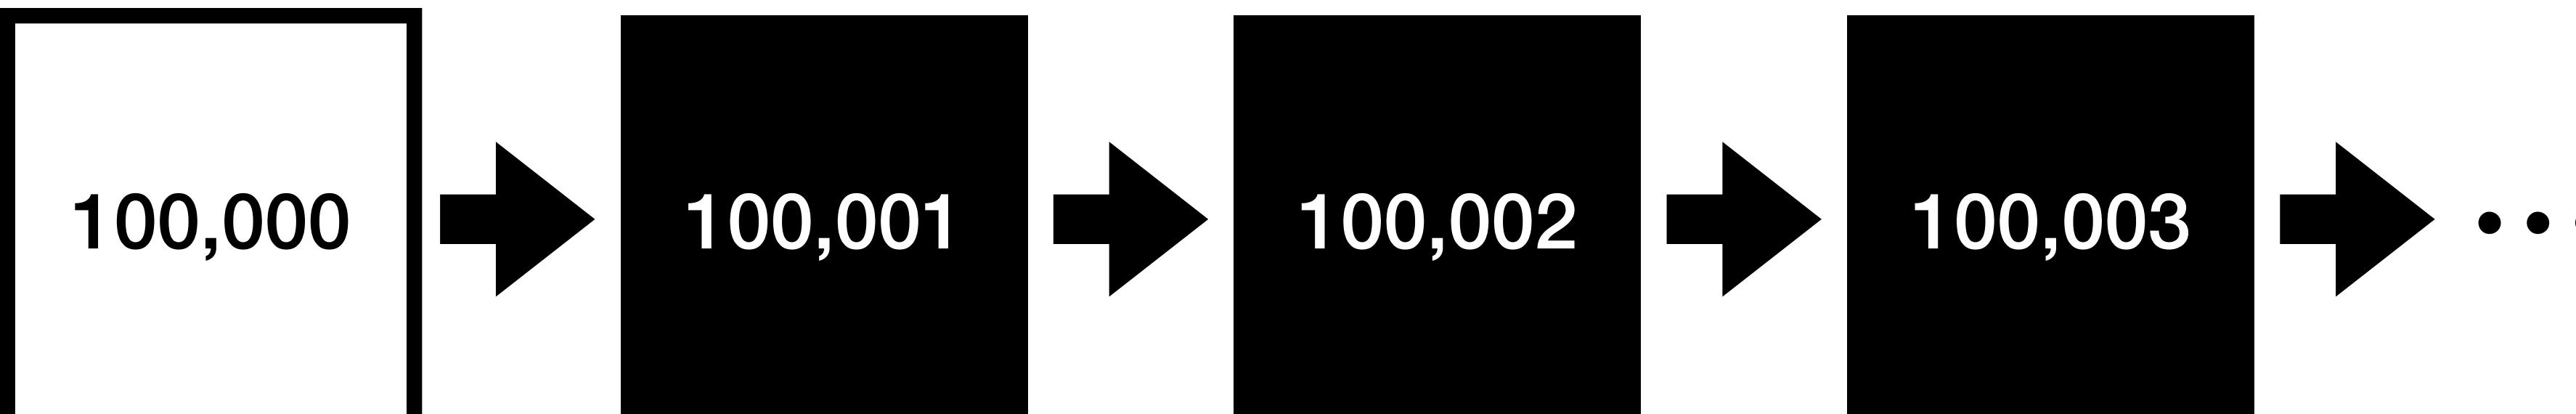
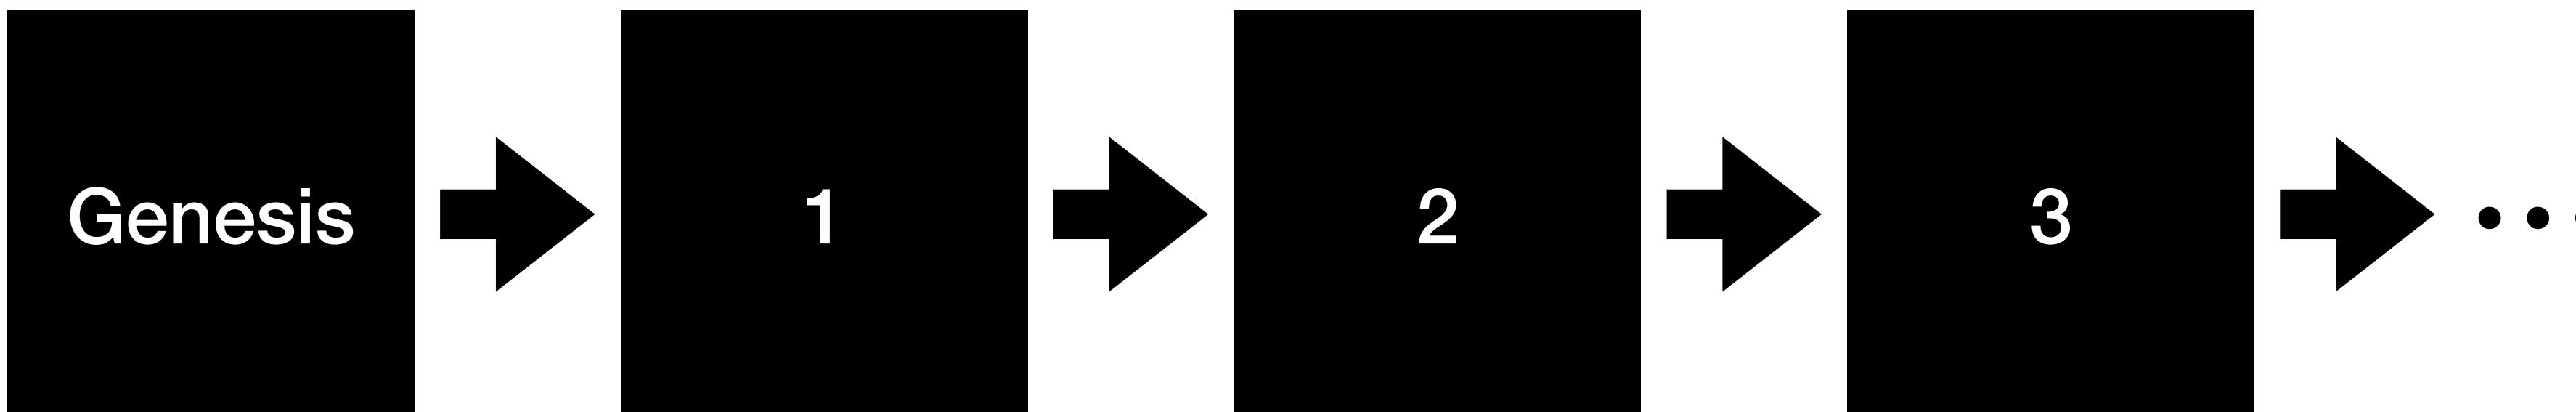
# Out of Order Validation



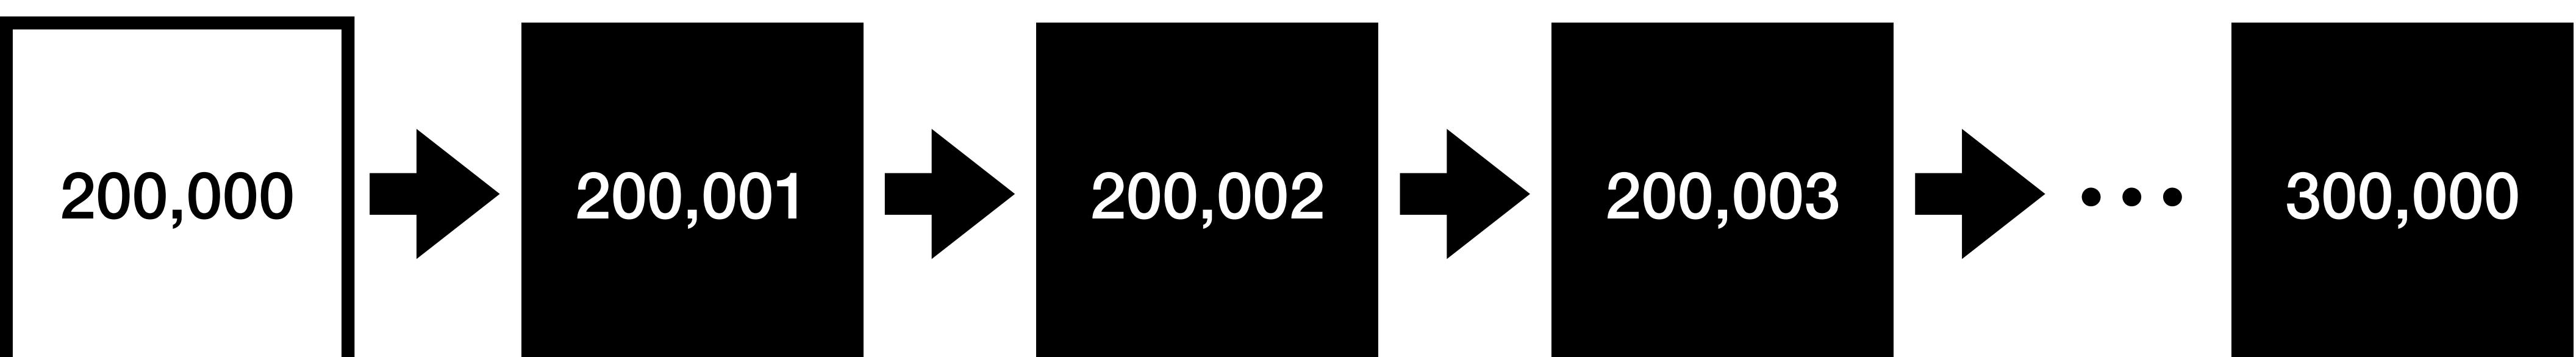
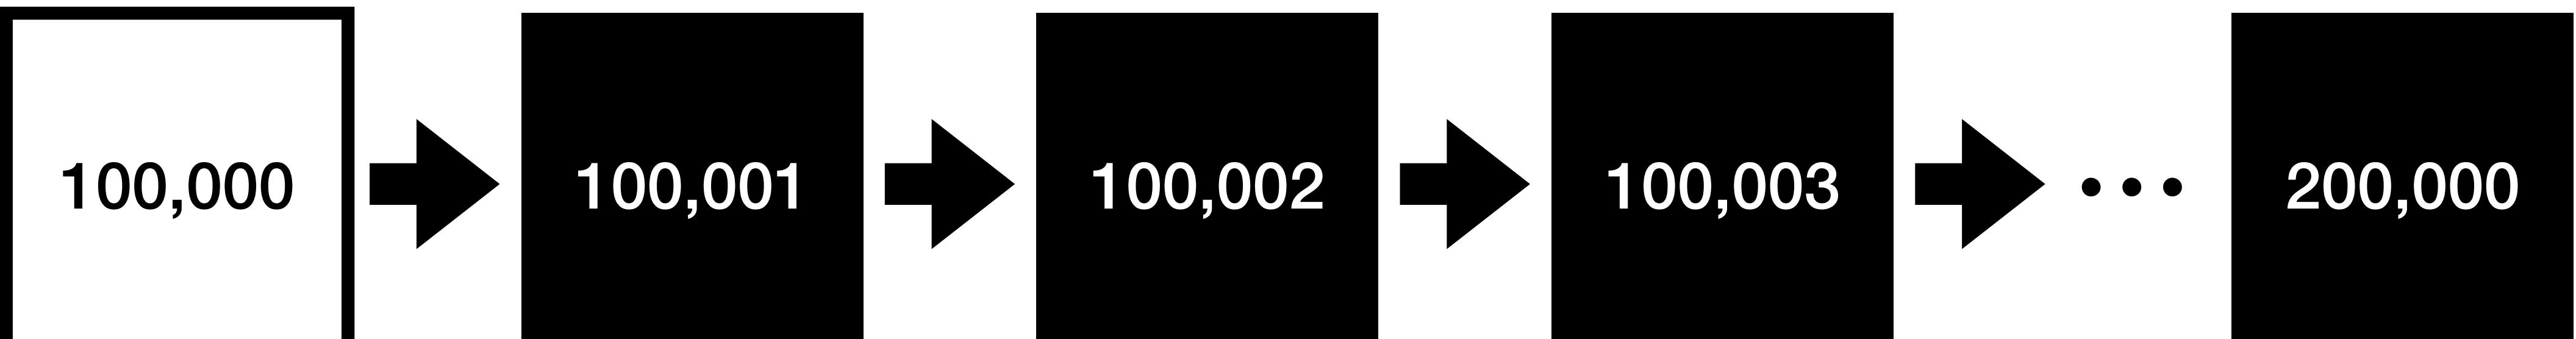
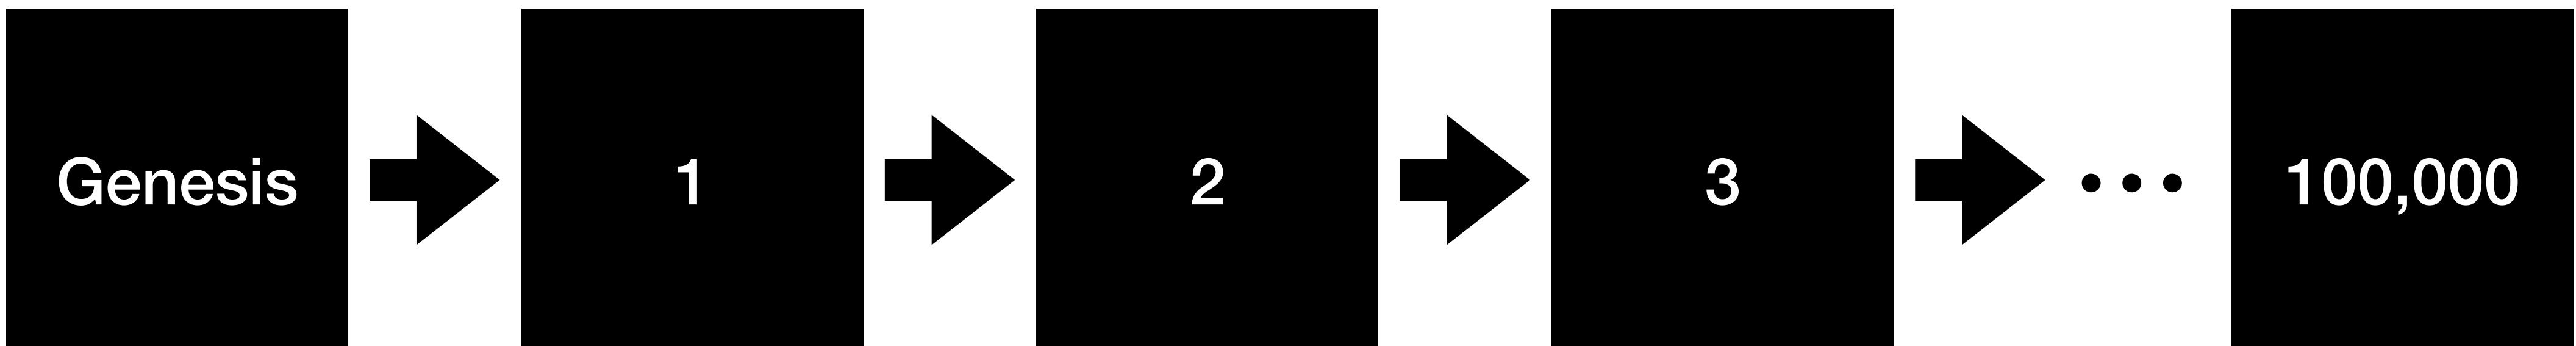
# Out of Order Validation



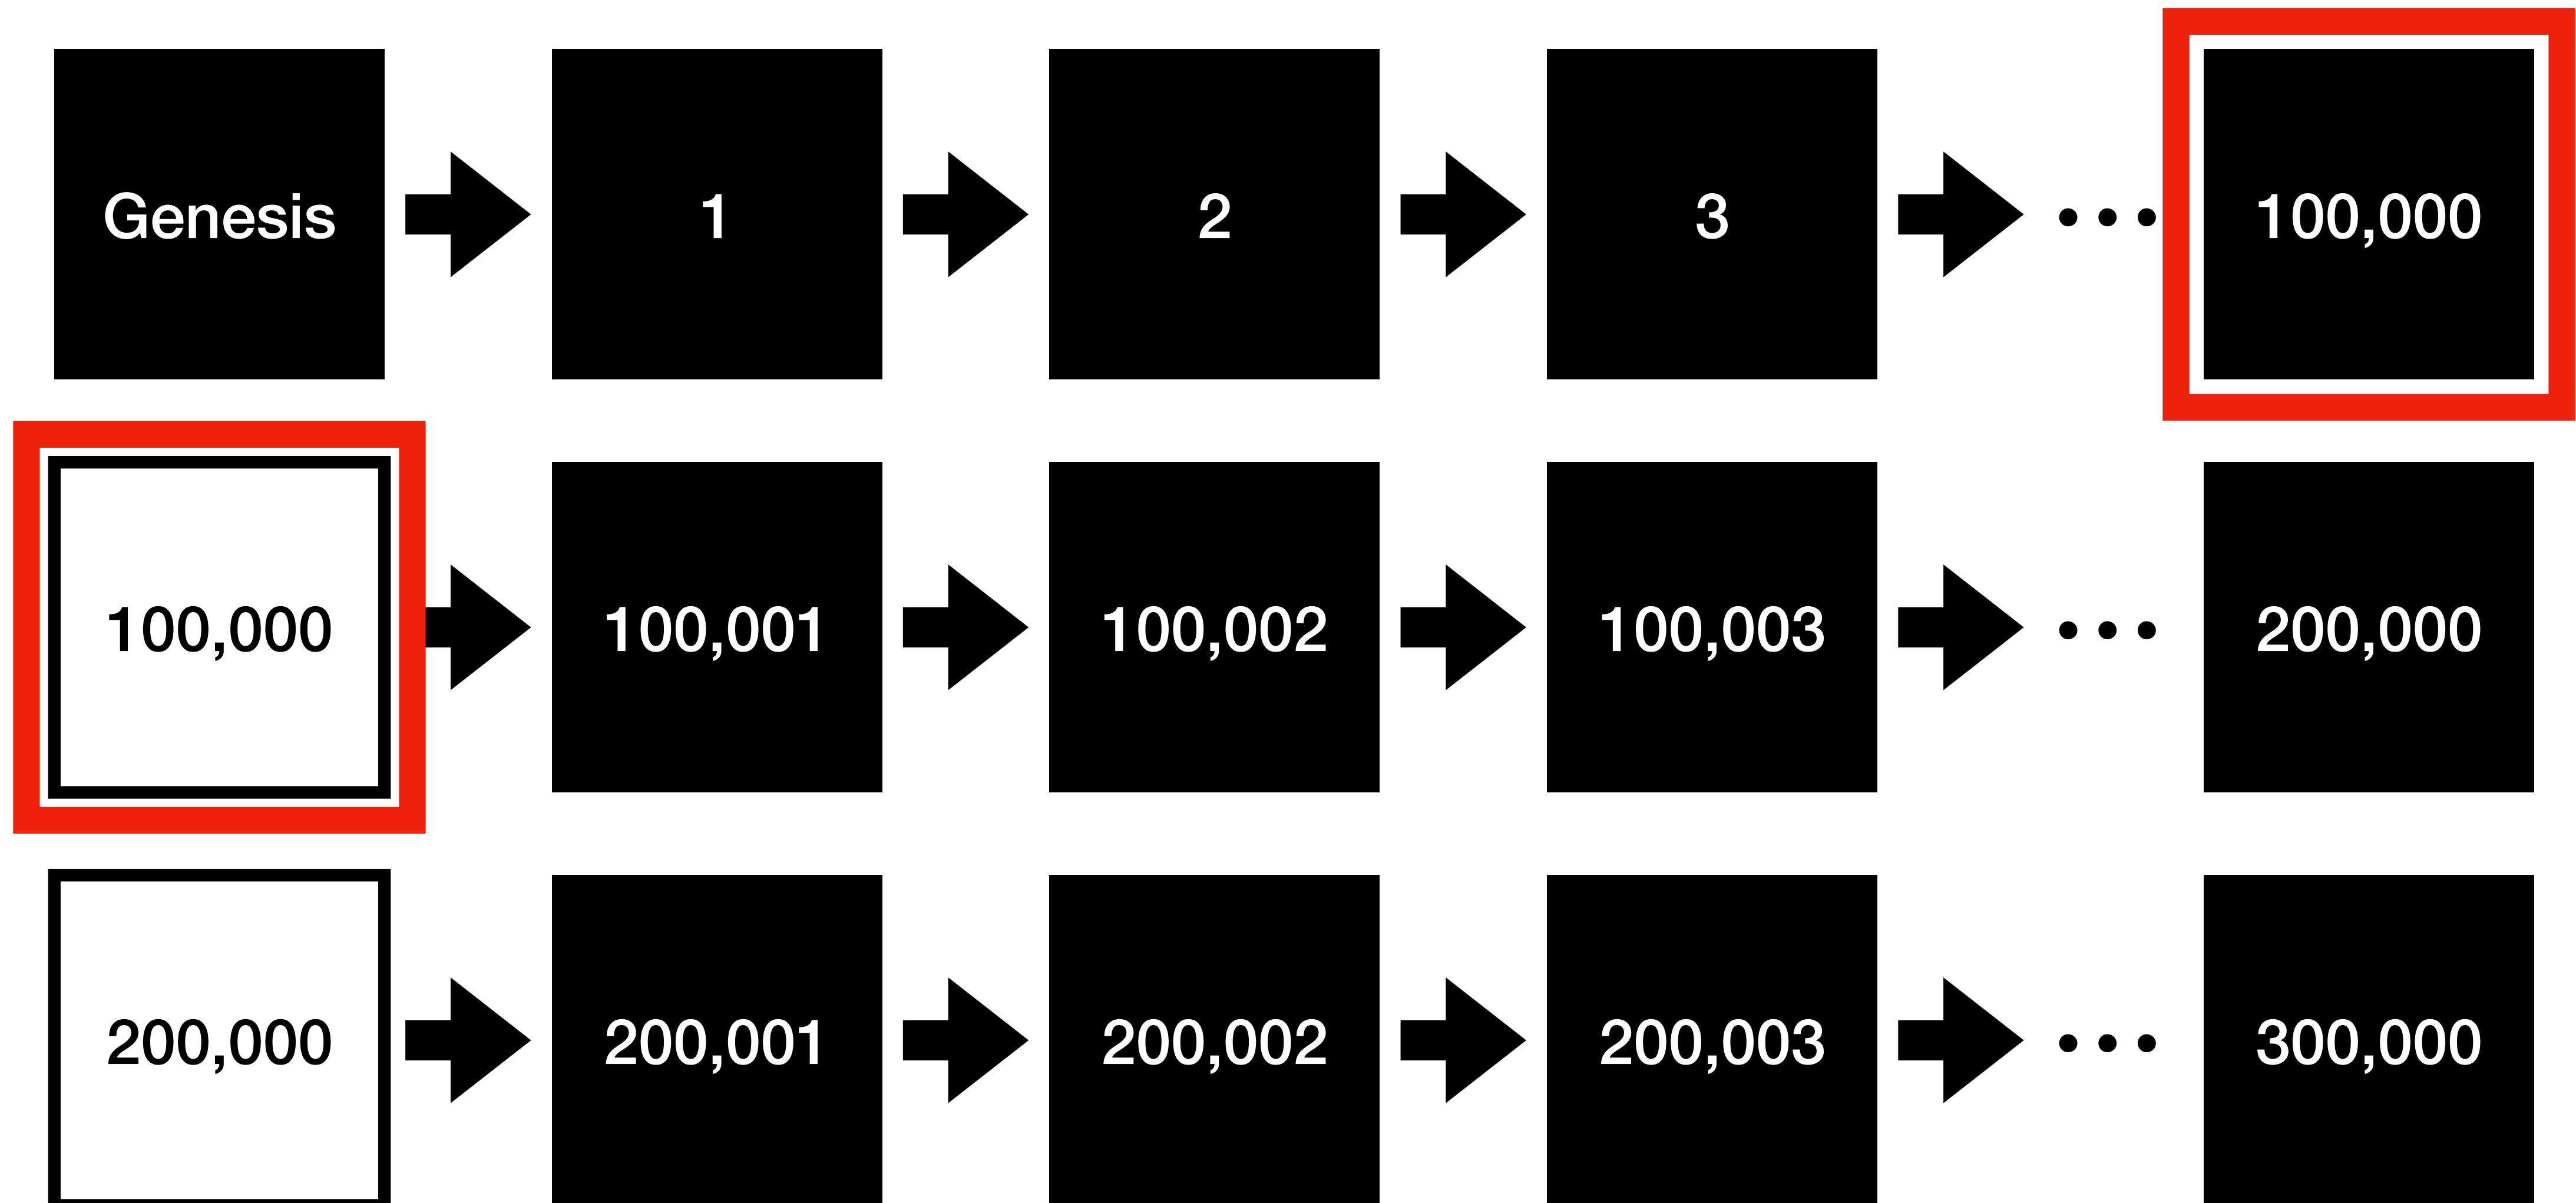
# Out of Order Validation



# Out of Order Validation

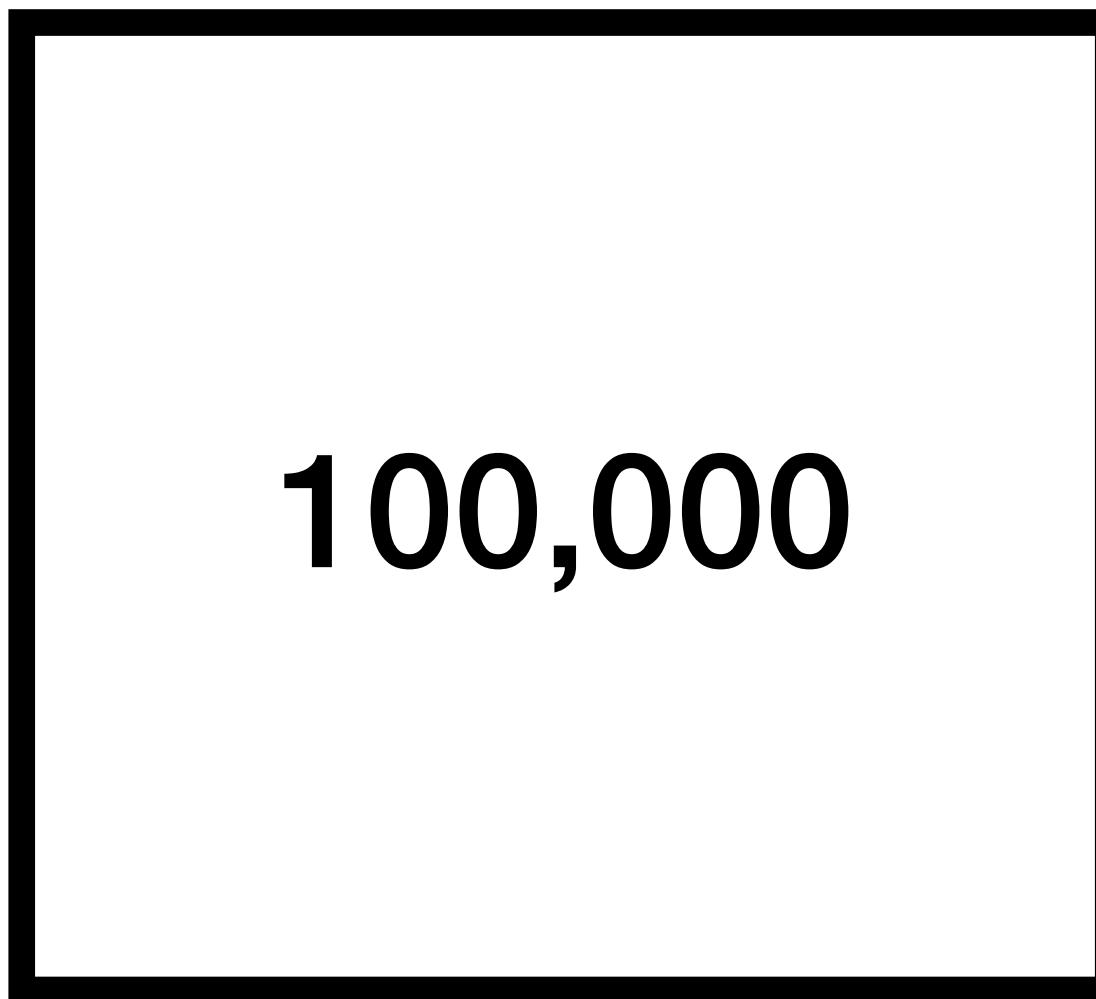


# Out of Order Validation



# Compare Merkle Roots

Check if the generated roots at 100,000 is the same

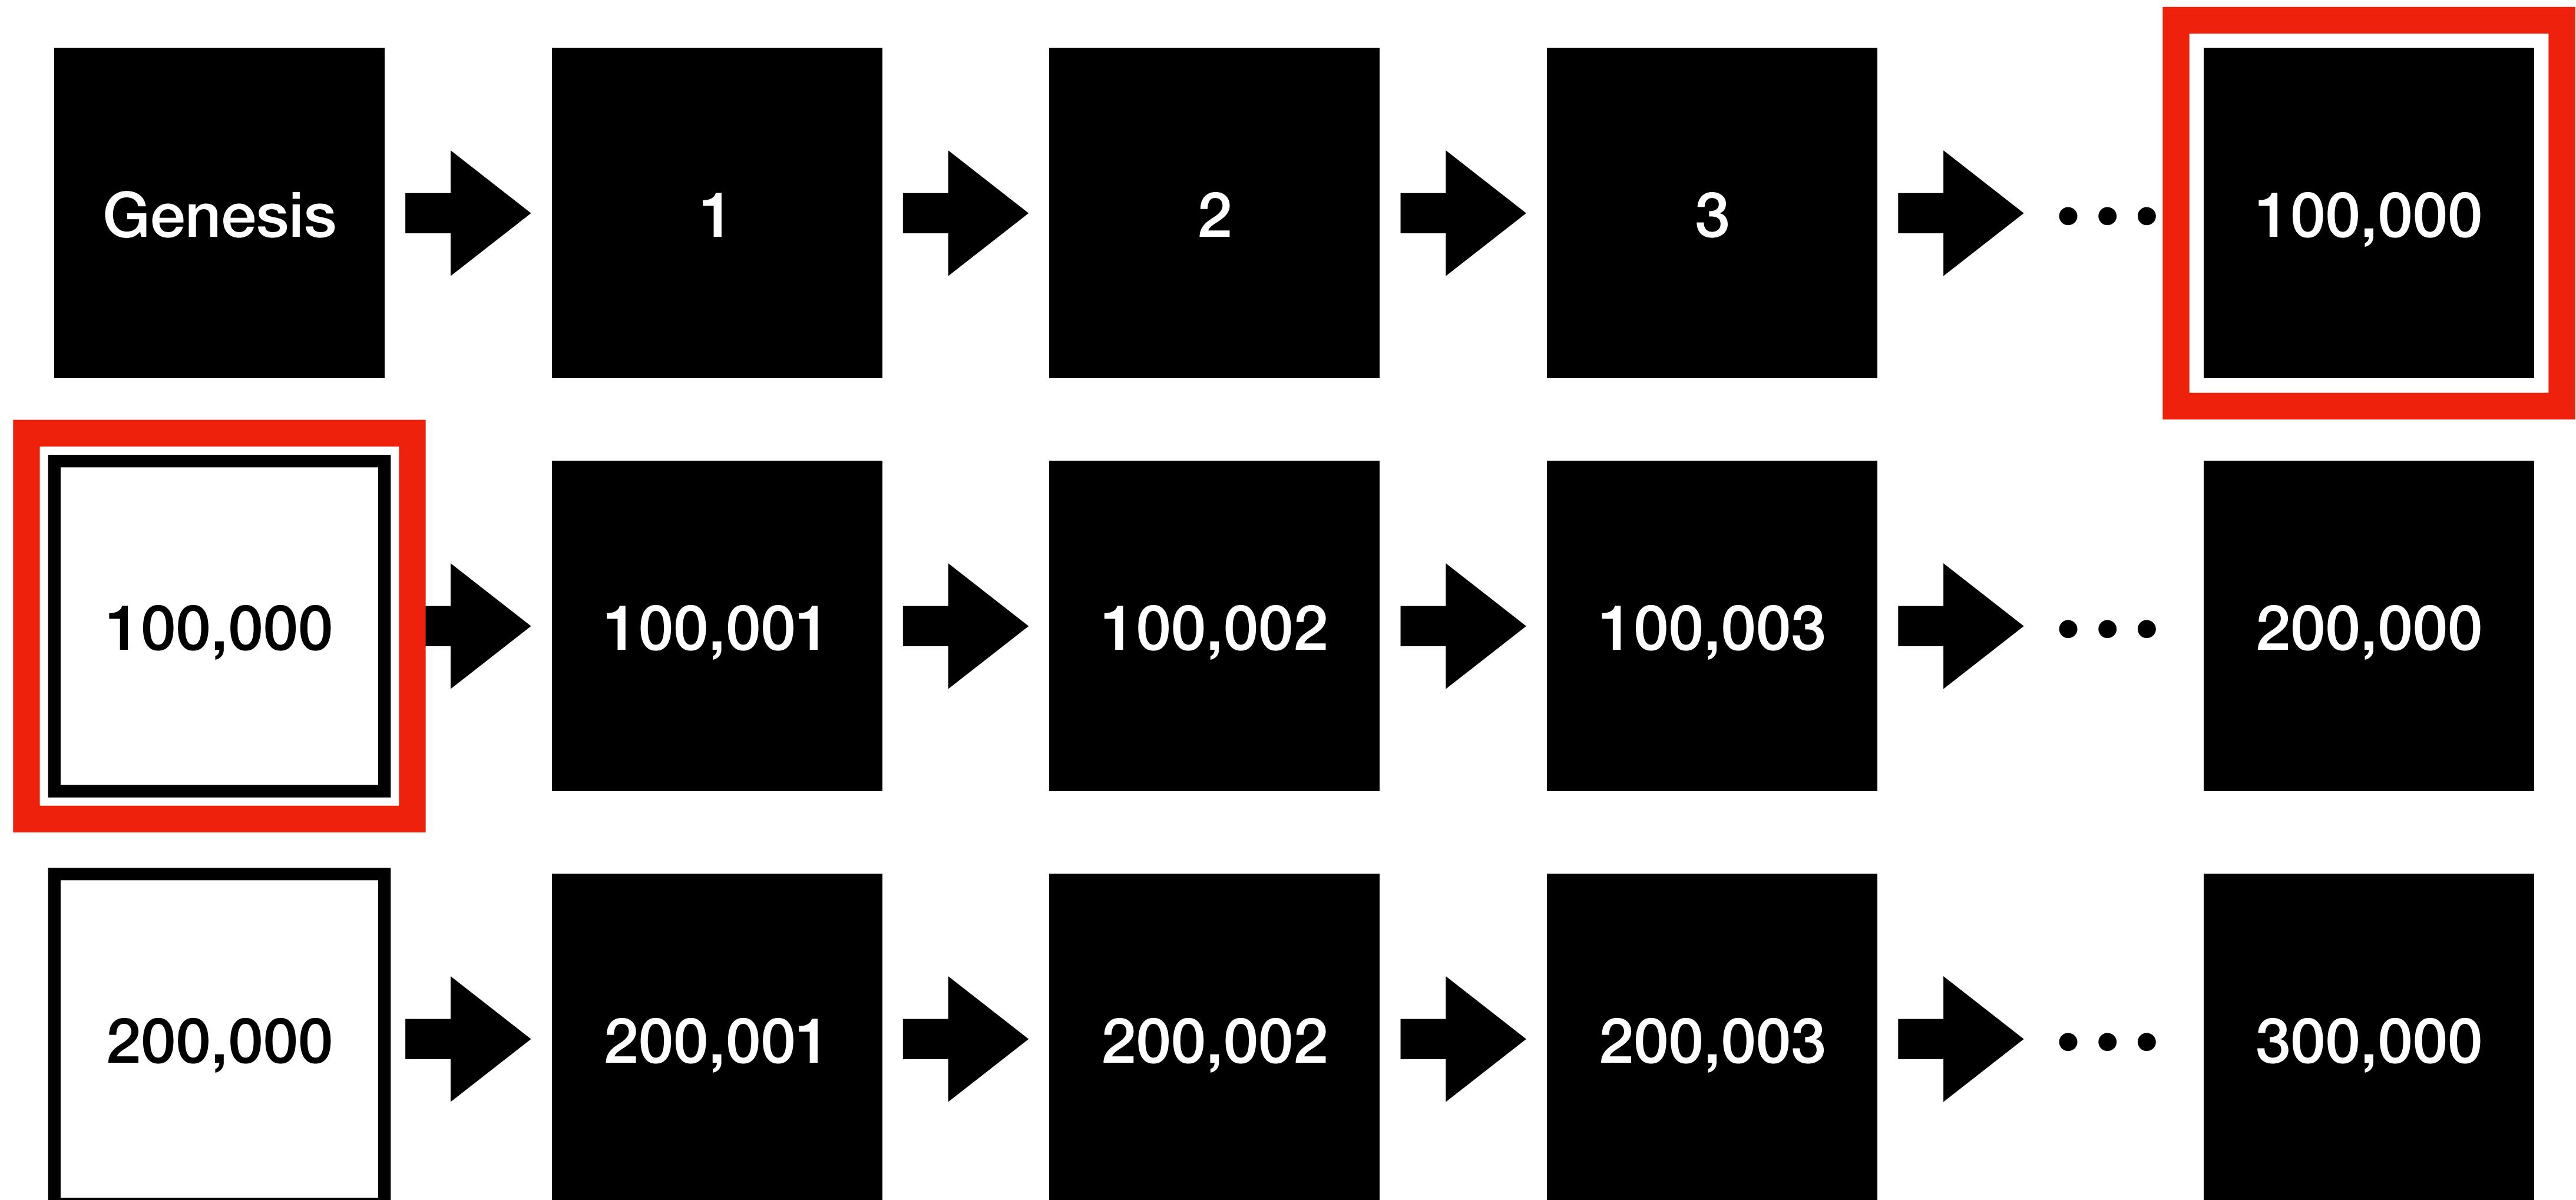


==



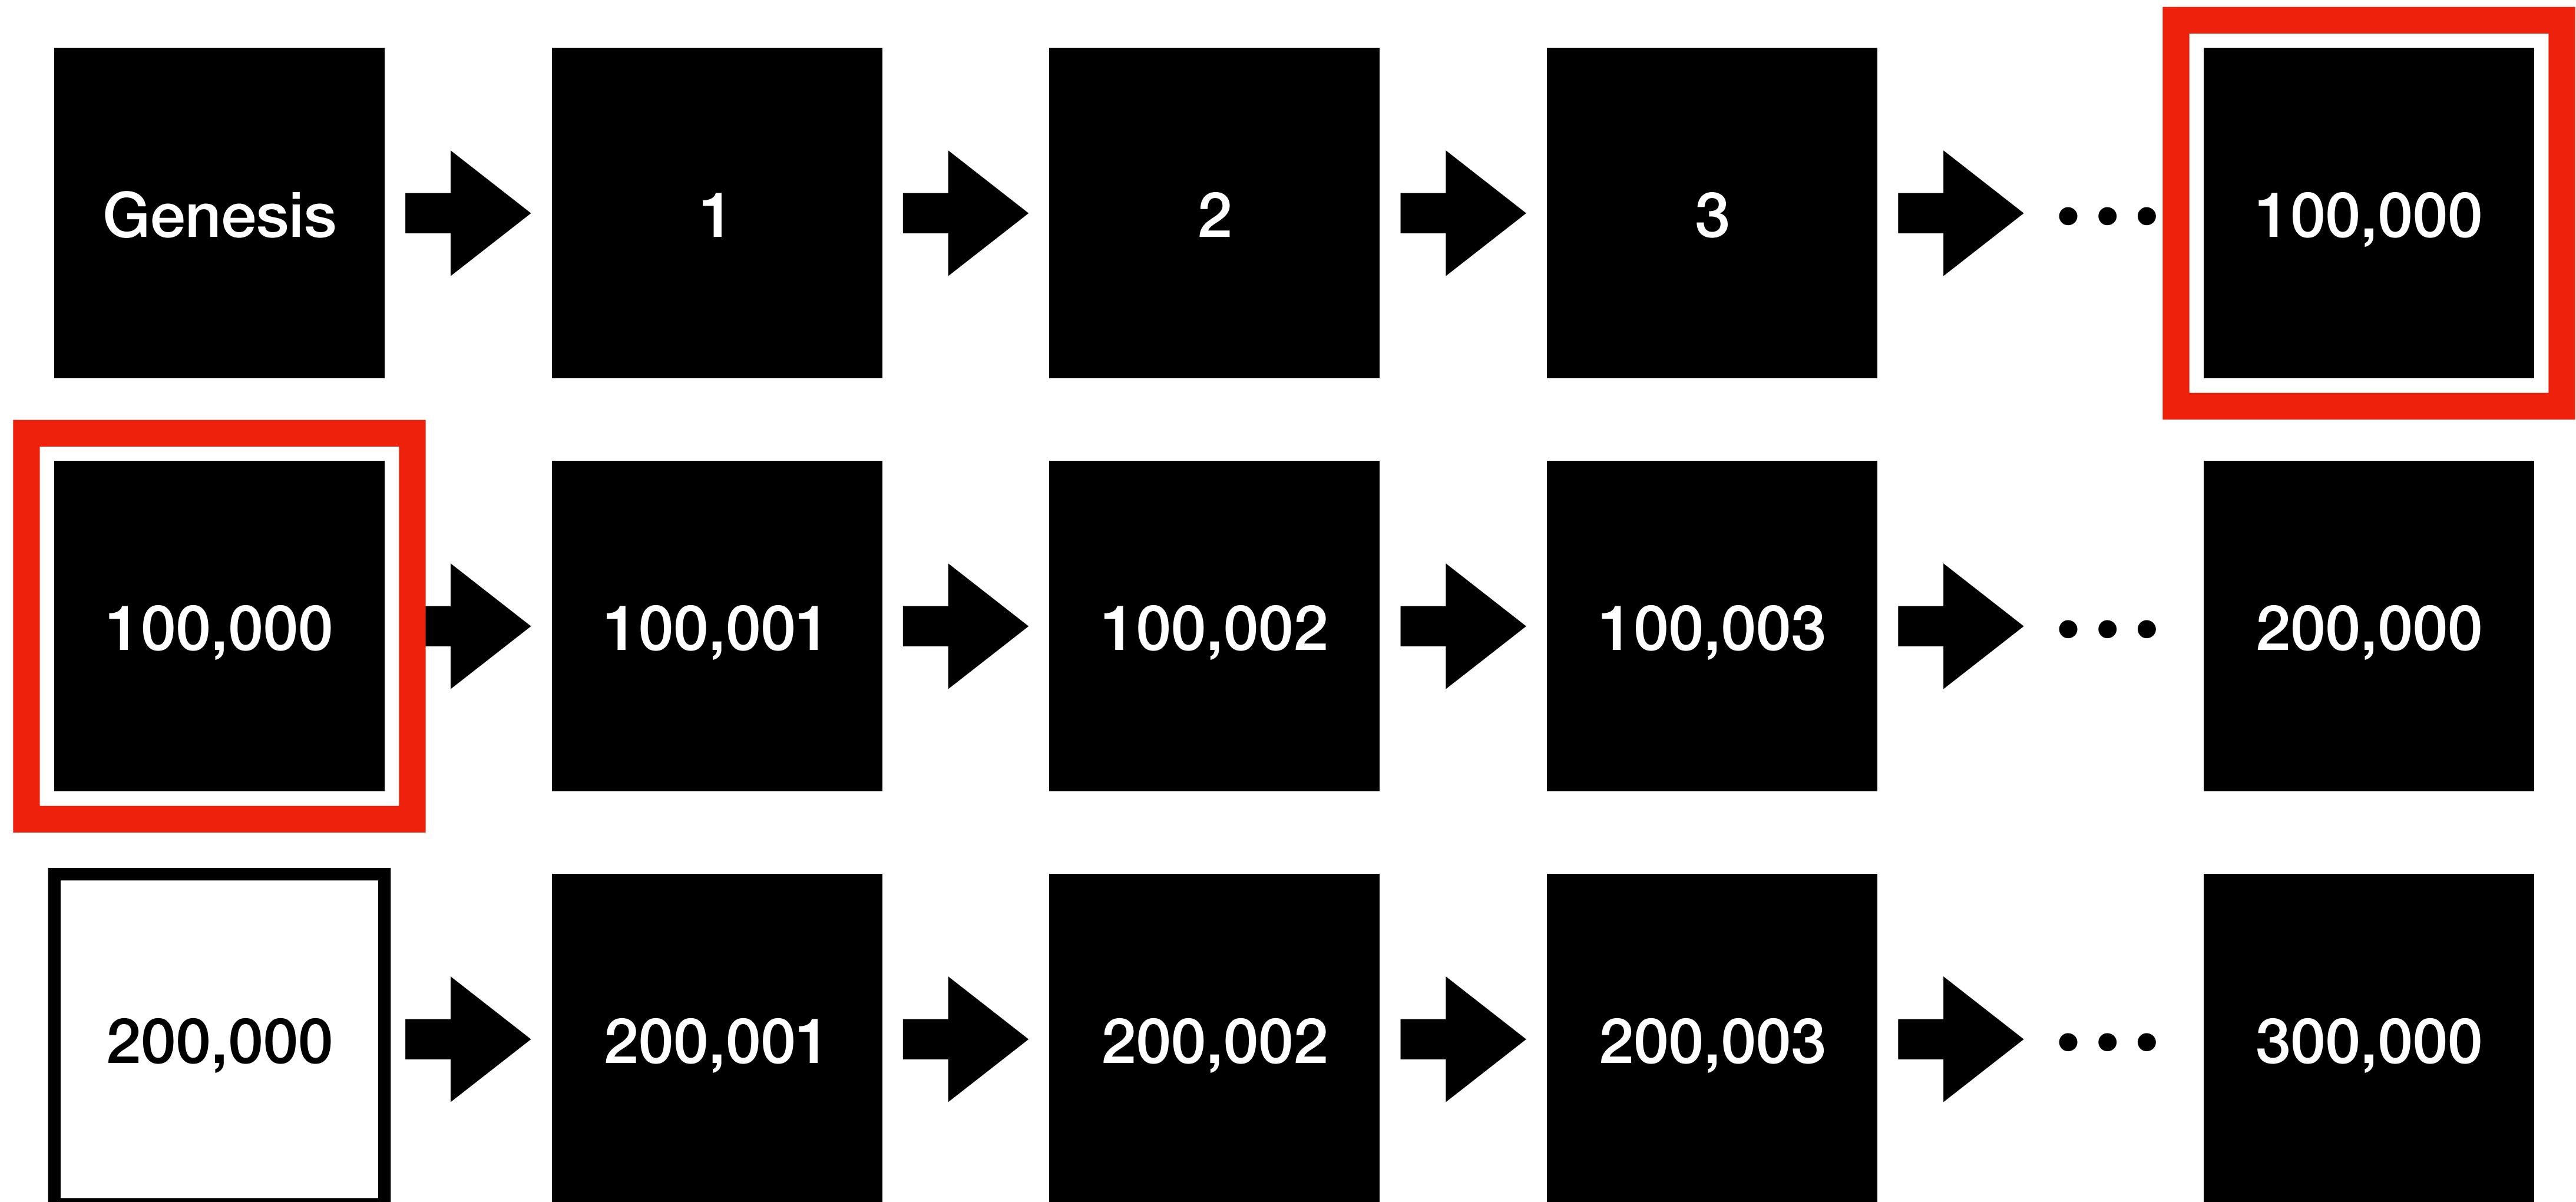
# With Utreexo

## Efficient parallel validation

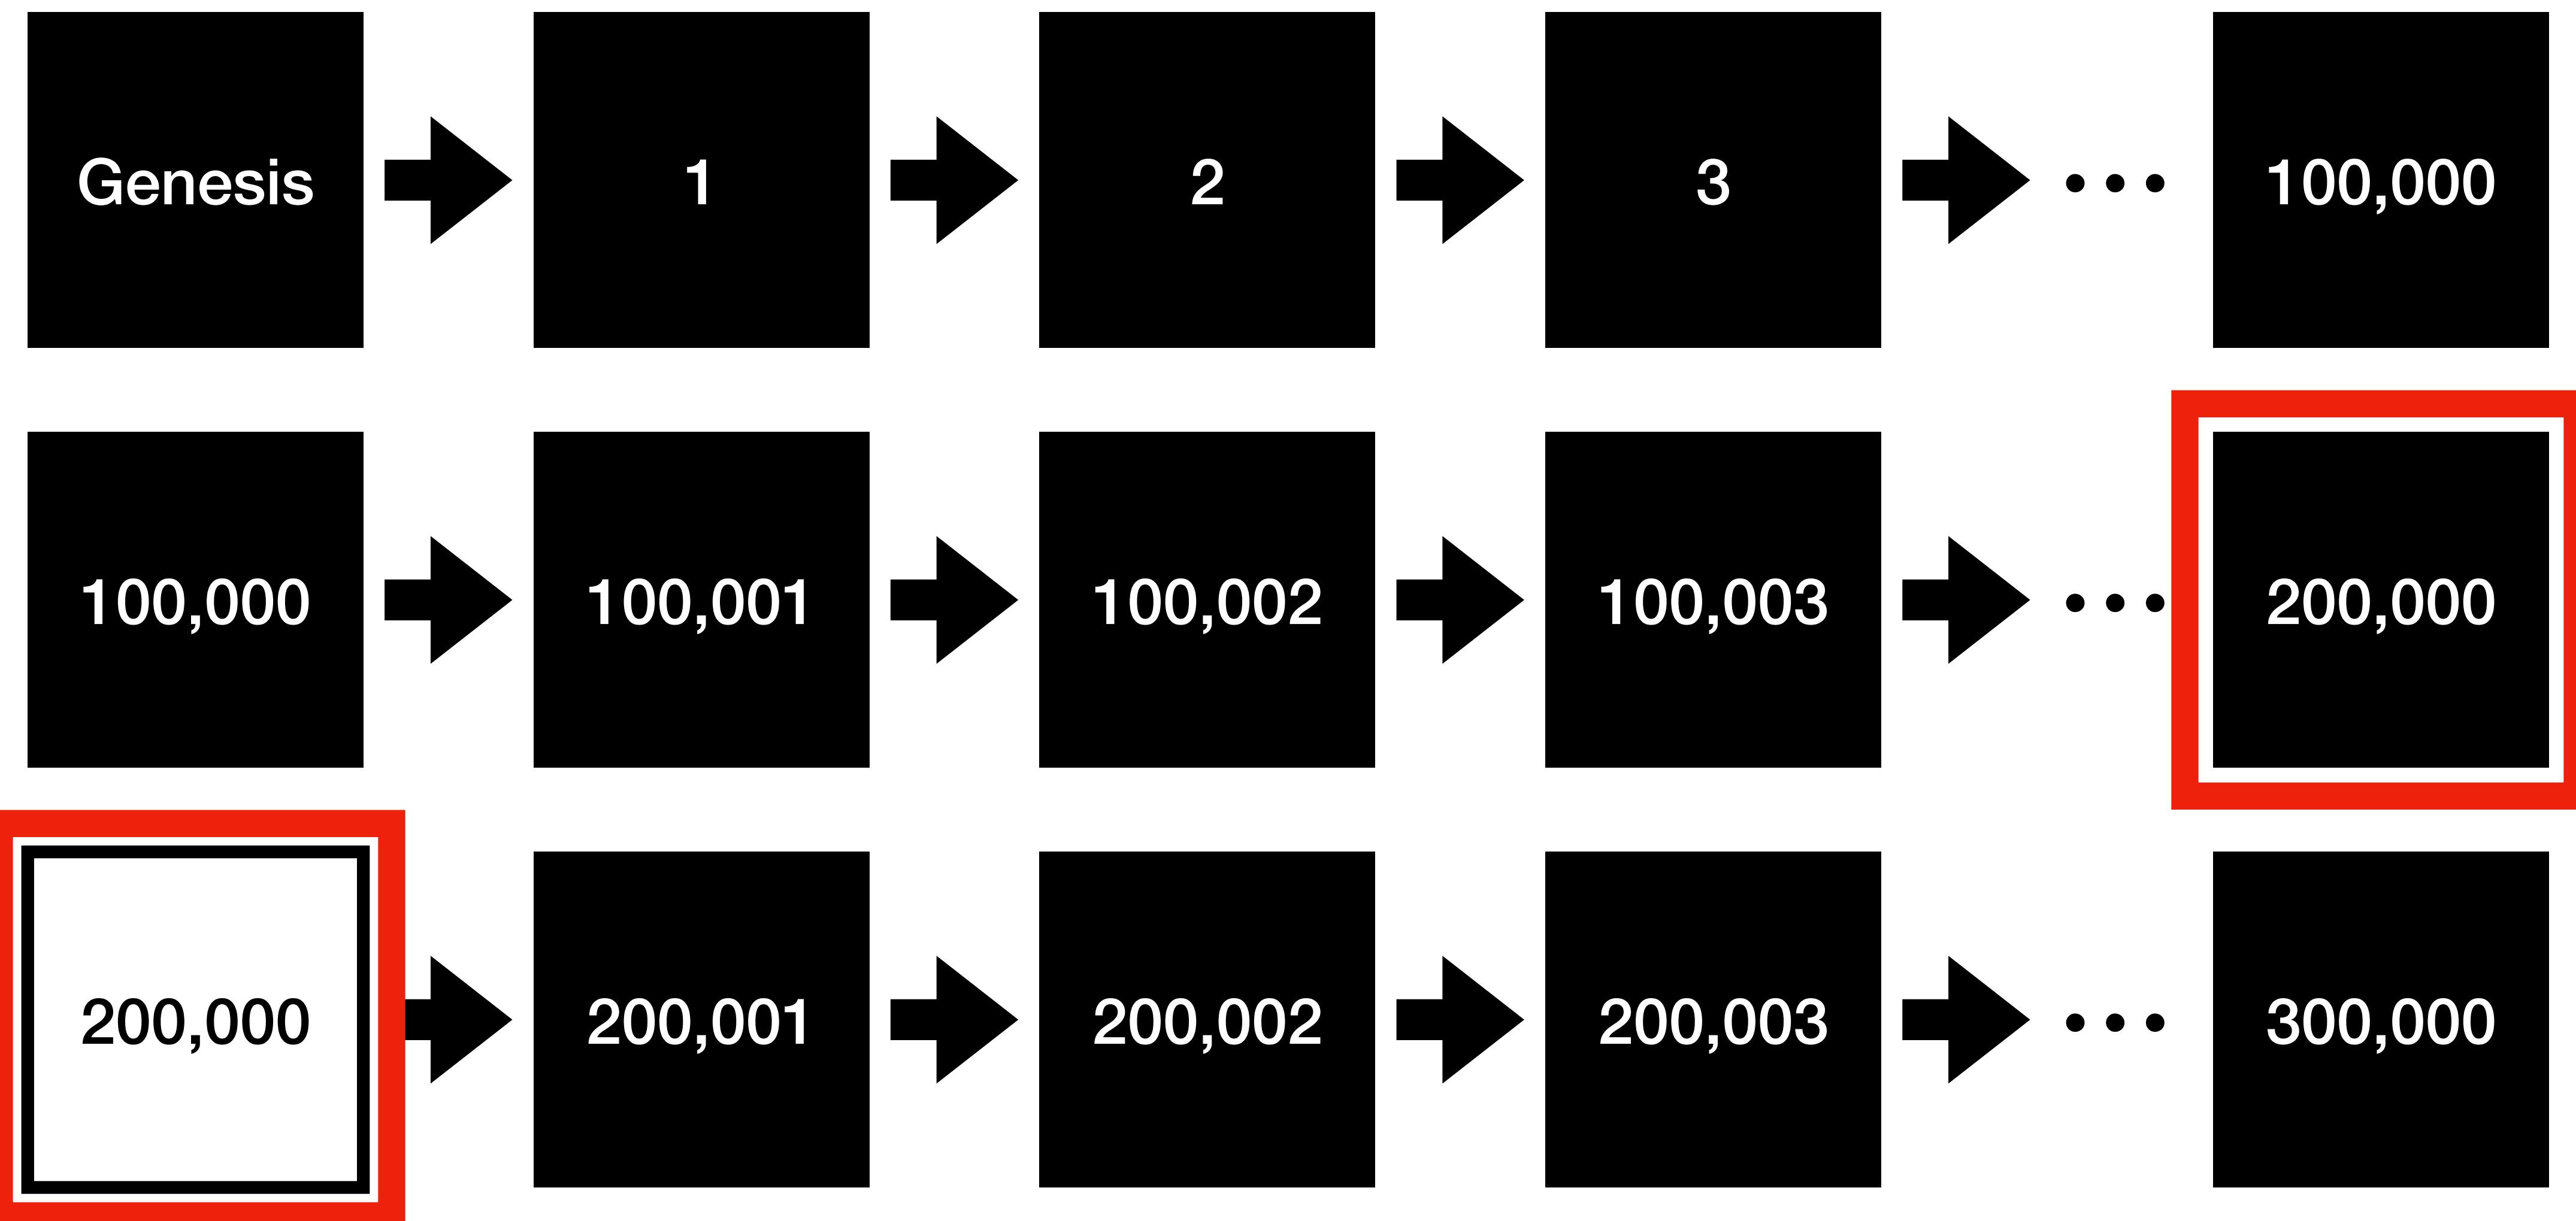


# With Utreexo

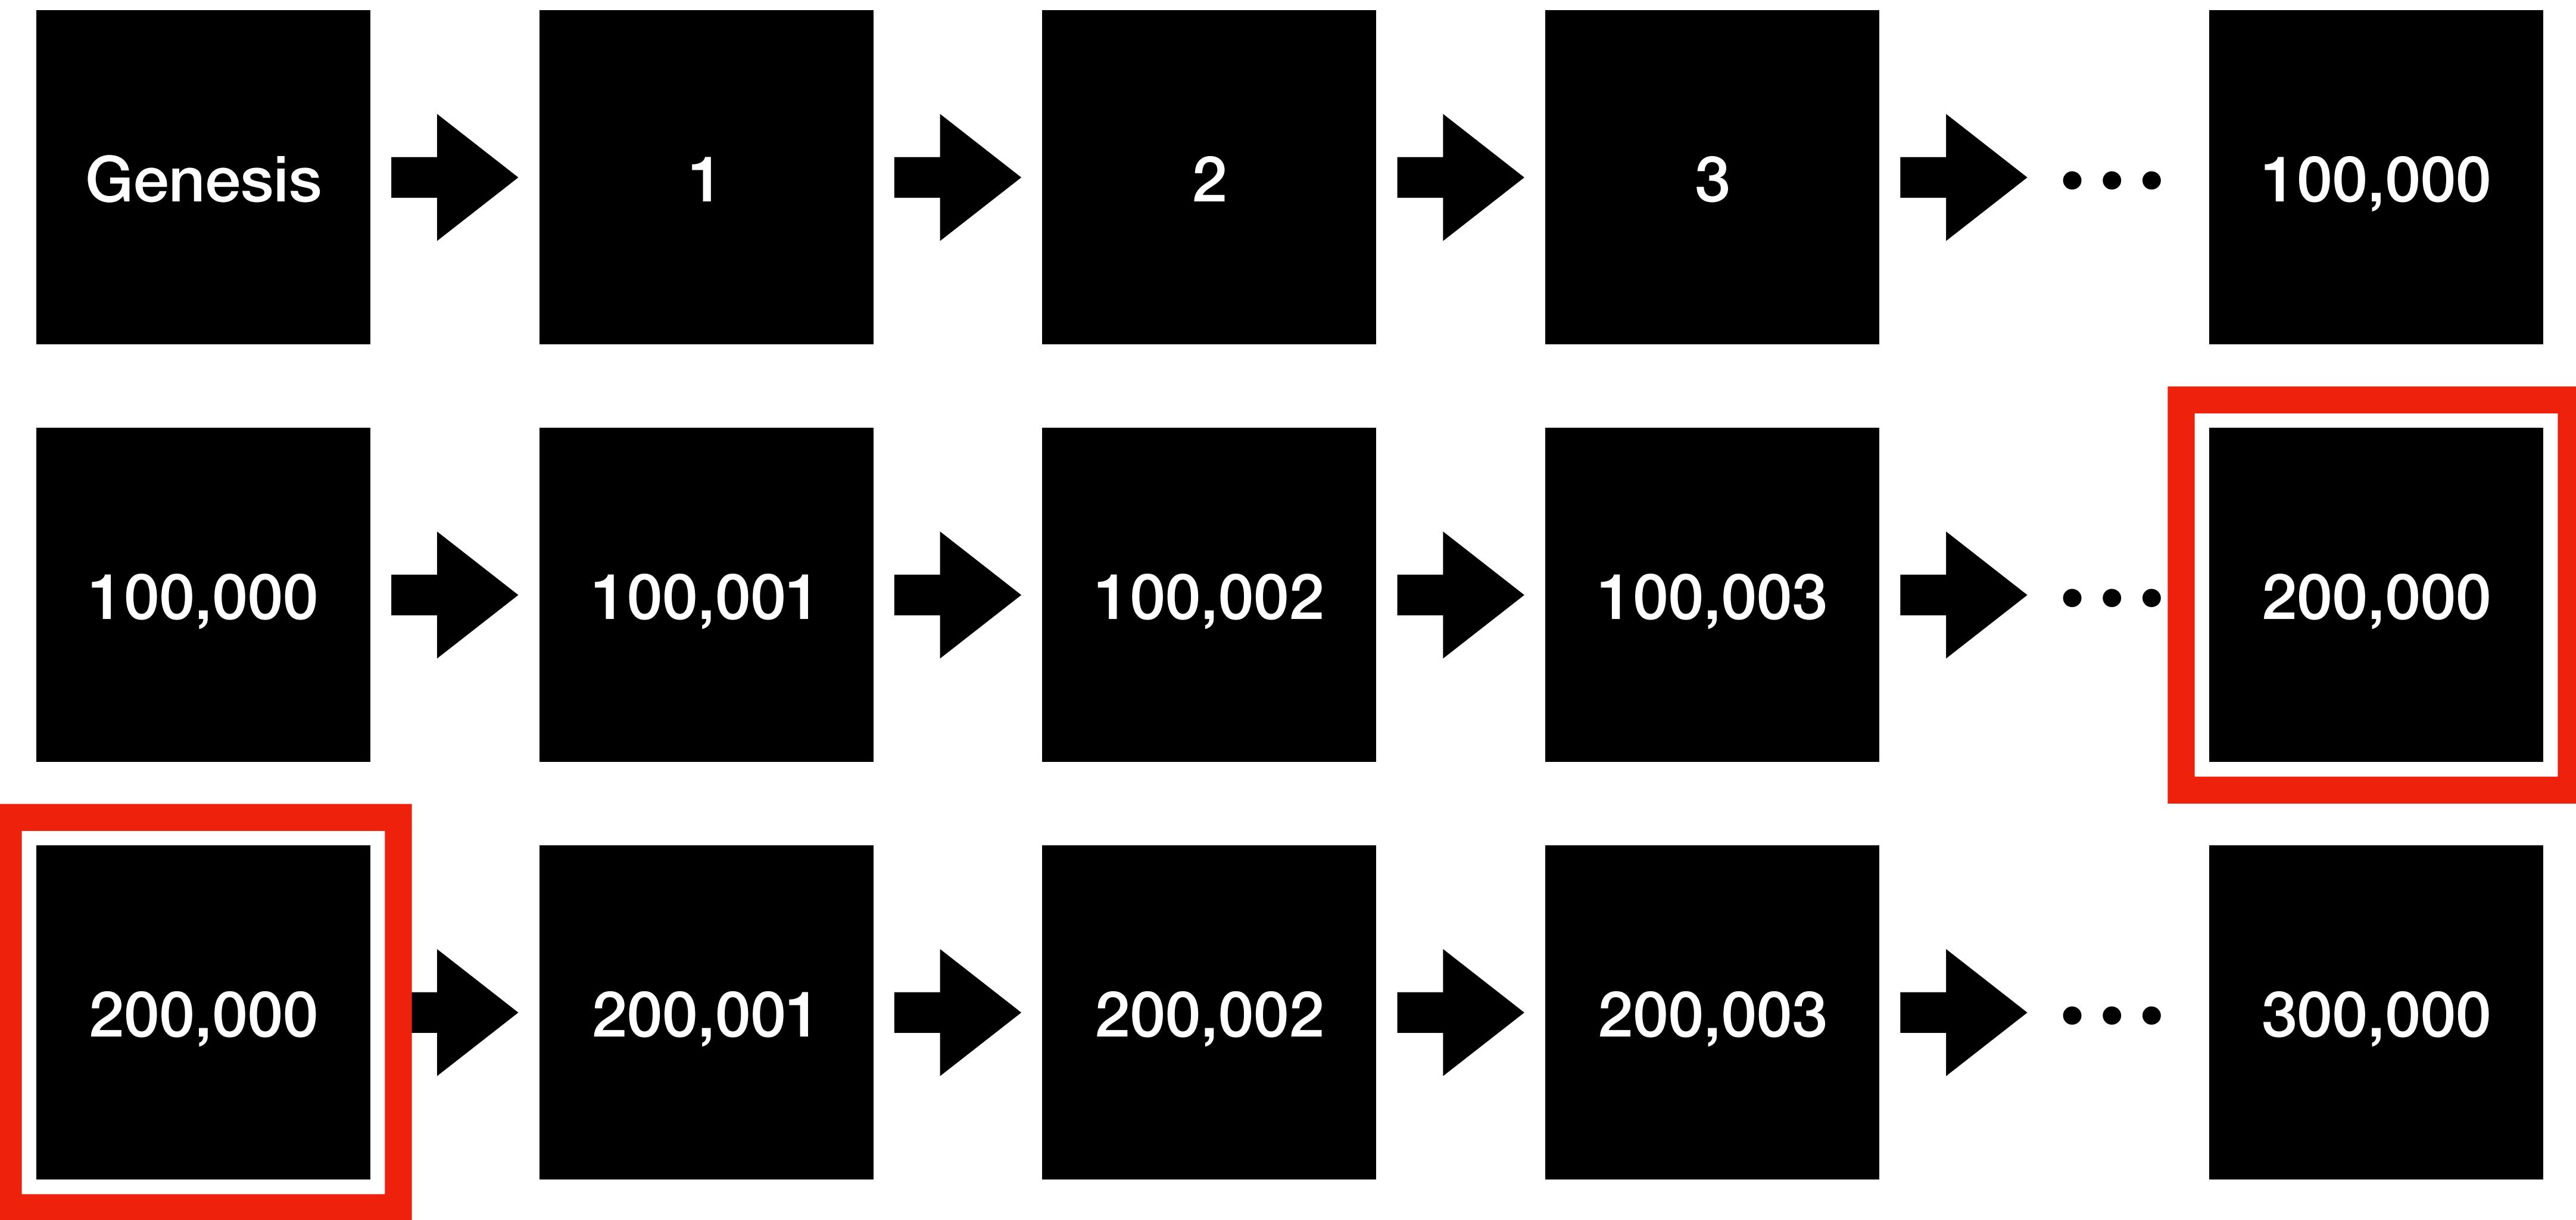
## Efficient parallel validation



# Out of Order Validation

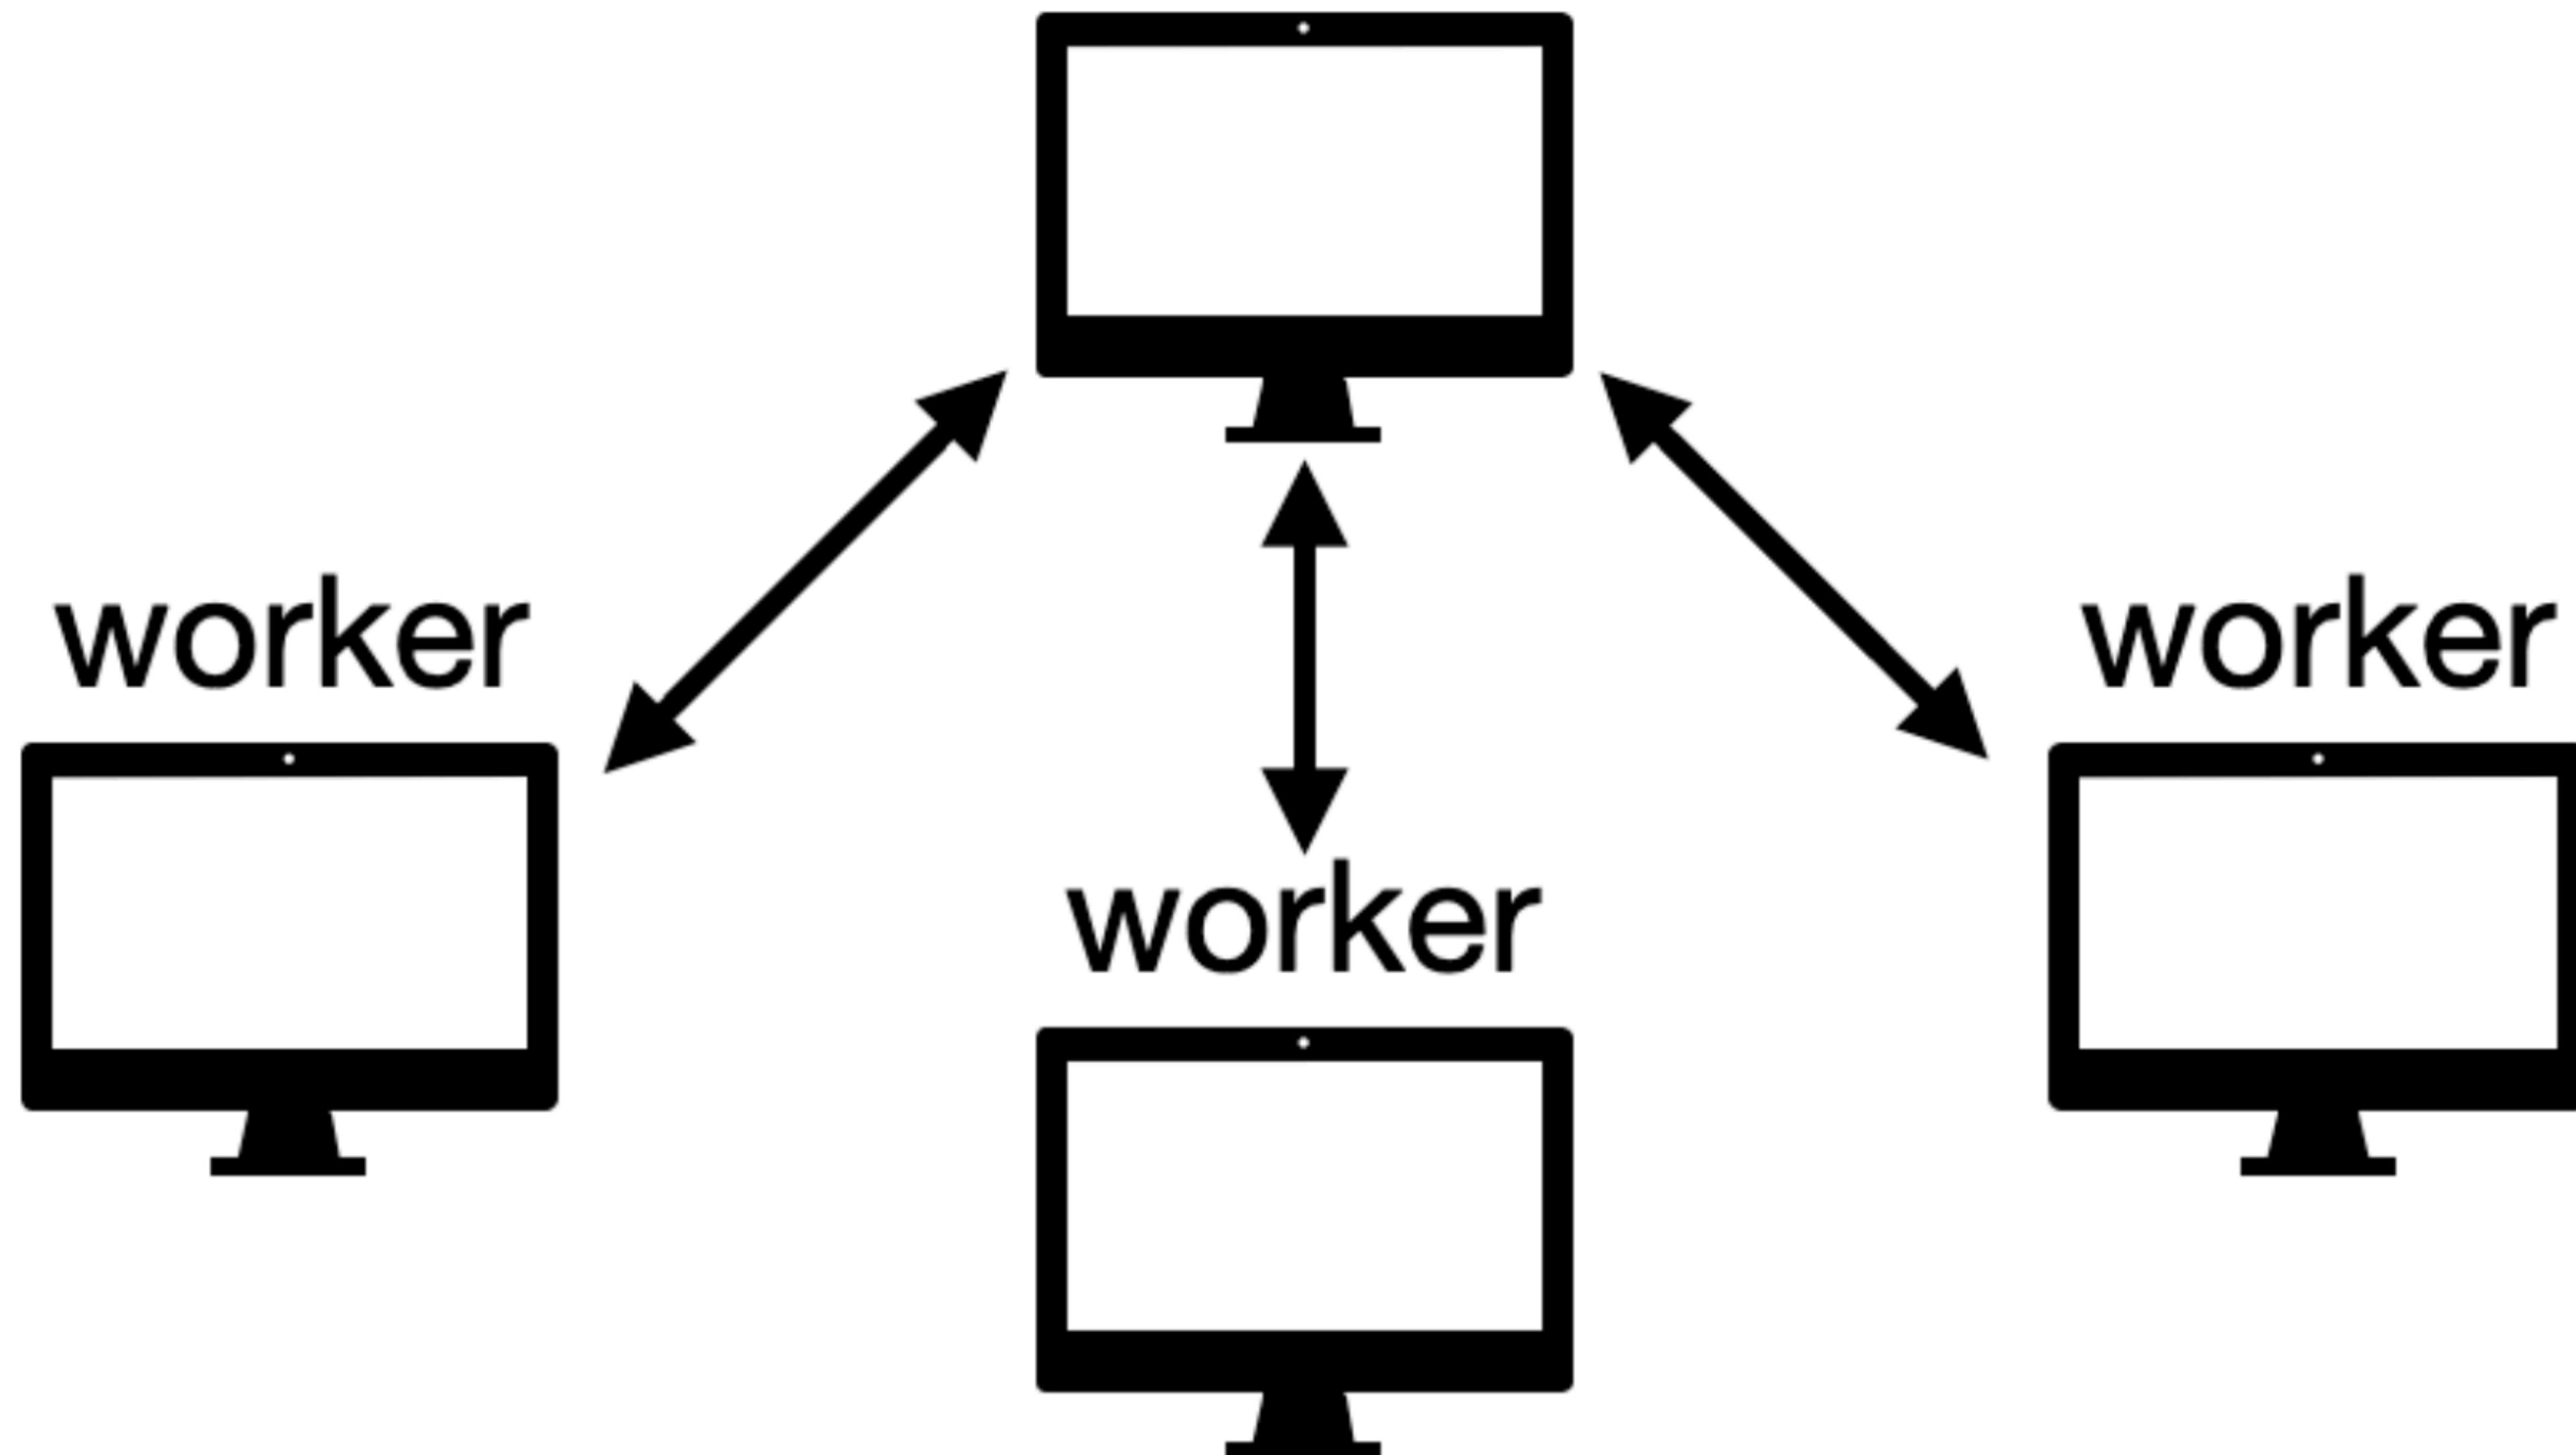


# Out of Order Validation



**~62% faster block  
validation**

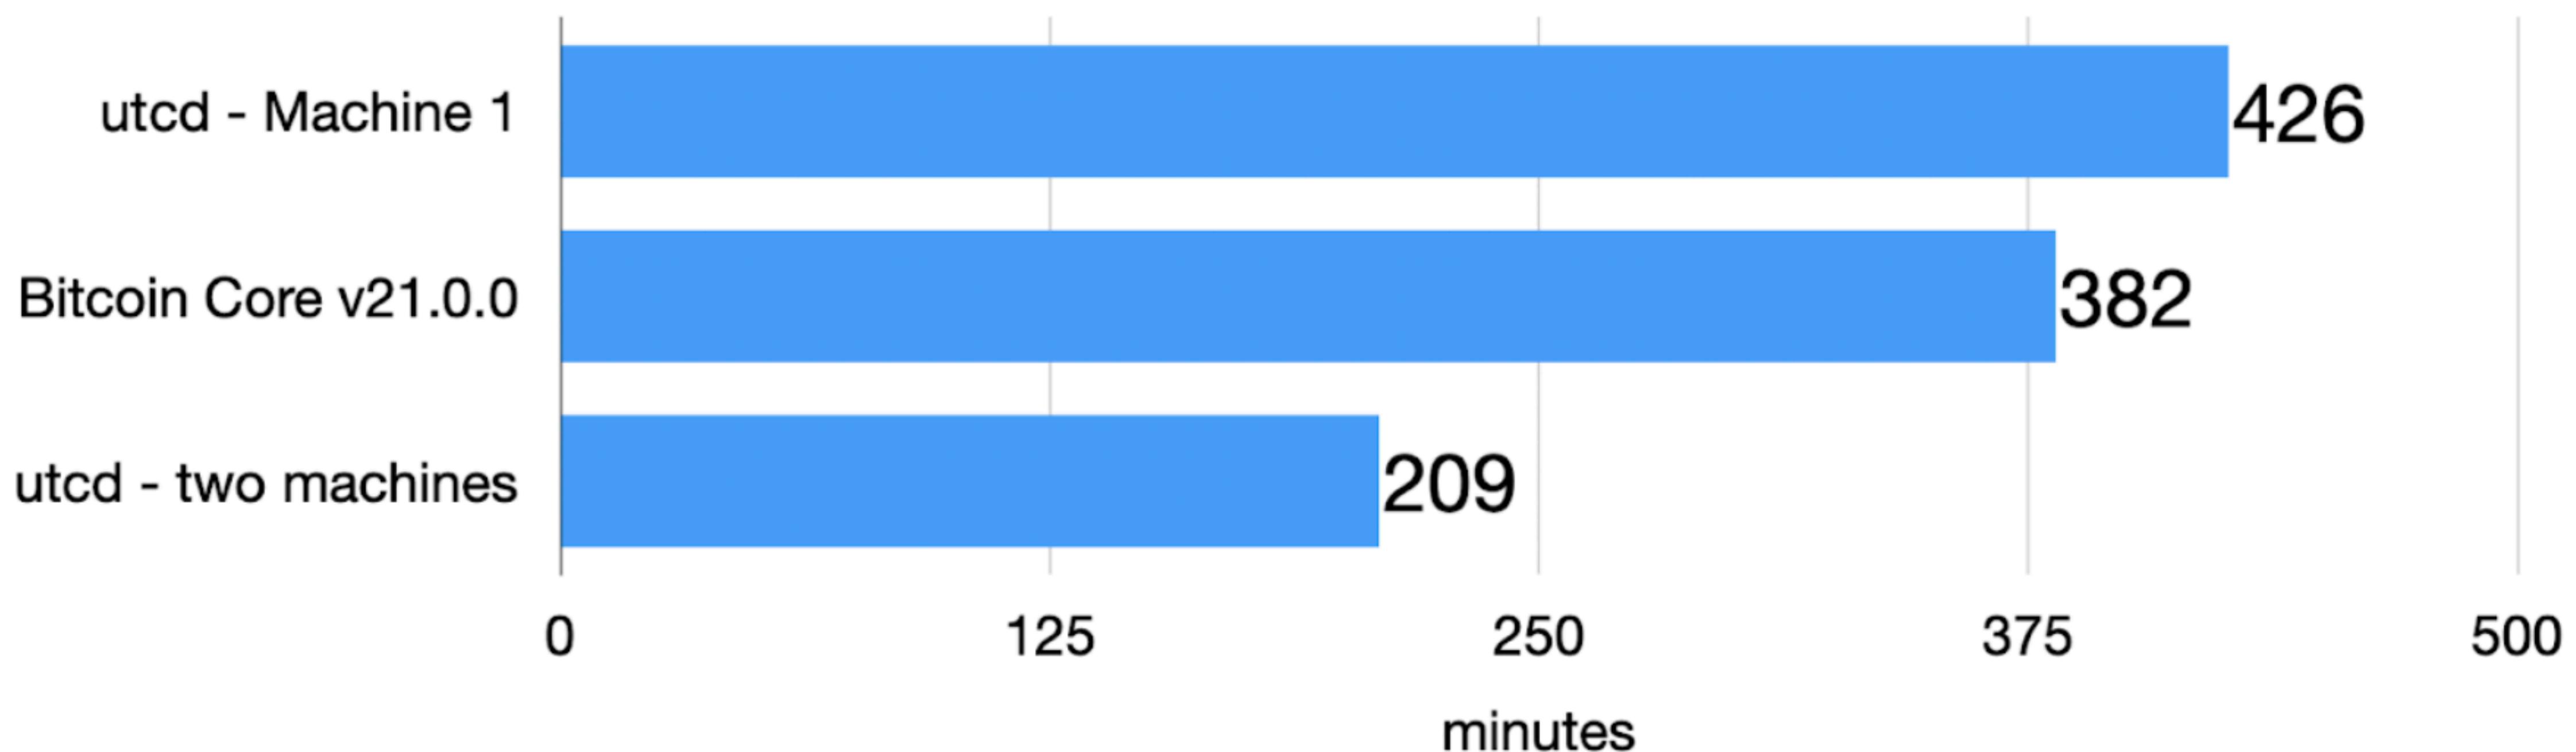
# Coordinator/worker



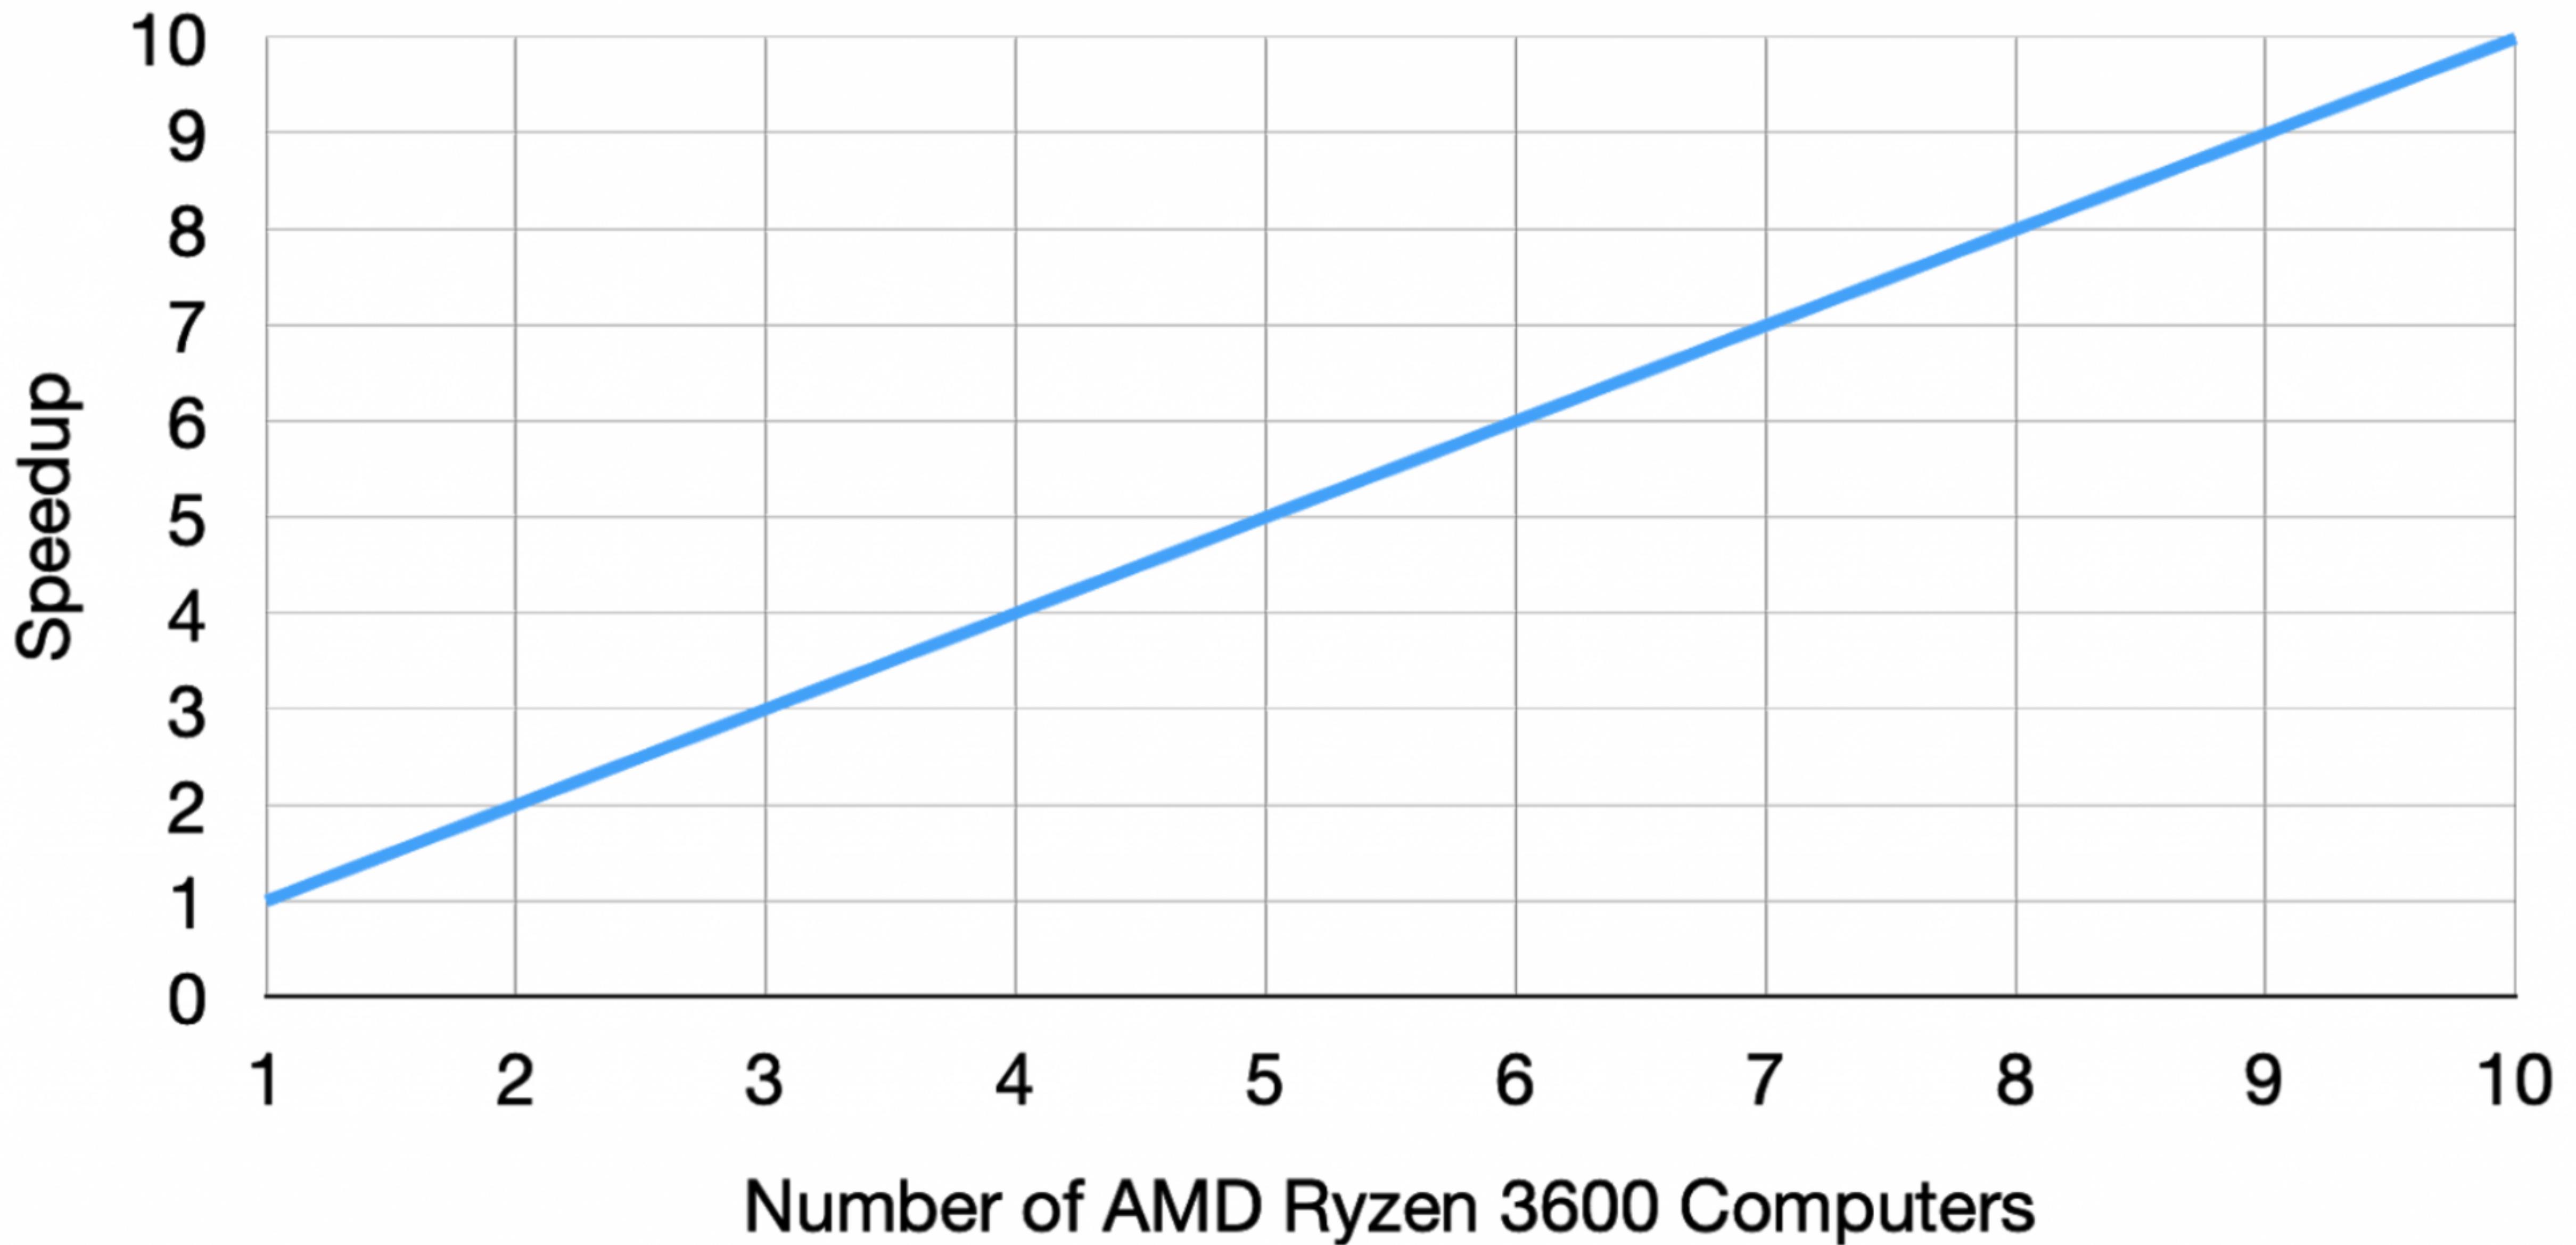
# **2 times faster block validation**

**When tested with 2 computers**

## Initial Block Download Speeds to block 671,000 (Full signature check)



# Utreexo Parallel Node Speedups (based on Amdahl's Law)



# Trade offs

# More internet usage

Extra data needed to be downloaded for Utreexo nodes

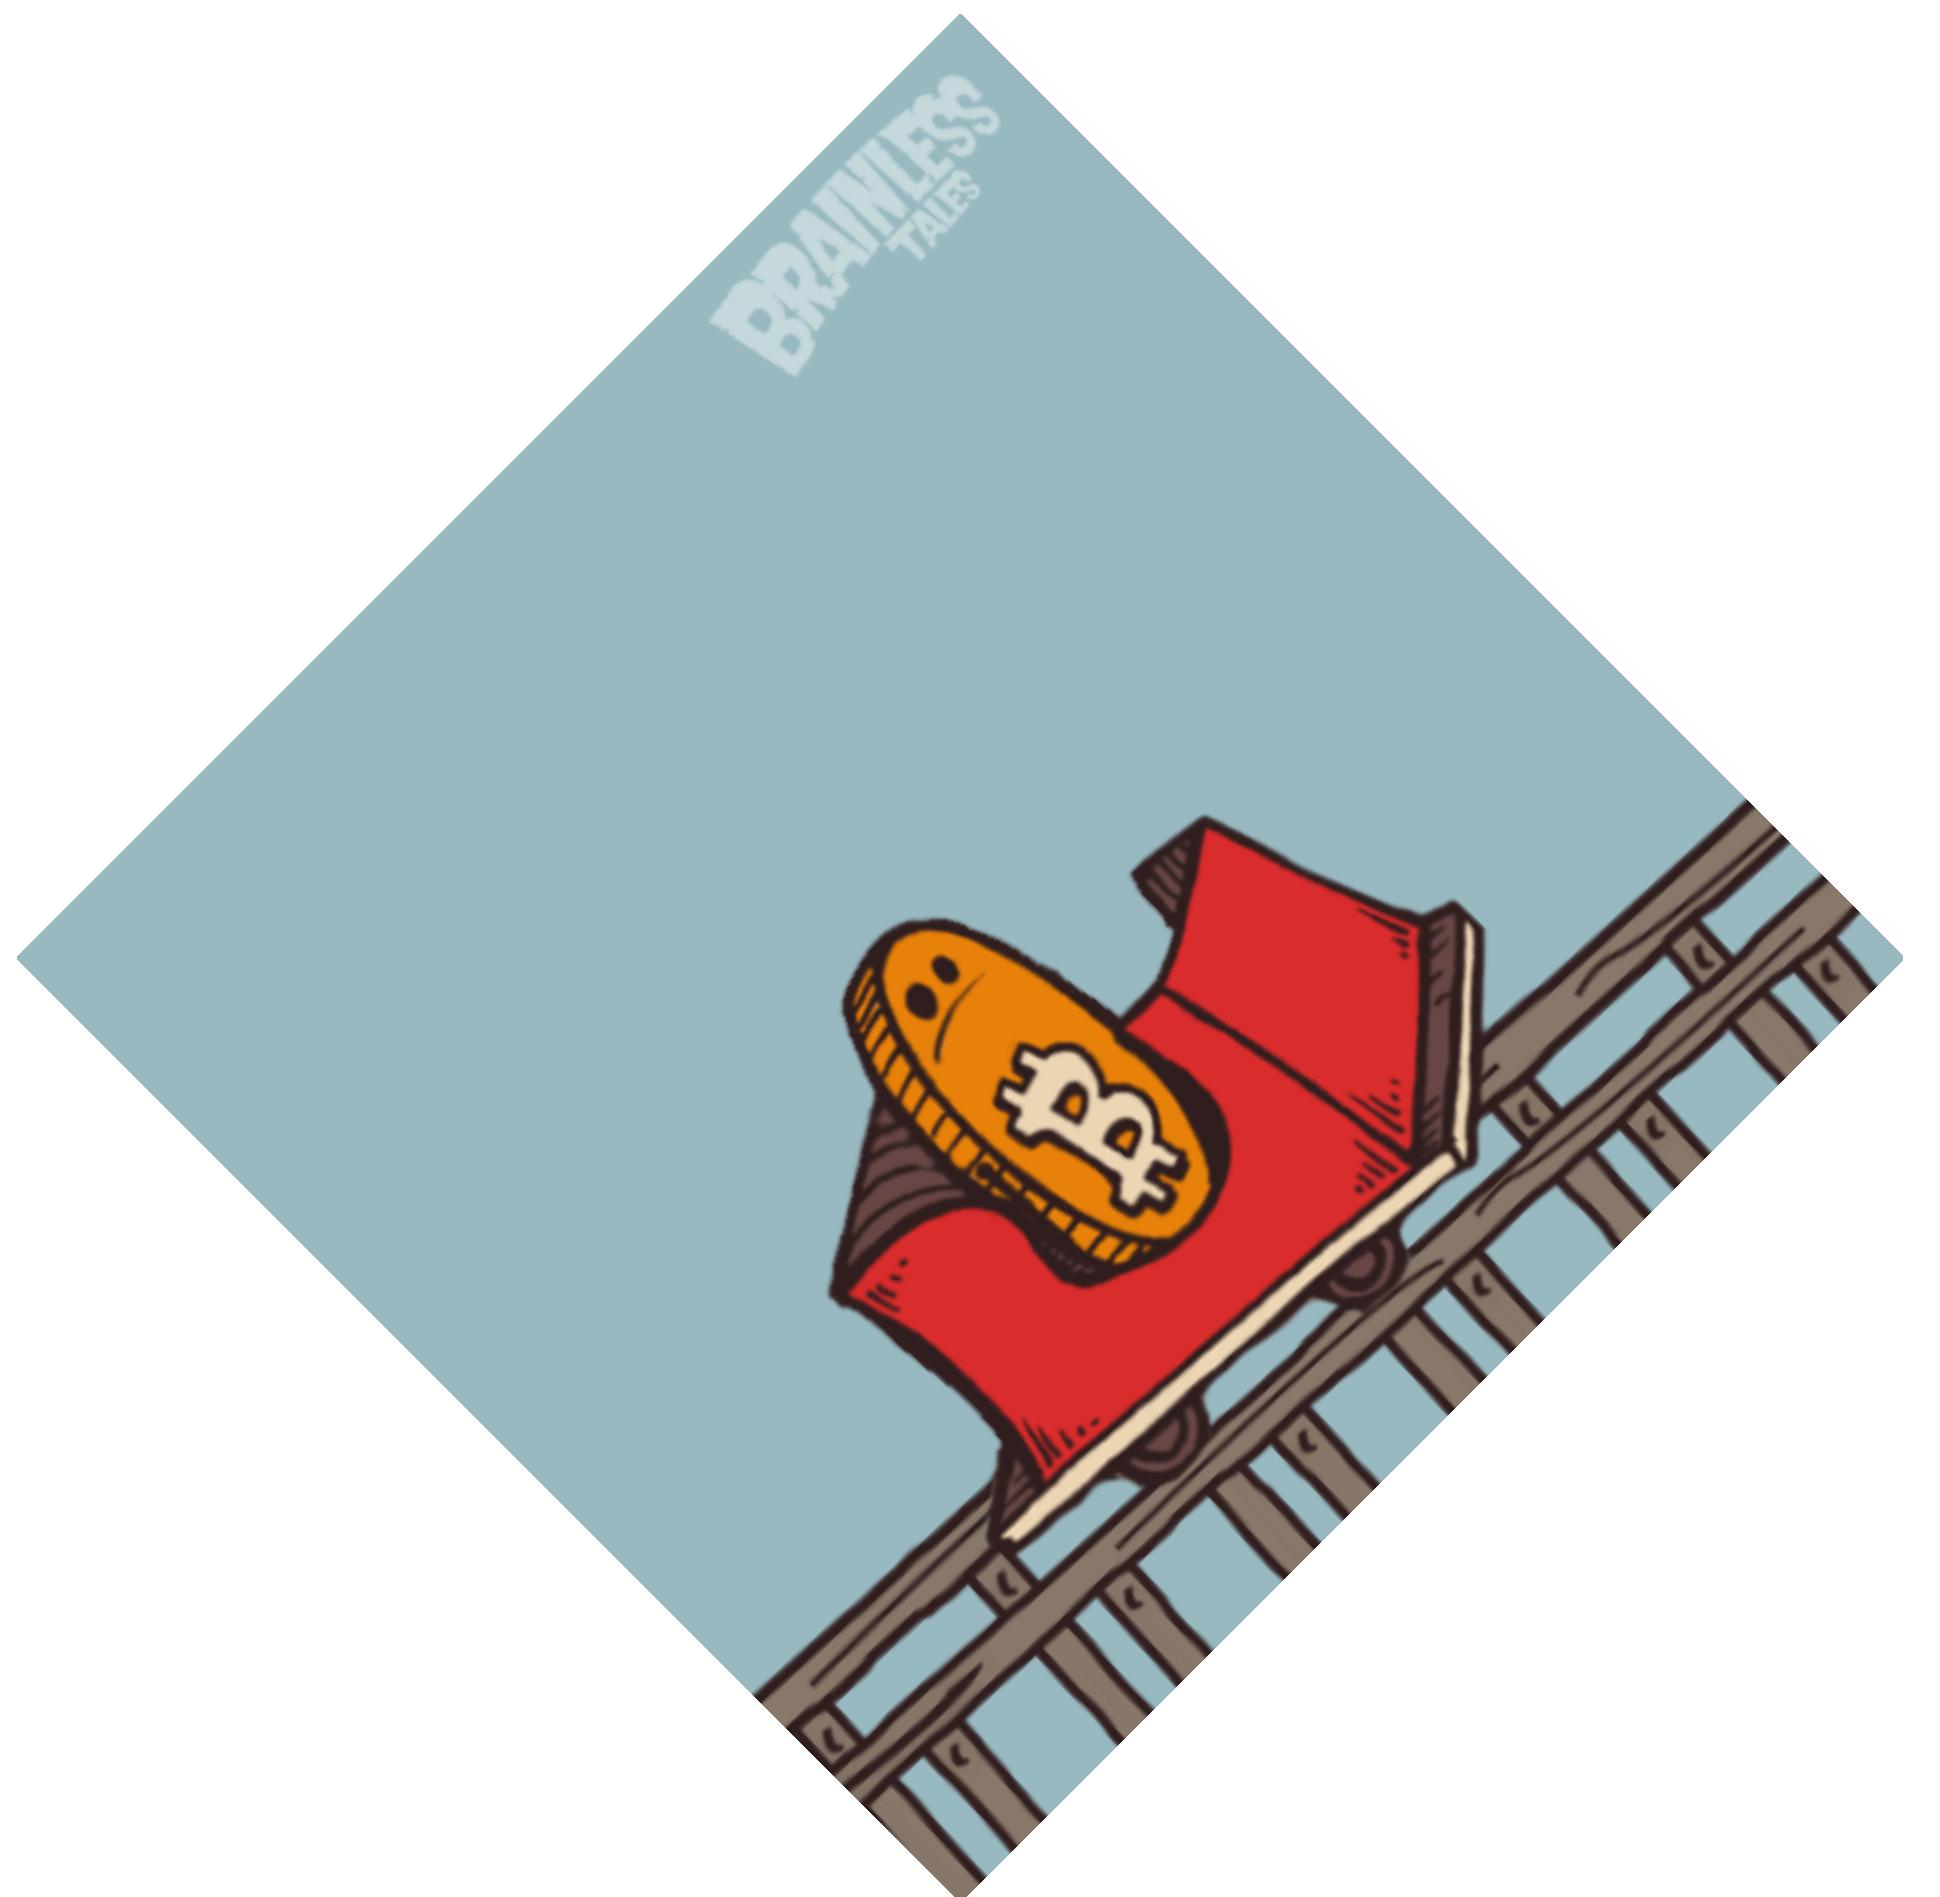
- Need to download 364GB Utreexo proofs (up until height 710,000)

# More internet usage

**Extra data needed to be downloaded for Utreexo nodes**

- Need to download 364GB Utreexo proofs (up until height 710,000)
- Working on caching to reduce the amount

# Internet CPU



# Disk read/write Memory Usage

